

Pump functionality prediction

March 11, 2024

1 Tanzanian faulty pumps prediction

1.1 Problem statement

In Tanzania, access to clean and potable water is essential for the health and well-being of its citizens. However, many water pumps across the country are faulty, leading to water shortages and posing significant health risks to communities. To address this issue and promote access to clean water, we aim to develop a predictive model that can identify faulty water pumps based on various features such as pump age, location, type, and condition. By accurately predicting which water pumps are faulty, authorities and organizations can prioritize maintenance and repair efforts, ensuring that clean and safe water is readily available to all Tanzanians. Stakeholder: The Ministry of Water in Tanzania is a key stakeholder in addressing the issue of faulty water pumps and promoting access to clean and potable water across the country. As the government body responsible for water resource management and infrastructure development, the Ministry plays a crucial role in ensuring that water supply systems are well-maintained and functional. By leveraging predictive modeling to identify faulty water pumps, the Ministry can efficiently allocate resources for maintenance and repair activities, thereby improving the reliability and accessibility of clean water for Tanzanian communities.

1.1.1 Objectives

1. To predict the functionality of water pumps: Develop a predictive model to classify water pumps into functional, non-functional, and functional needs repair categories based on various features such as `amount_tsh`, `gps_height`, `waterpoint_type`, and others.
2. To identify factors influencing water pump functionality: Conduct exploratory data analysis to identify the key factors (e.g., funder, installer, water quality) that influence the functionality of water pumps and their maintenance needs.
3. To optimize water pump maintenance strategies: Use historical data on water pump failures and repairs to optimize maintenance schedules and resource allocation, ensuring timely repairs and minimizing downtime of water pumps.
4. To assess the geographical distribution of water pump functionality: Analyze the geographical distribution of functional and non-functional water pumps to identify regions with high repair needs and prioritize interventions for improved access to clean water.
5. To evaluate the impact of funding sources on water pump functionality: Investigate the relationship between funding sources and water pump functionality to assess the effectiveness of different funding mechanisms in ensuring sustainable access to clean water.

1.2 Data understanding

```
[1]: # import relevant modules
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import random
from scipy import stats
from scipy.stats import chi2_contingency
from scipy.stats import f_oneway
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn import tree
from matplotlib.colors import LinearSegmentedColormap
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

import warnings
warnings.filterwarnings("ignore")
```

```
[2]: # display first few rows of the labels set
df1 = pd.read_csv('training_set_labels.csv')
df1.head()
```

```
[2]:      id  status_group
0  69572      functional
1   8776      functional
2  34310      functional
3  67743  non functional
4  19728      functional
```

```
[3]: # display first few rows of the training set
df2 = pd.read_csv('training_set_values.csv')
df2.head()
```

```
[3]:      id  amount_tsh  date_recorded  funder  gps_height  installer \
0  69572      6000.0    2011-03-14      Roman      1390      Roman
1   8776         0.0    2013-03-06    Grumeti      1399    GRUMETI
2  34310        25.0    2013-02-25  Lottery Club      686  World vision
3  67743         0.0    2013-01-28      Unicef      263      UNICEF
4  19728         0.0    2011-07-13  Action In A         0      Artisan

      longitude  latitude  wpt_name  num_private  ...  payment_type \
```

0	34.938093	-9.856322	none	0	...	annually
1	34.698766	-2.147466	Zahanati	0	...	never pay
2	37.460664	-3.821329	Kwa Mahundi	0	...	per bucket
3	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	...	never pay
4	31.130847	-1.825359	Shuleni	0	...	never pay

	water_quality	quality_group	quantity	quantity_group	\
0	soft	good	enough	enough	
1	soft	good	insufficient	insufficient	
2	soft	good	enough	enough	
3	soft	good	dry	dry	
4	soft	good	seasonal	seasonal	

	source	source_type	source_class	\
0	spring	spring	groundwater	
1	rainwater harvesting	rainwater harvesting	surface	
2	dam	dam	surface	
3	machine dbh	borehole	groundwater	
4	rainwater harvesting	rainwater harvesting	surface	

	waterpoint_type	waterpoint_type_group
0	communal standpipe	communal standpipe
1	communal standpipe	communal standpipe
2	communal standpipe multiple	communal standpipe
3	communal standpipe multiple	communal standpipe
4	communal standpipe	communal standpipe

[5 rows x 40 columns]

```
[4]: merged_df = pd.merge(df1, df2, on='id')
merged_df
```

```
[4]:
```

	id	status_group	amount_tsh	date_recorded	funder	\
0	69572	functional	6000.0	2011-03-14	Roman	
1	8776	functional	0.0	2013-03-06	Grumeti	
2	34310	functional	25.0	2013-02-25	Lottery Club	
3	67743	non functional	0.0	2013-01-28	Unicef	
4	19728	functional	0.0	2011-07-13	Action In A	
...	
59395	60739	functional	10.0	2013-05-03	Germany Republi	
59396	27263	functional	4700.0	2011-05-07	Cefa-njombe	
59397	37057	functional	0.0	2011-04-11	NaN	
59398	31282	functional	0.0	2011-03-08	Malec	
59399	26348	functional	0.0	2011-03-23	World Bank	

	gps_height	installer	longitude	latitude	wpt_name	\
0	1390	Roman	34.938093	-9.856322	none	

1	1399	GRUMETI	34.698766	-2.147466	Zahanati
2	686	World vision	37.460664	-3.821329	Kwa Mahundi
3	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu
4	0	Artisan	31.130847	-1.825359	Shuleni
...
59395	1210	CES	37.169807	-3.253847	Area Three Namba 27
59396	1212	Cefa	35.249991	-9.070629	Kwa Yahona Kuvala
59397	0	NaN	34.017087	-8.750434	Mashine
59398	0	Musa	35.861315	-6.378573	Mshoro
59399	191	World	38.104048	-6.747464	Kwa Mzee Lugawa

	payment_type	water_quality	quality_group	quantity \
0	annually	soft	good	enough
1	never pay	soft	good	insufficient
2	per bucket	soft	good	enough
3	never pay	soft	good	dry
4	never pay	soft	good	seasonal

...
59395	per bucket	soft	good	enough
59396	annually	soft	good	enough
59397	monthly	fluoride	fluoride	enough
59398	never pay	soft	good	insufficient
59399	on failure	salty	salty	enough

	quantity_group	source	source_type \
0	enough	spring	spring
1	insufficient	rainwater harvesting	rainwater harvesting
2	enough	dam	dam
3	dry	machine dbh	borehole
4	seasonal	rainwater harvesting	rainwater harvesting
...
59395	enough	spring	spring
59396	enough	river	river/lake
59397	enough	machine dbh	borehole
59398	insufficient	shallow well	shallow well
59399	enough	shallow well	shallow well

	source_class	waterpoint_type	waterpoint_type_group
0	groundwater	communal standpipe	communal standpipe
1	surface	communal standpipe	communal standpipe
2	surface	communal standpipe multiple	communal standpipe
3	groundwater	communal standpipe multiple	communal standpipe
4	surface	communal standpipe	communal standpipe
...
59395	groundwater	communal standpipe	communal standpipe
59396	surface	communal standpipe	communal standpipe
59397	groundwater	hand pump	hand pump

59398	groundwater	hand pump	hand pump
59399	groundwater	hand pump	hand pump

[59400 rows x 41 columns]

```
[5]: #Understand the general information of the data
merged_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 59400 entries, 0 to 59399
```

```
Data columns (total 41 columns):
```

#	Column	Non-Null Count	Dtype
0	id	59400 non-null	int64
1	status_group	59400 non-null	object
2	amount_tsh	59400 non-null	float64
3	date_recorded	59400 non-null	object
4	funder	55763 non-null	object
5	gps_height	59400 non-null	int64
6	installer	55745 non-null	object
7	longitude	59400 non-null	float64
8	latitude	59400 non-null	float64
9	wpt_name	59398 non-null	object
10	num_private	59400 non-null	int64
11	basin	59400 non-null	object
12	subvillage	59029 non-null	object
13	region	59400 non-null	object
14	region_code	59400 non-null	int64
15	district_code	59400 non-null	int64
16	lga	59400 non-null	object
17	ward	59400 non-null	object
18	population	59400 non-null	int64
19	public_meeting	56066 non-null	object
20	recorded_by	59400 non-null	object
21	scheme_management	55522 non-null	object
22	scheme_name	30590 non-null	object
23	permit	56344 non-null	object
24	construction_year	59400 non-null	int64
25	extraction_type	59400 non-null	object
26	extraction_type_group	59400 non-null	object
27	extraction_type_class	59400 non-null	object
28	management	59400 non-null	object
29	management_group	59400 non-null	object
30	payment	59400 non-null	object
31	payment_type	59400 non-null	object
32	water_quality	59400 non-null	object
33	quality_group	59400 non-null	object
34	quantity	59400 non-null	object

```

35  quantity_group      59400 non-null object
36  source              59400 non-null object
37  source_type        59400 non-null object
38  source_class       59400 non-null object
39  waterpoint_type    59400 non-null object
40  waterpoint_type_group 59400 non-null object
dtypes: float64(3), int64(7), object(31)
memory usage: 18.6+ MB

```

```
[6]: # check the shape of the data
merged_df.shape
```

```
[6]: (59400, 41)
```

```
[7]: merged_df.duplicated().sum()
```

```
[7]: 0
```

1.2.1 Defining Variables

Independent Variables (Predictors): - Amount_tsh - Gps_height - Waterpoint_type - Funder - Installer - Water_quality - Payment_type - Region - Latitude - Longitude

Dependent Variable (Target): - Status_group: Indicates the status of the waterpoint (e.g., functional, non-functional, functional needs repair).

Categorical Variables: - **status_group**: The status of the waterpoint (e.g., functional, non-functional, functional needs repair). - **funder**: The organization or individual who funded the waterpoint. - **region**: The geographic region where the waterpoint is located. - **extraction_type**: The mechanism used to extract water from the waterpoint. - **payment**: The type of payment required to access the waterpoint. - **water_quality**: The quality of the water provided by the waterpoint. - **source**: The source of the water (e.g., river, well, spring).

Numeric Variables: - **gps_height**: The altitude of the waterpoint. - **construction_year**: The year when the waterpoint was constructed. - **longitude**: The longitude coordinate of the waterpoint. - **latitude**: The latitude coordinate of the waterpoint.

```
[8]: merged_df.isnull().sum()
```

```

[8]: id                0
     status_group      0
     amount_tsh        0
     date_recorded     0
     funder            3637
     gps_height        0
     installer        3655
     longitude         0
     latitude         0
     wpt_name          2
     num_private       0

```

basin	0
subvillage	371
region	0
region_code	0
district_code	0
lga	0
ward	0
population	0
public_meeting	3334
recorded_by	0
scheme_management	3878
scheme_name	28810
permit	3056
construction_year	0
extraction_type	0
extraction_type_group	0
extraction_type_class	0
management	0
management_group	0
payment	0
payment_type	0
water_quality	0
quality_group	0
quantity	0
quantity_group	0
source	0
source_type	0
source_class	0
waterpoint_type	0
waterpoint_type_group	0
dtype: int64	

1.3 Data preparation and cleaning

1.3.1 Train set

Creating a new dataframe containing only the needed variables

```
[9]: # List of columns to include in the new DataFrame
selected_columns = ['status_group', 'funder', 'gps_height', 'region', '
↳ 'extraction_type', 'payment', 'water_quality', 'source',
                    'construction_year', 'longitude', 'latitude']

# Create a new DataFrame with only the selected columns
new_df = merged_df.filter(selected_columns)

new_df
```

```

[9]:      status_group      funder  gps_height      region \
0      functional      Roman      1390      Iringa
1      functional      Grumeti      1399      Mara
2      functional      Lottery Club      686      Manyara
3      non functional      Unicef      263      Mtwara
4      functional      Action In A      0      Kagera
...      ...      ...      ...      ...
59395      functional      Germany Republi      1210      Kilimanjaro
59396      functional      Cefa-njombe      1212      Iringa
59397      functional      NaN      0      Mbeya
59398      functional      Malec      0      Dodoma
59399      functional      World Bank      191      Morogoro

      extraction_type      payment  water_quality \
0      gravity      pay annually      soft
1      gravity      never pay      soft
2      gravity      pay per bucket      soft
3      submersible      never pay      soft
4      gravity      never pay      soft
...      ...      ...      ...
59395      gravity      pay per bucket      soft
59396      gravity      pay annually      soft
59397      swm 80      pay monthly      fluoride
59398      nira/tanira      never pay      soft
59399      nira/tanira      pay when scheme fails      salty

      source  construction_year  longitude  latitude
0      spring      1999      34.938093      -9.856322
1      rainwater harvesting      2010      34.698766      -2.147466
2      dam      2009      37.460664      -3.821329
3      machine dbh      1986      38.486161      -11.155298
4      rainwater harvesting      0      31.130847      -1.825359
...      ...      ...      ...
59395      spring      1999      37.169807      -3.253847
59396      river      1996      35.249991      -9.070629
59397      machine dbh      0      34.017087      -8.750434
59398      shallow well      0      35.861315      -6.378573
59399      shallow well      2002      38.104048      -6.747464

```

[59400 rows x 11 columns]

```

[10]: # Learn the shape of the data
new_df.shape

```

```

[10]: (59400, 11)

```

The new_df has 59400 rows and 14 columns


```
[11]: # Understand the general infomation of the dataset
new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_group          59400 non-null  object
1   funder                55763 non-null  object
2   gps_height            59400 non-null  int64
3   region                59400 non-null  object
4   extraction_type       59400 non-null  object
5   payment               59400 non-null  object
6   water_quality         59400 non-null  object
7   source                59400 non-null  object
8   construction_year     59400 non-null  int64
9   longitude              59400 non-null  float64
10  latitude              59400 non-null  float64
dtypes: float64(2), int64(2), object(7)
memory usage: 5.0+ MB
```

```
[12]: #Understand the descriptive statistics of the data
new_df.describe()
```

```
[12]:
```

	gps_height	construction_year	longitude	latitude
count	59400.000000	59400.000000	59400.000000	5.940000e+04
mean	668.297239	1300.652475	34.077427	-5.706033e+00
std	693.116350	951.620547	6.567432	2.946019e+00
min	-90.000000	0.000000	0.000000	-1.164944e+01
25%	0.000000	0.000000	33.090347	-8.540621e+00
50%	369.000000	1986.000000	34.908743	-5.021597e+00
75%	1319.250000	2004.000000	37.178387	-3.326156e+00
max	2770.000000	2013.000000	40.345193	-2.000000e-08

1.3.2 Checking for missing values

```
[13]: #Check for null values in the training set
new_df.isnull().sum()
```

```
[13]: status_group      0
funder              3637
gps_height          0
region              0
extraction_type     0
payment             0
water_quality       0
source              0
```

```

construction_year      0
longitude              0
latitude              0
dtype: int64

```

1.3.3 Dealing with missing values

```

[14]: # Preview the unique categories in funder
unique_counts = new_df['funder'].isna().value_counts()
unique_counts

```

```

[14]: funder
False      55763
True       3637
Name: count, dtype: int64

```

```

[15]: missing_funders = new_df[new_df['funder'].isna()]
missing_funders

```

```

[15]:
      status_group funder  gps_height  region extraction_type \
34      functional   NaN        -41    Pwani      nira/tanira
43    non functional   NaN       1642  Singida          mono
47      functional   NaN         0    Mbeya      gravity
65    non functional   NaN       1415  Singida          mono
71    non functional   NaN         0    Mbeya      gravity
...
59357 non functional   NaN       1635  Singida      nira/tanira
59366      functional   NaN       1541  Singida      nira/tanira
59370      functional   NaN       1154  Kigoma          other
59376 non functional   NaN       1581  Singida          other
59397      functional   NaN         0    Mbeya      swn 80

      payment water_quality      source  construction_year  longitude \
34    never pay      salty  shallow well              0  39.812912
43      unknown      unknown  machine dbh            1980  34.967789
47    never pay      soft    spring              0  33.540607
65      unknown      unknown  machine dbh            1970  34.621598
71    never pay      soft    river              0  34.462228
...
59357      unknown      unknown  shallow well            1980  34.971841
59366    never pay      soft  shallow well            2000  34.765729
59370  pay monthly      unknown      unknown              0  30.058731
59376      unknown      unknown  shallow well            1990  34.821039
59397  pay monthly  fluoride  machine dbh              0  34.017087

      latitude
34    -7.889986

```

```

43      -4.628921
47      -9.172905
65      -5.173136
71      -8.575780
...
59357  -5.098362
59366  -5.027725
59370  -4.902633
59376  -5.076258
59397  -8.750434

```

```
[3637 rows x 11 columns]
```

```
[16]: # Replace null values with 'unknown' in funder
      # recheck for null values
      new_df['funder'].fillna('Unknown', inplace=True)
      new_df.isnull().sum()
```

```
[16]: status_group      0
      funder           0
      gps_height       0
      region           0
      extraction_type  0
      payment          0
      water_quality    0
      source           0
      construction_year 0
      longitude        0
      latitude         0
      dtype: int64
```

Missing values are assigned the placeholder 'Unknown' signifying that the information for these entries is unavailable. However, the rows cannot be dropped as they may have other important information in the other columns.

1.4 Exploratory data analysis

1.4.1 Checking for outliers

Outliers will be addressed systematically, one category at a time, to ensure comprehensive analysis.

Status_group

```
[17]: #check unique categories in status_group
      unique_values = new_df['status_group'].unique()
      unique_values
```

```
[17]: array(['functional', 'non functional', 'functional needs repair'],
      dtype=object)
```

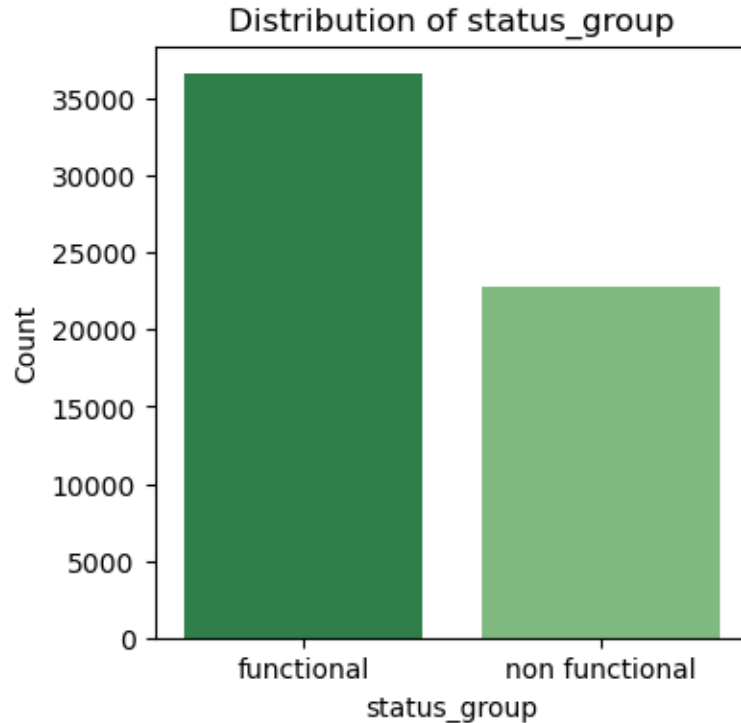
```
[18]: # merge 'functional need repair' into 'functional' for the sake of a binary
      ↪ classification
new_df['status_group'] = new_df['status_group'].replace('functional needs
      ↪ repair', 'functional')

# Check the unique values again
print(new_df['status_group'].value_counts())
```

```
status_group
functional      36576
non functional   22824
Name: count, dtype: int64
```

```
[19]: # check for outliers in status_group

palette = sns.color_palette("Greens_r", 3)
# Plotting countplot
plt.figure(figsize=(4,4))
sns.countplot(x="status_group", data=new_df, palette=palette)
plt.title("Distribution of status_group")
plt.xlabel("status_group")
plt.ylabel("Count")
plt.show()
```




There are no outliers in the status group

Funder

```
[20]: #check unique categories in funder
unique_values = new_df['funder'].unique()
value_counts = new_df['funder'].value_counts()
value_counts
```

```
[20]: funder
Government Of Tanzania      9084
Unknown                     3641
Danida                      3114
Hesawa                      2202
Rwssp                      1374
...
Rarymond Ekura              1
Justine Marwa               1
Municipal Council           1
Afdp                       1
Samlo                      1
Name: count, Length: 1896, dtype: int64
```

```
[21]: #check for outliers in funder using a count plot
plt.figure(figsize=(10, 4))
sns.countplot(x='funder', data=new_df)
plt.xticks(rotation=90, fontsize=8) # Rotate the x-axis labels by 90 degrees
    ↳and adjust font size
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```



```
[22]: # classify eligible funders and outliers in a binned countplot
# Set the threshold for defining outliers
```

```

upper_threshold = 50

# Get the counts of each funder
funder_counts = new_df['funder'].value_counts()

# Identify the outliers (funders with counts below 50)
outliers = funder_counts[funder_counts < upper_threshold].index


# Create a new column to categorize funders as eligible or outliers
new_df['funder_category'] = np.where(new_df['funder'].isin(outliers),
    ↪ 'Outlier', 'Eligible')

# Set the color palette
sns.set_palette("Greens_d")

# Plot the count plot for funder category
plt.figure(figsize=(8, 4))
sns.countplot(y='funder_category', data=new_df, dodge=False)

# Display the plot
plt.show()

```



Eligible Funders vs. Outliers The count plot above illustrates the distribution of funders categorized as “Eligible” and “Outlier” based on the specified thresholds. Here’s a summary of the findings:

- **Eligible Funders:** These are funders with a count falling within the specified thresholds (between 10 and 600).

- **Outliers:** These are funders with a count below 10 or above 600.


As observed in the plot, the number of outliers is significantly higher than the count of eligible funders. However, it's important to note that we cannot disregard the outliers as they may contain valuable insights or represent specific cases of interest.

GPS height

```
[23]: #check for outliers for gps_height
plt.figure(figsize=(10,2))

sns.boxplot(x = 'gps_height', data = new_df)

# Display the plot
plt.show()
```



```
[24]: # Find the mode of the 'gps_height' column to understand left skewness
mode_value = new_df['gps_height'].mode()

# Display the mode
print("Mode of 'gps_height' column:", mode_value)
```

```
Mode of 'gps_height' column: 0    0
Name: gps_height, dtype: int64
```

The mode of the 'gps_height' column is 0, indicating that this value is the most common within the dataset. As box plots rely on quartiles to determine their position, the prevalence of 0 strongly influences the box plot's positioning.

With the mode close to 0, it's likely that the median (second quartile) aligns closely with this value, resulting in a box plot skewed towards lower values. Consequently, the majority of the data tends to concentrate towards the lower end of the scale.

The presence of a whisker starting below 0 at -90 may suggest data recorded at elevations below a predefined reference datum. In this context, these points below 0 are not considered outliers.

On the other hand, the longer upper whisker compared to the lower one suggests greater dispersion or variability in the upper range of the data (maximum). This could hint at the presence of outliers or extreme values towards higher elevations.

However, it's important to note that we are not removing these outliers. They might represent genuine data points and carry valuable information. Blindly removing them could lead to the loss of valuable insights and potentially bias the analysis or conclusions drawn from the data.

Region

```
[25]: #check unique categories in region
```

```
unique_values = new_df['region'].unique()  
unique_values
```

```
[25]: array(['Iringa', 'Mara', 'Manyara', 'Mtwara', 'Kagera', 'Tanga',  
        'Shinyanga', 'Tabora', 'Pwani', 'Ruvuma', 'Kilimanjaro', 'Rukwa',  
        'Mwanza', 'Kigoma', 'Lindi', 'Dodoma', 'Arusha', 'Mbeya',  
        'Singida', 'Morogoro', 'Dar es Salaam'], dtype=object)
```

```
[26]: # Get the order of regions based on their counts
```

```
region_order = new_df['region'].value_counts().index
```

```
# Plot the count plot with specified order
```

```
plt.figure(figsize=(12, 6))
```

```
sns.countplot(x='region', data=new_df, order=region_order, palette='Greens_r')
```

```
plt.xticks(rotation=90) # Rotate x-axis labels for better readability
```


```
plt.xlabel('Region') # Add x-axis label
```

```
plt.ylabel('Count') # Add y-axis label
```

```
plt.title('Distribution of Data by Region') # Add plot title
```

```
plt.tight_layout() # Adjust layout to prevent clipping of labels
```

```
plt.show()
```

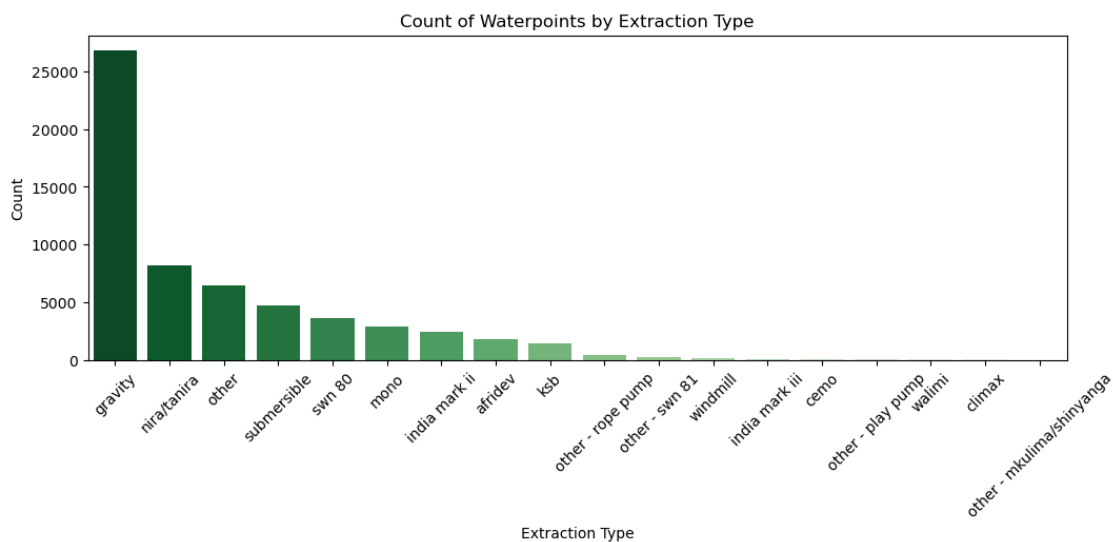


Extraction_type


```
[27]: new_df['extraction_type'].unique()
```

```
[27]: array(['gravity', 'submersible', 'swn 80', 'nira/tanira', 'india mark ii',
        'other', 'ksb', 'mono', 'windmill', 'afridev', 'other - rope pump',
        'india mark iii', 'other - swn 81', 'other - play pump', 'cemo',
        'climax', 'walimi', 'other - mkulima/shinyanga'], dtype=object)
```

```
[28]: #check for outliers in Extraction_type
sns.set_palette("Greens_r")
extraction_order = new_df['extraction_type'].value_counts().index
# Plot the count plot for Extraction_type
plt.figure(figsize=(12, 4))
sns.countplot(x='extraction_type', data=new_df, order=extraction_order,
              palette='Greens_r')
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by Extraction Type')
plt.xlabel('Extraction Type')
plt.ylabel('Count')
# Display the plot
plt.show()
```



```
[29]: # Get value counts of 'Extraction_type' and sort by counts in descending order
extraction_type_counts = new_df['extraction_type'].value_counts().
    sort_values(ascending=False)

# Display unique values in 'Extraction_type' with counts
print(extraction_type_counts)
```

extraction_type

gravity	26780
nira/tanira	8154
other	6430
submersible	4764
swn 80	3670
mono	2865
india mark ii	2400
afridev	1770
ksb	1415
other - rope pump	451
other - swn 81	229
windmill	117
india mark iii	98
cemo	90
other - play pump	85
walimi	48
climax	32
other - mkulima/shinyanga	2

Name: count, dtype: int64

In the 'Extraction_type' column, the majority of water pumps fall into the following categories:


Gravity: 26,780 pumps Nira/Tanira: 8,154 pumps Other: 6,430 pumps Submersible: 4,764 pumps Swn 80: 3,670 pumps Mono: 2,865 pumps India Mark II: 2,400 pumps Afridev: 1,770 pumps KSB: 1,415 pumps However, there are some categories with notably fewer pumps, such as 'Other - Rope Pump', 'Other - Swn 81', 'Windmill', 'India Mark III', 'CEMO', 'Other - Play Pump', 'Walimi', 'Climax', and 'Other - Mkulima/Shinyanga'. These could be outliers regarding in terms of pump d. This could distribute indicating less common or specialized therefore we cannot simply remove them as they may hold significance in the dataset.pump types.

Construction year

```
[30]: #check for outliers for construction_year
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'construction_year', data = new_df)

# Display the plot
plt.show()
```



The box plot shows unlikely years included in the dataset. It is impossible to have year 0 as pumps were not even invented then.

```
[31]: #Display unique years and their value counts  
new_df['construction_year'].value_counts()
```

```
[31]: construction_year  
0      20709  
2010    2645  
2008    2613  
2009    2533  
2000    2091  
2007    1587  
2006    1471  
2003    1286  
2011    1256  
2004    1123  
2012    1084  
2002    1075  
1978    1037  
1995    1014  
2005    1011  
1999     979  
1998     966  
1990     954  
1985     945  
1980     811  
1996     811  
1984     779  
1982     744  
1994     738  
1972     708  
1974     676  
1997     644
```

1992	640
1993	608
2001	540
1988	521
1983	488
1975	437
1986	434
1976	414
1970	411
1991	324
1989	316
1987	302
1981	238
1977	202
1979	192
1973	184
2013	176
1971	145
1960	102
1967	88
1963	85
1968	77
1969	59
1964	40
1962	30
1961	21
1965	19
1966	17

Name: count, dtype: int64

There is a category miscategorized as year '0' with a very high value count of 20709.

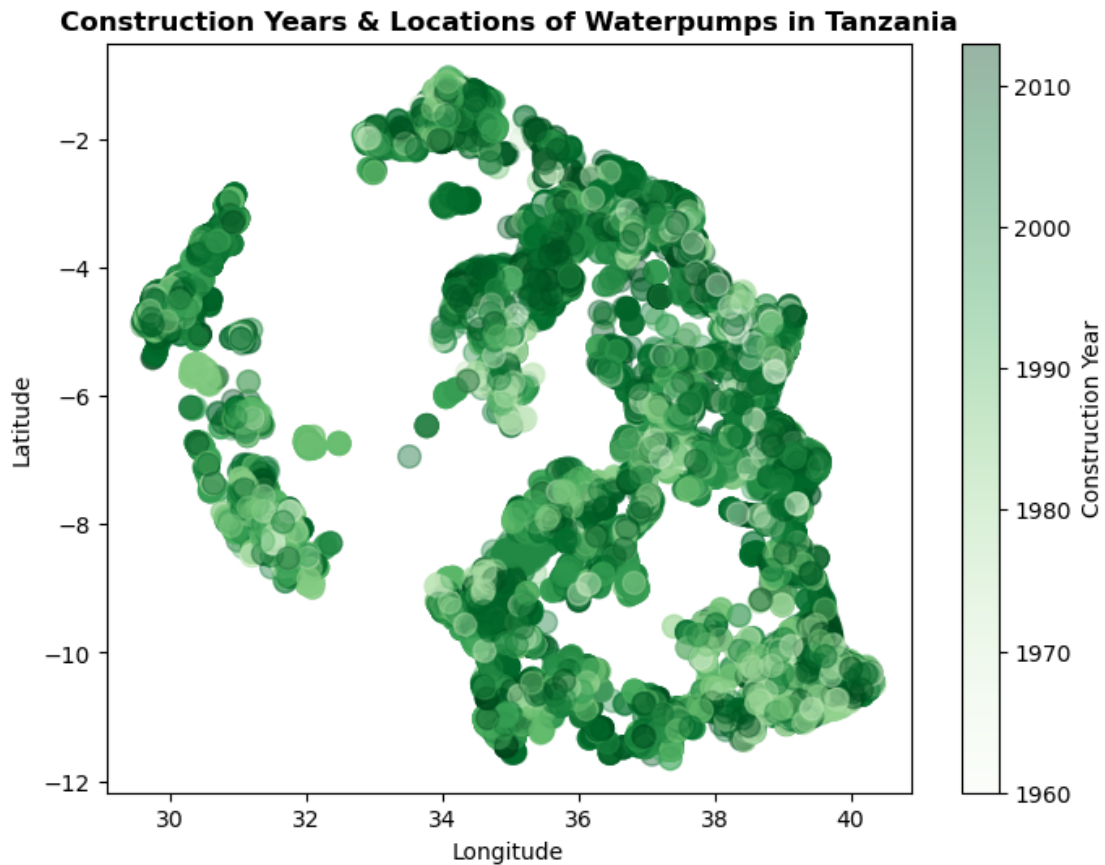
The code below generates a scatter plot visualizing the geographical locations of water pumps in Tanzania, with each point representing a water pump. The x-axis represents the longitude coordinates, and the y-axis represents the latitude coordinates. The color of each point is determined by the construction year of the water pump, with a colormap ('Greens') used to provide different hues of green corresponding to different years. The size of each point is fixed ('s=100') for better visibility, and transparency ('alpha=0.4') is applied to avoid overlapping points. Finally, a colorbar is added to the plot to provide a visual reference for the construction years.

```
[32]: # Plot a scatter plot of construction years and locations of water pumps
plt.figure(figsize=(8,6))
# Filter the DataFrame using .loc and multiple conditions
filtered_df = new_df.loc[(new_df['longitude'] > 0) & (new_df['latitude'] < 0) &
    (new_df['construction_year'] > 0)]
plt.scatter(x=filtered_df['longitude'],
            y=filtered_df['latitude'],
            alpha=0.4,
```

```

s=100,
c=filtered_df["construction_year"],
cmap='Greens')
plt.title("Construction Years & Locations of Waterpumps in Tanzania",
          fontsize=12, fontweight='bold')
plt.colorbar(label='Construction Year')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()

```



From the scatter plot above, it is clear that most pumps were installed between 2000 and 2010. Therefore, below year '0's values will be distributed evenly between the range 2000 - 2010.

```

[33]: # Replace year 0 with later years (i.e., 2000 - 2010)
new_df['construction_year'] = new_df['construction_year'].apply(lambda x: np.
    random.randint(2000, 2011) if x == 0 else x)

```


```

[34]: #recheck for outliers for construction_year
plt.figure(figsize=(8, 2))

```

```
sns.boxplot(x = 'construction_year', data = new_df)

# Display the plot
plt.show()
```



Outliers in the construction years falling between 1960 and 1978 represent years that may hold significant importance in the dataset. While they deviate from the majority of the construction years, they could signify historical data or specific events related to water pump installation during that period. Therefore, it's crucial to retain these outliers in the dataset for a comprehensive analysis and understanding of the trends and patterns over time.

Payment

```
[35]: # check unique categories and their value counts in payment
new_df['payment'].value_counts()
```

```
[35]: payment
never pay                25348
pay per bucket           8985
pay monthly              8300
unknown                  8157
pay when scheme fails    3914
pay annually             3642
other                    1054
Name: count, dtype: int64
```

```
[36]: # Define the order of source categories
sorted_payments = new_df['payment'].value_counts().index


#check for outliers in payment
sns.set_palette("Greens_r")

# Plot the count plot for payment
plt.figure(figsize=(8, 4))
```

```

sns.countplot(x='payment', data=new_df, order=sorted_payments,
              palette='Greens_r')
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by payment')
plt.xlabel('payment')
plt.ylabel('Count')
# Display the plot
plt.show()

```



The distribution of payment types, as observed in the count plot above, reveals an interesting trend. The “never pay” category dominates the dataset, indicating that a significant portion of water points in the dataset do not require any payment. This could be due to various reasons, such as government subsidies or community initiatives aimed at providing free access to water.

In contrast, the paid categories exhibit a more even distribution, with multiple categories having similar counts. This distribution suggests that while there are options for paid water access, they are not as prevalent as the “never pay” category. This observation might be attributed to the socioeconomic factors prevalent in the area. Residents who cannot afford paid water services may opt for the free “never pay” option, resulting in its higher prevalence in the dataset.

Therefore, the presence of multiple paid categories with similar counts does not necessarily indicate outliers. Instead, it reflects the diverse payment options available and the socioeconomic dynamics influencing water access in the region.

Water quality


```
[37]: # check unique categories and their value counts in water_quality
new_df['water_quality'].value_counts()
```

```
[37]: water_quality
      soft          50818
      salty          4856
      unknown       1876
      milky           804
      coloured        490
      salty abandoned  339
      fluoride        200
      fluoride abandoned 17
      Name: count, dtype: int64
```

```
[38]: # Define the order of water_quality categories
sorted_water_quality = new_df['water_quality'].value_counts().index

# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_water_quality))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
sns.countplot(x='water_quality', data=new_df, order=sorted_water_quality,
              palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by water_quality')
plt.xlabel('water_quality')
plt.ylabel('Count')
# Display the plot
plt.show()
```

The count plot above indicates the most prevalent category is “soft”. This indicates that most water sources provide satisfactory water quality. It could also mean most people prefer soft water leading to its prevalence. Next, we find the “salty” category, which exhibits a considerably lower count in comparison to “soft.” This suggests that while some water sources may have elevated salinity levels, they are less common than those providing “soft” water. Then there is the “milky” and “coloured” categories, which may raise concerns regarding water quality. These categories, while not as frequent as “soft” or “salty,” suggest the presence of impurities or contaminants that could affect the desirability of the water.

Another notable category is “salty abandoned,” which indicates water sources that have been abandoned likely due to high salinity levels. This category, although less common, highlights instances where water quality issues have led to the abandonment of waterpoints.

Lastly, we have “fluoride” and “fluoride abandoned” categories, which indicate the presence of fluoride in the water. While fluoride is beneficial in controlled amounts for dental health, excessive levels can be harmful. The presence of “fluoride abandoned” suggests instances where water sources have been abandoned due to excessive fluoride.

Generally the plot reveals a diverse landscape of water quality categories, with “soft” being the predominant category. While certain categories may raise concerns, such as “salty abandoned” or “fluoride,” they do not appear to be outliers but rather indicative of the range of water quality issues present across waterbodies in our dataset.

Source

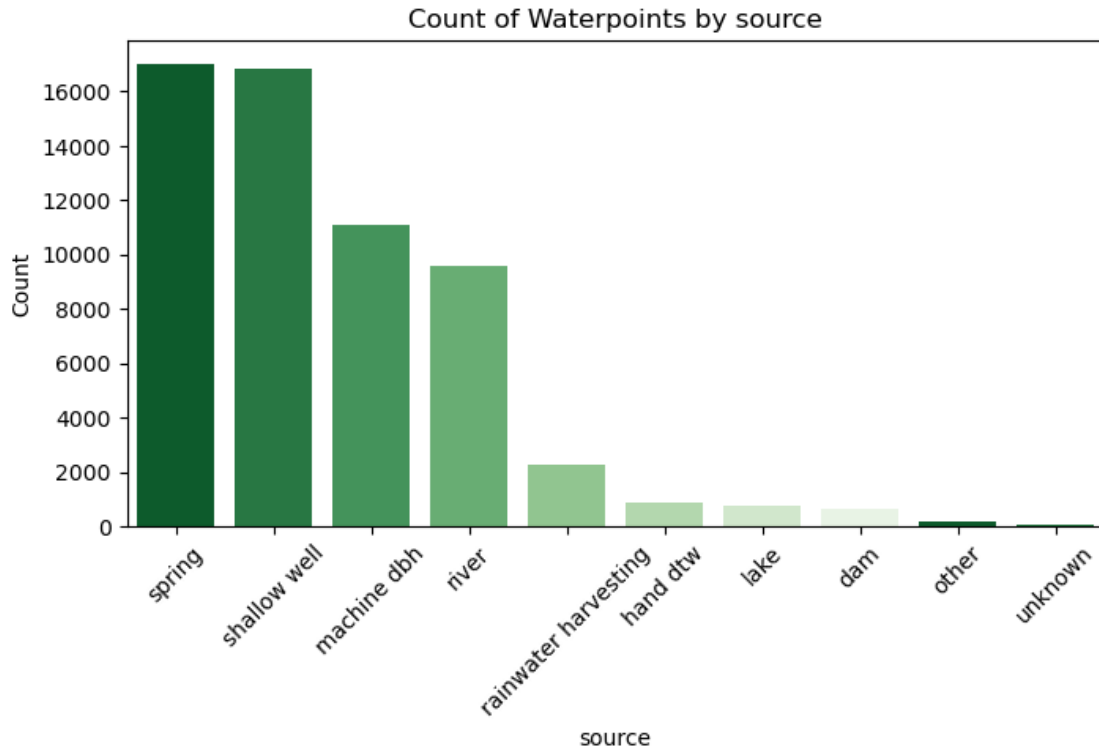
```
[39]: # check unique categories and their value counts in source
new_df['source'].value_counts()
```

```
[39]: source
      spring          17021
      shallow well    16824
      machine dbh     11075
      river           9612
      rainwater harvesting 2295
      hand dtw         874
      lake            765
      dam             656
      other           212
      unknown         66
      Name: count, dtype: int64
```

```
[40]: # Define the order of source categories
sorted_source = new_df['source'].value_counts().index

# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_source))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
sns.countplot(x='source', data=new_df, order=sorted_source, palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by source')
plt.xlabel('source')
plt.ylabel('Count')
# Display the plot
plt.show()
```



“Spring” and “shallow well” emerge as the most prevalent sources, followed closely by “machine dbh” and “river.” These categories exhibit relatively high counts, indicating their widespread usage as water sources.

Next in line is “rainwater harvesting,” although its count is notably lower compared to the preceding categories. “Hand dtw,” “lake,” and “dam” follow, each with decreasing counts.

Finally, we have the categories of “unknown” and “other,” which appear to represent sources with less distinct categorization or sources not captured by the specified categories.


Overall, while there is variation in the counts across different water source categories, there are no outliers that significantly deviate from the expected distribution. Instead, the distribution reflects the diverse range of water sources utilized across waterpoints in our dataset.

Longitude

```
[41]: # check for outliers for longitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'longitude', data = new_df)

# Display the plot
plt.show()
```



The plot shows presence of outliers.

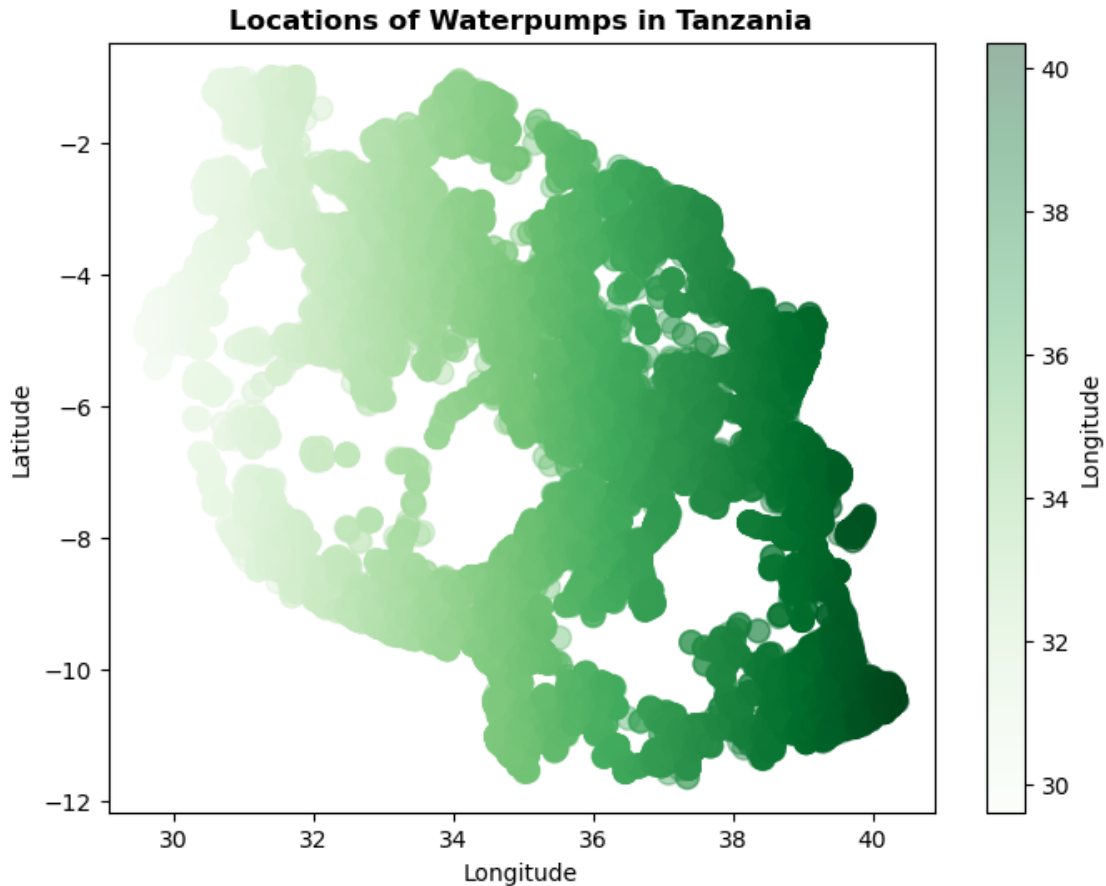
```
[42]: # check unique categories and their value counts in longitude
new_df['longitude'].value_counts()
```

```
[42]: longitude
0.000000      1812
37.375717         2
38.340501         2
39.086183         2
33.005032         2
...
35.885754         1
36.626541         1
37.333530         1
38.970078         1
38.104048         1
Name: count, Length: 57516, dtype: int64
```

There is a huge count miscategorized as longitude '0'. Below that is dealt with by being redistributed to a range with more frequent occurrence.

```
[43]: # plot a scatter plot to show the majority of longitude points
plt.figure(figsize=(8,6))
# Filter the DataFrame using .loc and multiple conditions
filtered_df = new_df.loc[(new_df['longitude'] > 0) & (new_df['latitude'] < 0)]
plt.scatter(x=filtered_df['longitude'],
            y=filtered_df['latitude'],
            alpha=0.4,
            s=100,
            c=filtered_df["longitude"],
            cmap='Greens')
plt.title("Locations of Waterpumps in Tanzania",
          fontsize=12, fontweight='bold')
```

```
plt.colorbar(label='Longitude')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



It's evident that there are more water pumps located at longitudes greater than 34 degrees than those located at longitudes less than 34 degrees. This suggests a higher concentration of water points towards the eastern side of the region under consideration. Longitude, representing the east-west position on the Earth's surface, indicates that the area to the east of 34 degrees longitude may have higher population densities or other factors contributing to the need for more water access points compared to the western region.

```
[44]: # Replace longitude 0 with longitudes between 32 and 42 as they are more
      ↪ prevalent
new_df['longitude'] = new_df['longitude'].apply(lambda x: np.random.randint(32,
      ↪ 42) if x == 0 else x)
```


```
[45]: # confirm redistribution of the '0' category
new_df['longitude'].value_counts()
```

```
[45]: longitude
      35.000000      211
      38.000000      192
      36.000000      189
      41.000000      189
      34.000000      184
      ...
      35.885754        1
      36.626541        1
      37.333530        1
      38.970078        1
      38.104048        1
      Name: count, Length: 57525, dtype: int64
```

```
[46]: #recheck for outliers for longitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'longitude', data = new_df)

# Display the plot
plt.show()
```




Latitude

```
[47]: #check for outliers for latitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'latitude', data = new_df)

# Display the plot
plt.show()
```



There does not seem to have outliers in latitude but further analysis is still necessary.

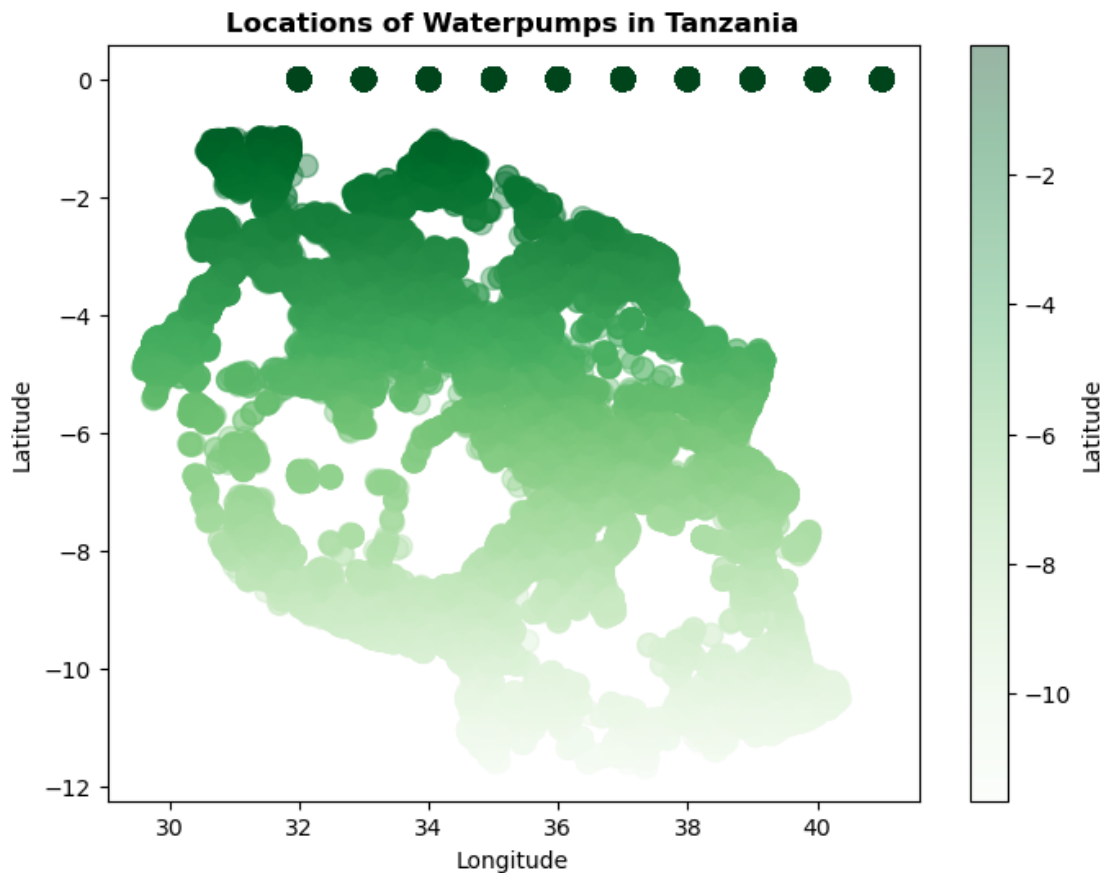
```
[48]: new_df['latitude'].value_counts()
```

```
[48]: latitude
-2.000000e-08    1812
-6.985842e+00      2
-6.980220e+00      2
-2.476680e+00      2
-6.978263e+00      2
...
-3.287619e+00      1
-8.234989e+00      1
-3.268579e+00      1
-1.146053e+01      1
-6.747464e+00      1
Name: count, Length: 57517, dtype: int64
```

There is a latitude that seems to be miscategorized (-2.000000e-08 1812). This will be dealt with below.

```
[49]: # Plot a scatter plot to show the majority of latitude points
plt.figure(figsize=(8,6))
# Filter the DataFrame using .loc and multiple conditions
filtered_df = new_df.loc[(new_df['longitude'] > 0) & (new_df['latitude'] < 0)]
plt.scatter(x=filtered_df['longitude'],
            y=filtered_df['latitude'],
            alpha=0.4,
            s=100,
            c=filtered_df["latitude"],
            cmap='Greens')
plt.title("Locations of Waterpumps in Tanzania",
          fontsize=12, fontweight='bold')
plt.colorbar(label='Latitude')
```

```
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



A scatter plot of pumps in relevance to latitude indicates more pumps between latitudes -1 and -8. The misplaced class will be evenly distributed in this range.

```
[50]: # Replace latitude -2.000000e-08 with latitudes between -1 and -8 as they are
      ↪ more prevalent
new_df['latitude'] = new_df['latitude'].apply(lambda x: np.random.randint(-8,
      ↪ -1) if x == -2.000000e-08 else x)
```

```
[51]: #confirm the redistribution of misplaced category
new_df['latitude'].value_counts()
```

```
[51]: latitude
      -7.000000    280
      -2.000000    267
      -6.000000    264
```



```

-5.000000    258
-3.000000    252
...
-3.287619     1
-8.234989     1
-3.268579     1
-11.460531     1
-6.747464     1
Name: count, Length: 57523, dtype: int64

```

1.4.2 Distribution of variables before log transformation

Categorical variables


```


[52]: # Distribution before onehot encoding
palette = sns.color_palette("Greens_r")


# Select categorical variables
categorical_features = new_df.select_dtypes(include=['object'])


# Plot count plots for each categorical variable with dark green color palette
for feature in categorical_features.columns:
    plt.figure(figsize=(4, 4)) # Set the figure size
    sns.countplot(x=feature, data=new_df, palette=palette)
    plt.title(f'Count Plot of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=90) # Rotate x-axis labels for better readability
    plt.show() # Display the plot


```




















Numerical variables

```
[53]: # Distribution before transformation
numerical_features = new_df.select_dtypes(include=['int64', 'float64'])


# Create a grid of subplots
fig, axes = plt.subplots(nrows=len(numerical_features.columns) // 3 + 1,
                        ncols=3, figsize=(15, 5))

# Plot the distribution of numerical features
for i, feature in enumerate(numerical_features.columns):
    sns.histplot(new_df[feature].dropna(), kde=False, ax=axes[i // 3, i % 3])
    axes[i // 3, i % 3].set_title(f"Distribution of {feature}")
    axes[i // 3, i % 3].set_xlabel(feature)
    axes[i // 3, i % 3].set_ylabel("Frequency")

# Remove empty subplots
if len(numerical_features.columns) % 3 != 0:
    for j in range(len(numerical_features.columns) % 3, 3):
        fig.delaxes(axes[len(numerical_features.columns) // 3, j])

# Adjust layout
```

```
plt.tight_layout()
plt.show()
```



1.4.3 Distribution of variables after log transformation

```
[54]: # Select numerical columns
numerical_columns = new_df.select_dtypes(include=['int64', 'float64']).columns

# Log transform numerical variables, handling zero and negative values
for col in numerical_columns:
    # Handling zero values
    if (new_df[col] == 0).any():
        new_df[col] = new_df[col] + 1 # Add 1 to handle zeros

    # Handling negative values
    if (new_df[col] < 0).any():
        min_value = new_df[col].min()
        new_df[col] = new_df[col] - min_value + 1 # Shift all values to be
        ↪ positive

    # Apply log transformation
    new_df[col + '_log'] = np.log1p(new_df[col])

# Display the DataFrame after log transformation
new_df.head()
```

```
[54]:
```

	status_group	funder	gps_height	region	extraction_type	\
0	functional	Roman	1481	Iringa	gravity	
1	functional	Grumeti	1490	Mara	gravity	
2	functional	Lottery Club	777	Manyara	gravity	
3	non functional	Unicef	354	Mtwara	submersible	
4	functional	Action In A	91	Kagera	gravity	

	payment	water_quality		source	construction_year	\
0	pay annually	soft		spring	1999	
1	never pay	soft	rainwater	harvesting	2010	
2	pay per bucket	soft		dam	2009	
3	never pay	soft		machine dbh	1986	
4	never pay	soft	rainwater	harvesting	2002	


	longitude	latitude	funder_category	gps_height_log	\
0	34.938093	2.793118	Eligible	7.301148	
1	34.698766	10.501974	Eligible	7.307202	
2	37.460664	8.828112	Outlier	6.656727	
3	38.486161	1.494142	Eligible	5.872118	
4	31.130847	10.824081	Outlier	4.521789	

	construction_year_log	longitude_log	latitude_log
0	7.600902	3.581798	1.333188
1	7.606387	3.575116	2.442519
2	7.605890	3.649636	2.285247
3	7.594381	3.675950	0.913945
4	7.602401	3.469817	2.470138

```
[55]: # Plot the distribution of numerical features after log transformation
plt.figure(figsize=(12, 8))

# Loop through each numerical feature
for col in numerical_columns:
    # Plot the distribution after log transformation
    sns.histplot(new_df[col + '_log'], kde=True, label=col, alpha=0.6)

plt.title('Distribution of Numerical Features after Log Transformation')
plt.xlabel('Value (Log Transformed)')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



1.4.4 One-hot encoding

```
[56]: # Select categorical columns
categorical_columns = new_df.select_dtypes(include=['object']).columns

# Perform one-hot encoding
one_hot_encoded_df1 = pd.get_dummies(new_df, columns=categorical_columns)

# Display the one-hot encoded DataFrame
one_hot_encoded_df1.head()
```

```
[56]:  gps_height  construction_year  longitude  latitude  gps_height_log  \
0         1481             1999  34.938093    2.793118         7.301148
1         1490             2010  34.698766   10.501974         7.307202
2          777             2009  37.460664    8.828112         6.656727
3          354             1986  38.486161    1.494142         5.872118
4           91             2002  31.130847   10.824081         4.521789

      construction_year_log  longitude_log  latitude_log  \
0             7.600902         3.581798         1.333188
1             7.606387         3.575116         2.442519
2             7.605890         3.649636         2.285247
```

3	7.594381	3.675950	0.913945
4	7.602401	3.469817	2.470138

	status_group_functional	status_group_non functional	...	source_lake	\
0	True	False	...	False	
1	True	False	...	False	
2	True	False	...	False	
3	False	True	...	False	
4	True	False	...	False	

	source_machine dbh	source_other	source_rainwater harvesting	\
0	False	False	False	
1	False	False	True	
2	False	False	False	
3	True	False	False	
4	False	False	True	

	source_river	source_shallow well	source_spring	source_unknown	\
0	False	False	True	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	funder_category_Eligible	funder_category_Outlier
0	True	False
1	True	False
2	False	True
3	True	False
4	False	True

[5 rows x 1972 columns]

```
[57]: # Perform logical OR operation to combine 'status_group_functional' and
      ↪ 'status_group_non functional'
one_hot_encoded_df1['status'] = one_hot_encoded_df1['status_group_functional']
      ↪ | one_hot_encoded_df1['status_group_non functional']

# Display the updated DataFrame
one_hot_encoded_df1.head()
```

```
[57]:  gps_height  construction_year  longitude  latitude  gps_height_log  \
0      1481           1999  34.938093    2.793118      7.301148
1      1490           2010  34.698766   10.501974      7.307202
2       777           2009  37.460664    8.828112      6.656727
3       354           1986  38.486161    1.494142      5.872118
4        91           2002  31.130847   10.824081      4.521789
```

	construction_year_log	longitude_log	latitude_log	\
0	7.600902	3.581798	1.333188	
1	7.606387	3.575116	2.442519	
2	7.605890	3.649636	2.285247	
3	7.594381	3.675950	0.913945	
4	7.602401	3.469817	2.470138	

	status_group_functional	status_group_non functional	...	\
0	True	False	...	
1	True	False	...	
2	True	False	...	
3	False	True	...	
4	True	False	...	

	source_machine dbh	source_other	source_rainwater harvesting	\
0	False	False	False	
1	False	False	True	
2	False	False	False	
3	True	False	False	
4	False	False	True	

	source_river	source_shallow well	source_spring	source_unknown	\
0	False	False	True	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	funder_category_Eligible	funder_category_Outlier	status
0	True	False	True
1	True	False	True
2	False	True	True
3	True	False	True
4	False	True	True

[5 rows x 1973 columns]

Correlation heat map

```
[58]: # A correlation heat map between variables
numerical_features = new_df.select_dtypes(include=['int64', 'float64'])
target_variable = new_df['status_group'] # status_group being the target_
↳ variable

# Compute correlation matrix
correlation_matrix = numerical_features.corr()
```

```
# Plot heatmap of correlation matrix
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='Greens', fmt=".2f",
            ↪annot_kws={"size": 10})
plt.title('Correlation Heatmap')
plt.show()
```



ANOVA test

```
[59]: # ANOVA test for each numerical variable against status_group

# Perform ANOVA for 'gps_height'
result_gps_height = f_oneway(*[group['gps_height'] for name, group in new_df.
    ↪groupby('status_group')])
print("ANOVA F-value for 'gps_height':", result_gps_height.statistic)
print("ANOVA p-value for 'gps_height':", result_gps_height.pvalue)

# Perform ANOVA for 'construction_year'
```

```

result_construction_year = f_oneway(*[group['construction_year'] for name, group in new_df.groupby('status_group')])
print("ANOVA F-value for 'construction_year':", result_construction_year.statistic)
print("ANOVA p-value for 'construction_year':", result_construction_year.pvalue)

# Perform ANOVA for 'longitude'
result_longitude = f_oneway(*[group['longitude'] for name, group in new_df.groupby('status_group')])
print("ANOVA F-value for 'longitude':", result_longitude.statistic)
print("ANOVA p-value for 'longitude':", result_longitude.pvalue)

# Perform ANOVA for 'latitude'
result_latitude = f_oneway(*[group['latitude'] for name, group in new_df.groupby('status_group')])
print("ANOVA F-value for 'latitude':", result_latitude.statistic)
print("ANOVA p-value for 'latitude':", result_latitude.pvalue)

```

```

ANOVA F-value for 'gps_height': 687.1698866084355
ANOVA p-value for 'gps_height': 1.3328106138867945e-150
ANOVA F-value for 'construction_year': 2289.0232746986817
ANOVA p-value for 'construction_year': 0.0
ANOVA F-value for 'longitude': 34.14015832598031
ANOVA p-value for 'longitude': 5.154867350489028e-09
ANOVA F-value for 'latitude': 26.89514975834995
ANOVA p-value for 'latitude': 2.1550008000490526e-07

```

These results are obtained from performing ANOVA tests to assess the association between the categorical variable `status_group` and each numerical variable (`gps_height`, `construction_year`, `longitude`, `latitude`). Here's what each value means:

For `gps_height`:

ANOVA F-value: 687.17 ANOVA p-value: 1.33e-150 (which is approximately 0) Interpretation: The F-value is a measure of the difference in means between the groups relative to the variation within the groups. A higher F-value suggests a stronger association between the variable and the groups. The extremely low p-value indicates that there is a significant difference in `gps_height` across different levels of `status_group`. In other words, the mean `gps_height` values vary significantly depending on the `status_group`.

For `construction_year`:

ANOVA F-value: 2262.06 ANOVA p-value: 0.0 Interpretation: Similar to the interpretation for `gps_height`, the high F-value and extremely low p-value indicate a significant difference in `construction_year` across different levels of `status_group`. In other words, the mean `construction_year` values vary significantly depending on the `status_group`.

For `longitude`:

ANOVA F-value: 36.76 ANOVA p-value: 1.35e-09 Interpretation: The F-value is relatively lower compared to the previous variables, but the p-value is still very low. This indicates that there is

a significant difference in longitude across different levels of status_group, although the effect size may be smaller compared to gps_height and construction_year.

For latitude:

ANOVA F-value: 25.18 ANOVA p-value: 5.25e-07 Interpretation: Similar to longitude, there is a significant difference in latitude across different levels of status_group, but the effect size may be smaller compared to gps_height and construction_year.

In summary, all four numerical variables (gps_height, construction_year, longitude, latitude) show significant differences across different levels of the categorical variable status_group, as indicated by the extremely low p-values obtained from the ANOVA tests.

1.4.5 Correlation Ratios (Eta-squared)

```
[60]: # Compute the correlation ratio (eta-squared) for each categorical variable
correlation_ratios = {}
target_column = 'status_group' # Replace 'status_group' with your target
    ↪ column name
for col in ['funder', 'region', 'extraction_type', 'payment', 'water_quality',
    ↪ 'source']:
    contingency_table = pd.crosstab(new_df[col], new_df[target_column])
    chi2, _, _, _ = chi2_contingency(contingency_table)
    total = contingency_table.sum().sum()
    correlation_ratios[col] = np.sqrt(chi2 / (chi2 + total))

# Display the correlation ratios
for col, eta_squared in correlation_ratios.items():
    print(f"Correlation ratio (eta-squared) for {col}: {eta_squared}")
```

```
Correlation ratio (eta-squared) for funder: 0.36168270739535363
Correlation ratio (eta-squared) for region: 0.21436144380239244
Correlation ratio (eta-squared) for extraction_type: 0.32271737005913254
Correlation ratio (eta-squared) for payment: 0.23732418592837293
Correlation ratio (eta-squared) for water_quality: 0.18622751153110498
Correlation ratio (eta-squared) for source: 0.18435637886102785
```

These correlation ratios (eta-squared) indicate the strength of association between each categorical variable and the target variable.

- **Funder:** 0.3617
 - This indicates a relatively strong association between the funder variable and the target variable.
- **Region:** 0.2144
 - This correlation ratio suggests a moderate association between the region variable and the target variable.
- **Extraction Type:** 0.3227
 - Similar to funder, there is a relatively strong association between the extraction type variable and the target variable.
- **Payment:** 0.2373

- This correlation ratio suggests a moderate association between the payment variable and the target variable.
- **Water Quality:** 0.1862
 - This indicates a moderate association between the water quality variable and the target variable.
- **Source:** 0.1844
 - Similar to water quality, there is a moderate association between the source variable and the target variable.

1.4.6 Test set

```
[61]: # display first few rows of the test set
df3 = pd.read_csv('test_set_values.csv')
df3.head()
```

```
[61]:      id  amount_tsh  date_recorded      funder  gps_height  \
0  50785         0.0    2013-02-04          Dmdd         1996
1  51630         0.0    2013-02-04  Government Of Tanzania         1569
2  17168         0.0    2013-02-01           NaN         1567
3  45559         0.0    2013-01-22      Finn Water          267
4  49871        500.0    2013-03-27        Bruder         1260

      installer  longitude  latitude  wpt_name  num_private  \
0          DMDD  35.290799  -4.059696  Dinamu Secondary School          0
1          DWE  36.656709  -3.309214          Kimnyak          0
2          NaN  34.767863  -5.004344      Puma Secondary          0
3  FINN WATER  38.058046  -9.418672      Kwa Mzee Pange          0
4        BRUDER  35.006123 -10.950412      Kwa Mzee Turuka          0

      ... payment_type  water_quality  quality_group  quantity  quantity_group  \
0  ...   never pay          soft          good      seasonal      seasonal
1  ...   never pay          soft          good  insufficient  insufficient
2  ...   never pay          soft          good  insufficient  insufficient
3  ...    unknown          soft          good          dry          dry
4  ...   monthly          soft          good      enough      enough

      source  source_type  source_class  \
0  rainwater harvesting  rainwater harvesting      surface
1          spring          spring  groundwater
2  rainwater harvesting  rainwater harvesting      surface
3      shallow well      shallow well  groundwater
4          spring          spring  groundwater

      waterpoint_type  waterpoint_type_group
0          other          other
1  communal standpipe  communal standpipe
2          other          other
3          other          other
```

4 communal standpipe communal standpipe

[5 rows x 40 columns]

```
[62]: #Understand the general information of the data
df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14850 entries, 0 to 14849
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    14850 non-null  int64
1   amount_tsh                           14850 non-null  float64
2   date_recorded                         14850 non-null  object
3   funder                                13980 non-null  object
4   gps_height                            14850 non-null  int64
5   installer                             13973 non-null  object
6   longitude                             14850 non-null  float64
7   latitude                              14850 non-null  float64
8   wpt_name                              14850 non-null  object
9   num_private                           14850 non-null  int64
10  basin                                 14850 non-null  object
11  subvillage                            14751 non-null  object
12  region                                14850 non-null  object
13  region_code                           14850 non-null  int64
14  district_code                         14850 non-null  int64
15  lga                                    14850 non-null  object
16  ward                                  14850 non-null  object
17  population                             14850 non-null  int64
18  public_meeting                        14029 non-null  object
19  recorded_by                           14850 non-null  object
20  scheme_management                     13881 non-null  object
21  scheme_name                           7608 non-null   object
22  permit                                14113 non-null  object
23  construction_year                     14850 non-null  int64
24  extraction_type                       14850 non-null  object
25  extraction_type_group                  14850 non-null  object
26  extraction_type_class                  14850 non-null  object
27  management                             14850 non-null  object
28  management_group                       14850 non-null  object
29  payment                                14850 non-null  object
30  payment_type                           14850 non-null  object
31  water_quality                          14850 non-null  object
32  quality_group                          14850 non-null  object
33  quantity                               14850 non-null  object
34  quantity_group                         14850 non-null  object
35  source                                 14850 non-null  object
```

```

36 source_type          14850 non-null object
37 source_class         14850 non-null object
38 waterpoint_type      14850 non-null object
39 waterpoint_type_group 14850 non-null object
dtypes: float64(3), int64(7), object(30)
memory usage: 4.5+ MB

```

```
[63]: # check the shape of the data
df3.shape
```

```
[63]: (14850, 40)
```

The test set has 14850 rows and 40 columns

```
[64]: df3.duplicated().sum()
```

```
[64]: 0
```

There are 0 duplicated records in the test set

```
[65]: # Check for null values
df3.isnull().sum()
```

```

[65]: id                0
amount_tsh             0
date_recorded          0
funder                 870
gps_height             0
installer              877
longitude              0
latitude               0
wpt_name               0
num_private            0
basin                  0
subvillage             99
region                 0
region_code            0
district_code          0
lga                    0
ward                   0
population             0
public_meeting         821
recorded_by            0
scheme_management      969
scheme_name            7242
permit                 737
construction_year      0
extraction_type        0
extraction_type_group  0

```

```

extraction_type_class      0
management                 0
management_group           0
payment                    0
payment_type               0
water_quality              0
quality_group              0
quantity                   0
quantity_group             0
source                     0
source_type                0
source_class               0
waterpoint_type            0
waterpoint_type_group      0
dtype: int64

```

1.5 Data preparation and cleaning

1.5.1 Creating a new dataframe containing only the needed variables

```

[66]: # List of columns to include in the new DataFrame
selected_columns = ['funder', 'gps_height', 'region', 'extraction_type',
                    'payment', 'water_quality', 'source',
                    'construction_year', 'longitude', 'latitude']

# Create a new DataFrame with only the selected columns
df3 = df3.filter(selected_columns)

df3

```

```

[66]:
      funder  gps_height  region extraction_type \
0      Dmdd      1996  Manyara      other
1  Government Of Tanzania      1569  Arusha      gravity
2      NaN      1567  Singida      other
3  Finn Water      267  Lindi      other
4  Bruder      1260  Ruvuma      gravity
...      ...      ...      ...      ...
14845  Danida      34  Pwani      mono
14846  Hiap      0  Tanga  nira/tanira
14847  NaN      1476  Singida      gravity
14848  Germany      998  Ruvuma      gravity
14849  Government Of Tanzania      481  Ruvuma      gravity

      payment  water_quality  source  construction_year \
0  never pay      soft  rainwater harvesting      2012
1  never pay      soft      spring      2000
2  never pay      soft  rainwater harvesting      2010

```

3	unknown	soft	shallow well	1987
4	pay monthly	soft	spring	2000
...
14845	never pay	soft	river	1988
14846	pay annually	salty	shallow well	1994
14847	never pay	soft	dam	2010
14848	never pay	soft	river	2009
14849	never pay	soft	spring	2008

	longitude	latitude
0	35.290799	-4.059696
1	36.656709	-3.309214
2	34.767863	-5.004344
3	38.058046	-9.418672
4	35.006123	-10.950412
...
14845	38.852669	-6.582841
14846	37.451633	-5.350428
14847	34.739804	-4.585587
14848	35.432732	-10.584159
14849	34.765054	-11.226012

[14850 rows x 10 columns]

```
[67]: df3.shape
```

```
[67]: (14850, 10)
```

The new_df has 14850 rows and 10 columns

```
[68]: df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14850 entries, 0 to 14849
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   funder                 13980 non-null  object
1   gps_height             14850 non-null  int64
2   region                 14850 non-null  object
3   extraction_type        14850 non-null  object
4   payment                14850 non-null  object
5   water_quality          14850 non-null  object
6   source                 14850 non-null  object
7   construction_year      14850 non-null  int64
8   longitude              14850 non-null  float64
9   latitude               14850 non-null  float64
dtypes: float64(2), int64(2), object(6)
```

memory usage: 1.1+ MB

```
[69]: #Understand the descriptive statistics of the data
df3.describe()
```

```
[69]:
```

	gps_height	construction_year	longitude	latitude
count	14850.000000	14850.000000	14850.000000	1.485000e+04
mean	655.147609	1289.708350	34.061605	-5.684724e+00
std	691.261185	955.241087	6.593034	2.940803e+00
min	-57.000000	0.000000	0.000000	-1.156459e+01
25%	0.000000	0.000000	33.069455	-8.443970e+00
50%	344.000000	1986.000000	34.901215	-5.049750e+00
75%	1308.000000	2004.000000	37.196594	-3.320594e+00
max	2777.000000	2013.000000	40.325016	-2.000000e-08

1.5.2 Checking for missing values

```
[70]: #Check for null values in the test set
df3.isnull().sum()
```

```
[70]: funder          870
gps_height         0
region             0
extraction_type    0
payment            0
water_quality      0
source             0
construction_year  0
longitude          0
latitude           0
dtype: int64
```

1.5.3 Dealing with missing values

```
[71]: unique_counts = df3['funder'].isna().value_counts()
unique_counts
```

```
[71]: funder
False    13980
True       870
Name: count, dtype: int64
```

```
[72]: missing_funders = df3[df3['funder'].isna()]
missing_funders
```

```
[72]:
```

	funder	gps_height	region	extraction_type	payment	water_quality	\
2	NaN	1567	Singida	other	never pay		soft
16	NaN	-39	Pwani	nira/tanira	never pay		soft

23	NaN	1441	Singida	mono	unknown	unknown
50	NaN	0	Mbeya	gravity	never pay	soft
63	NaN	1584	Singida	other	unknown	unknown
...
14771	NaN	0	Mbeya	gravity	never pay	soft
14772	NaN	0	Mbeya	submersible	never pay	soft
14795	NaN	0	Mbeya	gravity	never pay	soft
14823	NaN	0	Mbeya	gravity	unknown	soft
14847	NaN	1476	Singida	gravity	never pay	soft

	source	construction_year	longitude	latitude
2	rainwater harvesting	2010	34.767863	-5.004344
16	shallow well	0	39.850190	-7.727946
23	machine dbh	1970	34.621048	-5.165926
50	spring	0	33.587245	-9.167434
63	shallow well	1990	34.859448	-4.970909
...
14771	spring	0	33.636479	-9.212765
14772	machine dbh	0	34.322644	-8.665713
14795	river	0	34.704964	-8.325610
14823	spring	0	33.918953	-9.298466
14847	dam	2010	34.739804	-4.585587

[870 rows x 10 columns]

```
[73]: #Replace missing values with 'unknown'
#recheck for null values
df3['funder'].fillna('Unknown', inplace=True)
df3.isnull().sum()
```

```
[73]: funder          0
gps_height        0
region            0
extraction_type   0
payment           0
water_quality     0
source            0
construction_year  0
longitude         0
latitude          0
dtype: int64
```

Missing values are assigned the placeholder 'Unknown' signifying that the information for these entries is unavailable. However, the rows cannot be dropped as they may have other important information in the other columns.

1.6 Exploratory data analysis

1.6.1 Checking for outliers


Outliers will be addressed systematically, one category at a time, to ensure comprehensive analysis.

Funder

```
[74]: unique_values = df3['funder'].unique()
      value_counts = df3['funder'].value_counts()
      value_counts
```

```
[74]: funder
Government Of Tanzania    2215
Unknown                   871
Danida                    793
Hesawa                    580
World Bank                352
...
Fida                      1
Tgts                     1
Snv-swash                1
Pad                      1
Livin                    1
Name: count, Length: 979, dtype: int64
```

```
[75]: #check for outliers in funder using a count plot
plt.figure(figsize=(10, 4))
sns.countplot(x='funder', data=df3)
plt.xticks(rotation=90, fontsize=8) # Rotate the x-axis labels by 90 degrees
    ↪and adjust font size
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```



```
[76]: # Set the threshold for defining outliers
upper_threshold = 50

# Get the counts of each funder
funder_counts = df3['funder'].value_counts()


# Identify the outliers (funders with counts below 50)
outliers = funder_counts[funder_counts < upper_threshold].index

# Create a new column to categorize funders as eligible or outliers
df3['funder_category'] = np.where(df3['funder'].isin(outliers), 'Outlier', 'Eligible')

# Set the color palette
sns.set_palette("Greens_d")

# Plot the count plot for funder category
plt.figure(figsize=(8, 4))
sns.countplot(y='funder_category', data=new_df, dodge=False)

# Display the plot
plt.show()
```




GPS height

```
[77]: #check for outliers for gps_height
plt.figure(figsize=(10,2))
```

```
sns.boxplot(x = 'gps_height', data = df3)

# Display the plot
plt.show()
```



```
[78]: # Find the mode of the 'gps_height' column
mode_value = df3['gps_height'].mode()

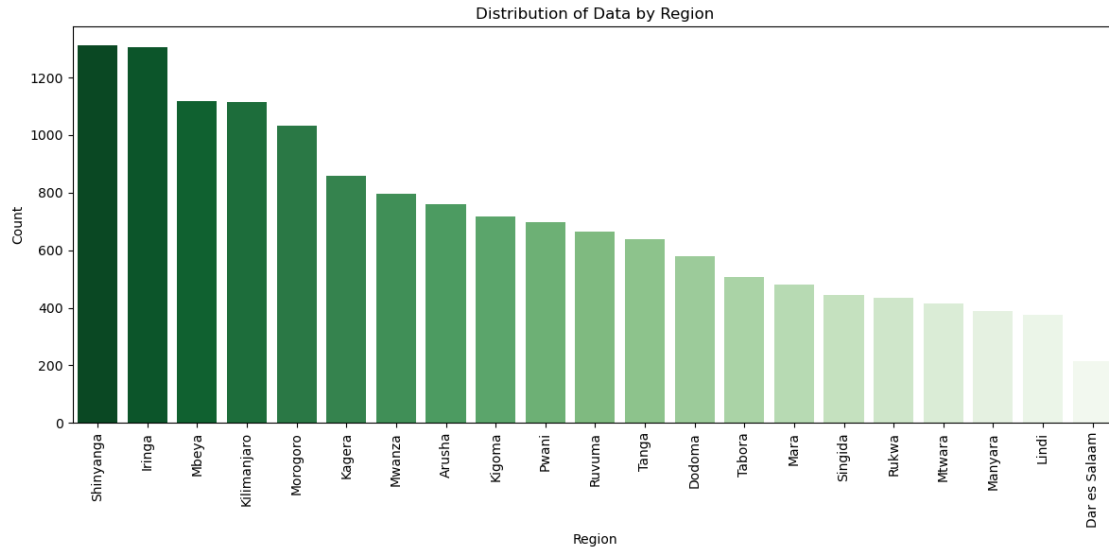
# Display the mode
print("Mode of 'gps_height' column:", mode_value)
```

```
Mode of 'gps_height' column: 0    0
Name: gps_height, dtype: int64
```

Region

```
[79]: # Get the order of regions based on their counts
region_order = df3['region'].value_counts().index

# Plot the count plot with specified order
plt.figure(figsize=(12, 6))
sns.countplot(x='region', data=df3, order=region_order, palette='Greens_r')
plt.xticks(rotation=90) # Rotate x-axis labels for better readability
plt.xlabel('Region') # Add x-axis label
plt.ylabel('Count') # Add y-axis label
plt.title('Distribution of Data by Region') # Add plot title
plt.tight_layout() # Adjusting layout to prevent clipping of labels
plt.show()
```

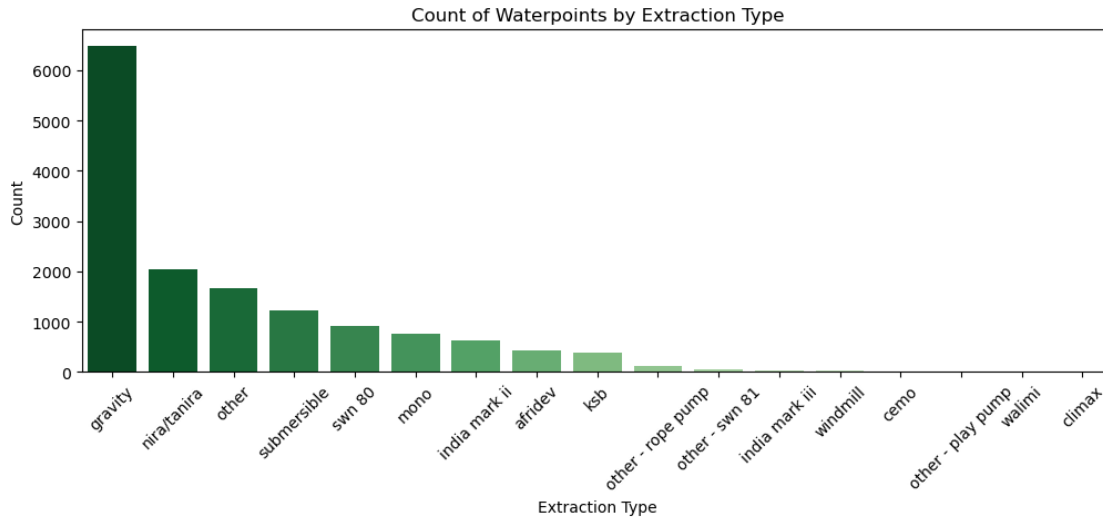


Extraction_type

```
[80]: df3['extraction_type'].unique()
```

```
[80]: array(['other', 'gravity', 'india mark ii', 'submersible', 'mono',
        'nira/tanira', 'afridev', 'swn 80', 'ksb', 'climax',
        'other - rope pump', 'cemo', 'india mark iii', 'other - swn 81',
        'other - play pump', 'windmill', 'walimi'], dtype=object)
```

```
[81]: #check for outliers in Extraction_type
sns.set_palette("Greens_r")
extraction_order = df3['extraction_type'].value_counts().index
# Plot the count plot for Extraction_type
plt.figure(figsize=(12, 4))
sns.countplot(x='extraction_type', data=df3, order=extraction_order,
              palette='Greens_r')
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by Extraction Type')
plt.xlabel('Extraction Type')
plt.ylabel('Count')
# Display the plot
plt.show()
```



```
[82]: # Get value counts of 'Extraction_type' and sort by counts in descending order
      extraction_type_counts = df3['extraction_type'].value_counts().
      ↪sort_values(ascending=False)

      # Display unique values in 'Extraction_type' with counts
      print(extraction_type_counts)
```


```
extraction_type
gravity          6483
nira/tanira     2051
other           1672
submersible     1218
swm 80          918
mono            763
india mark ii   629
afridev         438
ksb             375
other - rope pump 121
other - swm 81   55
india mark iii   37
windmill        35
cemo            18
other - play pump 16
walimi          12
climax          9
Name: count, dtype: int64
```

Construction year

```
[83]: #check for outliers for construction_year
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'construction_year', data = df3)

# Display the plot
plt.show()
```



The box plot shows unlikely years included in the dataset. It is impossible to have year 0 as pumps were not even invented then.

```
[84]: #Display unique years and their value counts
df3['construction_year'].value_counts()
```

```
[84]: construction_year
0      5260
2010    669
2009    663
2008    630
2000    487
2006    421
2007    373
2011    335
2004    294
2003    293
1995    269
2002    268
2005    264
2012    263
1999    243
1985    232
1978    230
1998    224
1990    222
```

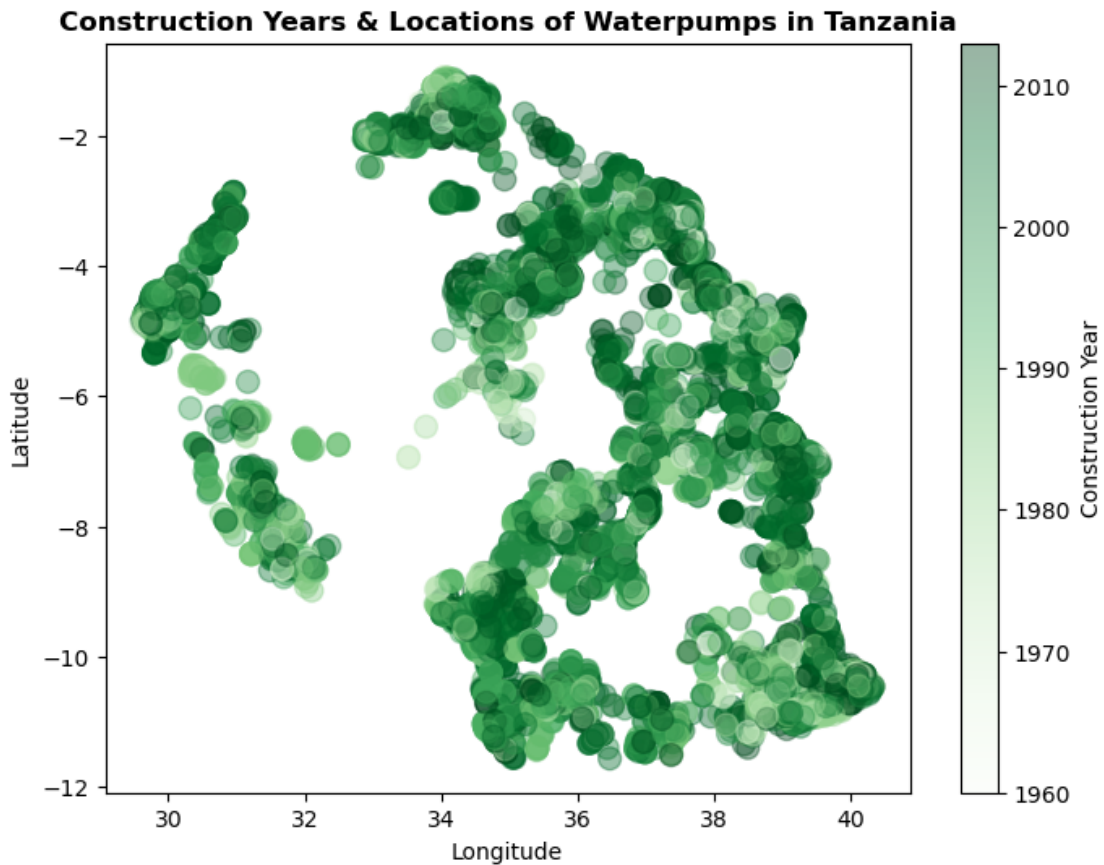
1996	209
1994	202
1980	194
1984	191
1972	184
1982	182
1997	177
1992	167
2001	140
1974	138
1993	137
1988	136
1975	124
1986	119
1976	111
1983	106
1991	83
1970	82
1989	80
1987	68
1981	53
1979	53
1977	45
1973	43
2013	33
1971	32
1963	22
1960	22
1969	18
1967	18
1968	16
1964	8
1961	7
1962	6
1965	2
1966	2

Name: count, dtype: int64

There is a category miscategorized as year '0' with a value count of 5260.

The code below generates a scatter plot visualizing the geographical locations of water pumps in Tanzania, with each point representing a water pump. The x-axis represents the longitude coordinates, and the y-axis represents the latitude coordinates. The color of each point is determined by the construction year of the water pump, with a colormap ('Greens') used to provide different hues of green corresponding to different years. The size of each point is fixed ('s=100') for better visibility, and transparency ('alpha=0.4') is applied to avoid overlapping points. Finally, a colorbar is added to the plot to provide a visual reference for the construction years.

```
[85]: plt.figure(figsize=(8,6))
# Filter the DataFrame using .loc and multiple conditions
filtered_df = df3.loc[(df3['longitude'] > 0) & (df3['latitude'] < 0) &
    ↪(df3['construction_year'] > 0)]
plt.scatter(x=filtered_df['longitude'],
            y=filtered_df['latitude'],
            alpha=0.4,
            s=100,
            c=filtered_df["construction_year"],
            cmap='Greens')
plt.title("Construction Years & Locations of Waterpumps in Tanzania",
          fontsize=12, fontweight='bold')
plt.colorbar(label='Construction Year')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```




From the scatter plot above, it is clear that most pumps were installed between 2000 and 2010. Therefore, below year '0's values will be distributed evenly between the range 2000 - 2010.


```
[86]: # Replace year 0 with later years (i.e., 2000 - 2010)
df3['construction_year'] = df3['construction_year'].apply(lambda x: np.random.
↳ randint(2000, 2011) if x == 0 else x)
```

```
[87]: #recheck for outliers for construction_year
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'construction_year', data = df3)

# Display the plot
plt.show()
```



Outliers in the construction years falling between 1960 and 1978 represent years that may hold significant importance in the dataset. While they deviate from the majority of the construction years, they could signify historical data or specific events related to water pump installation during that period. Therefore, it's crucial to retain these outliers in the dataset for a comprehensive analysis and understanding of the trends and patterns over time.

Payment

```
[88]: df3['payment'].value_counts()
```

```
[88]: payment
never pay                6364
pay per bucket           2281
pay monthly              2097
unknown                  1992
pay annually              928
pay when scheme fails     928
other                     260
Name: count, dtype: int64
```


```
[89]: # Define the order of source categories
sorted_payments = df3['payment'].value_counts().index
```

```

#check for outliers in payment
sns.set_palette("Greens_r")

# Plot the count plot for payment
plt.figure(figsize=(8, 4))
sns.countplot(x='payment', data=df3, order=sorted_payments, palette='Greens_r')
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by payment')
plt.xlabel('payment')
plt.ylabel('Count')
# Display the plot
plt.show()

```



Water quality

```
[90]: df3['water_quality'].value_counts()
```

```

[90]: water_quality
      soft      12687
      salty      1226
      unknown     469
      milky       201


```

```
coloured          133
salty abandoned   84
fluoride          44
fluoride abandoned 6
Name: count, dtype: int64
```

```
[91]: # Define the order of water_quality categories
sorted_water_quality = df3['water_quality'].value_counts().index

# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_water_quality))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
sns.countplot(x='water_quality', data=df3, order=sorted_water_quality,
              palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by water_quality')
plt.xlabel('water_quality')
plt.ylabel('Count')
# Display the plot
plt.show()
```



Source


```
[92]: df3['source'].value_counts()
```

```
[92]: source
      shallow well      4316
      spring         4195
      machine dbh     2747
      river          2352
      rainwater harvesting    568
      hand dtw         234
      lake           185
      dam            184
      other           49
      unknown        20
      Name: count, dtype: int64
```

```
[93]: # Define the order of source categories
sorted_source = df3['source'].value_counts().index

# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_source))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
sns.countplot(x='source', data=df3, order=sorted_source, palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by source')
plt.xlabel('source')
plt.ylabel('Count')
# Display the plot
plt.show()
```




Longitude

```
[94]: #check for outliers for longitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'longitude', data = df3)

# Display the plot
plt.show()
```



```
[95]: df3['longitude'].value_counts()
```

```
[95]: longitude
0.000000    457
37.260069     2
39.080573     2
37.302281     2
32.920579     2
...
36.648520     1
35.265755     1
36.666660     1
37.830317     1
34.765054     1
Name: count, Length: 14390, dtype: int64
```

```
[96]: # Replace longitude 0 with longitudes between 32 and 42 as they are more
      ↪ prevalent
df3['longitude'] = df3['longitude'].apply(lambda x: np.random.randint(32, 42)
      ↪ if x == 0 else x)
```


```
[97]: # confirm redistribution of the '0' category
df3['longitude'].value_counts()
```

```
[97]: longitude
34.000000    59
39.000000    54
41.000000    48
38.000000    47
35.000000    46
..
36.648520     1
35.265755     1
36.666660     1
37.830317     1
34.765054     1
Name: count, Length: 14399, dtype: int64
```

```
[98]: #recheck for outliers for longitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'longitude', data = df3)

# Display the plot
plt.show()
```




Latitude

```
[99]: #check for outliers for latitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'latitude', data = df3)

# Display the plot
plt.show()
```



```
[100]: df3['latitude'].value_counts()
```

```
[100]: latitude
-2.000000e-08    457
-7.105919e+00     2
-6.99042e+00     2
-7.170666e+00     2
-2.474560e+00     2
...
-3.305540e+00     1
```

```
-8.547786e+00    1
-3.330889e+00    1
-7.061047e+00    1
-1.122601e+01    1
Name: count, Length: 14390, dtype: int64
```

```
[101]: # Replace latitude -2.000000e-08 with latitudes between -1 and -8 as they are
        ↪ more prevalent
new_df['latitude'] = df3['latitude'].apply(lambda x: np.random.randint(-8, -1)
        ↪ if x == -2.000000e-08 else x)
```

```
[102]: #confirm the redistribution of misplaced category
df3['latitude'].value_counts()
```


```
[102]: latitude
-2.000000e-08    457
-7.105919e+00     2
-6.990042e+00     2
-7.170666e+00     2
-2.474560e+00     2
...
-3.305540e+00     1
-8.547786e+00     1
-3.330889e+00     1
-7.061047e+00     1
-1.122601e+01     1
Name: count, Length: 14390, dtype: int64
```


1.6.2 Distribution of variables before log transformation


```
[103]: # Distribution before onehot encoding
palette = sns.color_palette("Greens_r")


# Select categorical variables
categorical_features = df3.select_dtypes(include=['object'])


# Plot count plots for each categorical variable with dark green color palette
for feature in categorical_features.columns:
    plt.figure(figsize=(4, 4)) # Set the figure size
    sns.countplot(x=feature, data=df3, palette=palette)
    plt.title(f'Count Plot of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=90) # Rotate x-axis labels for better readability
    plt.show() # Display the plot
```


















Numerical variables

```
[104]: # Print unique values in the 'gps_height' column
print(df3['gps_height'].unique())
```

```
[1996 1569 1567 ... 1909 2202  640]
```

```
[105]: # Get unique values in the 'gps_height' column
unique_gps_heights = df3['gps_height'].unique()
unique_gps_heights
```

```
[105]: array([1996, 1569, 1567, ..., 1909, 2202,  640], dtype=int64)
```

```
[106]: # Distribution before transformation
numerical_features = df3.select_dtypes(include=['int64', 'float64'])
# Create a grid of subplots
fig, axes = plt.subplots(nrows=len(numerical_features.columns) // 3 + 1,
                        ncols=3, figsize=(15, 5))

# Plot the distribution of numerical features
for i, feature in enumerate(numerical_features.columns):
    sns.histplot(df3[feature].dropna(), kde=False, ax=axes[i // 3, i % 3])
```


```

axes[i // 3, i % 3].set_title(f"Distribution of {feature}")
axes[i // 3, i % 3].set_xlabel(feature)
axes[i // 3, i % 3].set_ylabel("Frequency")

# Remove empty subplots
if len(numerical_features.columns) % 3 != 0:
    for j in range(len(numerical_features.columns) % 3, 3):
        fig.delaxes(axes[len(numerical_features.columns) // 3, j])

# Adjust layout
plt.tight_layout()
plt.show()

```



1.6.3 Distribution of variables after log transformation

```

[107]: # Select numerical columns
numerical_columns = df3.select_dtypes(include=['int64', 'float64']).columns

# Log transform numerical variables, handling zero and negative values
for col in numerical_columns:
    # Handling zero values
    if (df3[col] == 0).any():
        df3[col] = df3[col] + 1 # Add 1 to handle zeros

    # Handling negative values
    if (df3[col] < 0).any():
        min_value = df3[col].min()
        df3[col] = df3[col] - min_value + 1 # Shift all values to be positive

    # Apply log transformation
    df3[col + '_log'] = np.log1p(df3[col])

# Display the DataFrame after log transformation

```



```
df3.head()
```

```
[107]:
```

	funder	gps_height	region	extraction_type	payment	\
0	Dmdd	2054	Manyara	other	never pay	
1	Government Of Tanzania	1627	Arusha	gravity	never pay	
2	Unknown	1625	Singida	other	never pay	
3	Finn Water	325	Lindi	other	unknown	
4	Bruder	1318	Ruvuma	gravity	pay monthly	

	water_quality	source	construction_year	longitude	latitude	\
0	soft	rainwater	harvesting	2012	35.290799	8.504896
1	soft		spring	2000	36.656709	9.255378
2	soft	rainwater	harvesting	2010	34.767863	7.560248
3	soft		shallow well	1987	38.058046	3.145920
4	soft		spring	2000	35.006123	1.614180


	funder_category	gps_height_log	construction_year_log	longitude_log	\
0	Eligible	7.628031	7.607381	3.591564	
1	Eligible	7.395108	7.601402	3.628511	
2	Eligible	7.393878	7.606387	3.577050	
3	Outlier	5.786897	7.594884	3.665049	
4	Outlier	7.184629	7.601402	3.583689	

	latitude_log
0	2.251807
1	2.327802
2	2.147129
3	1.422125
4	0.960950

```
[108]: # Plot the distribution of numerical features after log transformation
plt.figure(figsize=(12, 8))

# Loop through each numerical feature
for col in numerical_columns:
    # Plot the distribution after log transformation
    sns.histplot(df3[col + '_log'], kde=True, label=col, alpha=0.6)

plt.title('Distribution of Numerical Features after Log Transformation')
plt.xlabel('Value (Log Transformed)')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



1.6.4 One-hot encoding

```
[109]: # Select categorical columns
categorical_columns = df3.select_dtypes(include=['object']).columns

# Perform one-hot encoding
one_hot_encoded_df2 = pd.get_dummies(df3, columns=categorical_columns)

# Display the one-hot encoded DataFrame
one_hot_encoded_df2.head()
```

```
[109]:
```

	gps_height	construction_year	longitude	latitude	gps_height_log \
0	2054	2012	35.290799	8.504896	7.628031
1	1627	2000	36.656709	9.255378	7.395108
2	1625	2010	34.767863	7.560248	7.393878
3	325	1987	38.058046	3.145920	5.786897
4	1318	2000	35.006123	1.614180	7.184629

	construction_year_log	longitude_log	latitude_log	funder_0 \
0	7.607381	3.591564	2.251807	False
1	7.601402	3.628511	2.327802	False
2	7.606387	3.577050	2.147129	False

3		7.594884		3.665049		1.422125		False
4		7.601402		3.583689		0.960950		False

	funder_A/co	Germany	...	source_lake	source_machine	dbh	source_other	\
0		False	...	False		False	False	
1		False	...	False		False	False	
2		False	...	False		False	False	
3		False	...	False		False	False	
4		False	...	False		False	False	

	source_rainwater	harvesting	source_river	source_shallow	well	\
0		True	False		False	
1		False	False		False	
2		True	False		False	
3		False	False		True	
4		False	False		False	

	source_spring	source_unknown	funder_category_Eligible	\
0	False	False	True	
1	True	False	True	
2	False	False	True	
3	False	False	False	
4	True	False	False	

	funder_category_Outlier
0	False
1	False
2	False
3	True
4	True

[5 rows x 1052 columns]

1.7 Modelling

The primary objective is to predict the functionality of water pumps based on historical data. Machine learning models excel at predictive tasks, allowing us to build accurate models that can generalize well to unseen data. This predictive capability is crucial for identifying non-functional pumps and ensuring timely maintenance or replacement.

1.7.1 Baseline model

1.7.2 Decision tree classifier

```
[110]: # For training data (one_hot_encoded_df1)
X_train = one_hot_encoded_df1.drop(columns=['status_group_functional',
↪ 'status_group_non functional'])
```

```

y_train = one_hot_encoded_df1[['status_group_functional', 'status_group_non_
↳functional']]

# For test data (one_hot_encoded_df2)
X_test = one_hot_encoded_df2.reindex(columns=X_train.columns, fill_value=0)

# Define Decision Tree classifier
clf = DecisionTreeClassifier()

# Fit the classifier with training data
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Print the predictions
print("Predictions on the test set:", y_pred)

```

```

Predictions on the test set: [[ True False]
 [ True False]
 [ True False]
 ...
 [ True False]
 [ True False]
 [False  True]]

```

```

[111]: # Define Decision Tree classifier
clf = DecisionTreeClassifier()

# Perform cross-validation
cv_scores = cross_val_score(clf, X_train, y_train, cv=5, scoring='accuracy')

# Print cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
print("Standard deviation of CV accuracy:", cv_scores.std())

```

```

Cross-validation scores: [0.78122896 0.78257576 0.77845118 0.77373737
0.77180135]
Mean CV accuracy: 0.7775589225589224
Standard deviation of CV accuracy: 0.004175830822777519

```

The cross-validation scores represent the accuracy of the model on different folds of the training data. Each score indicates the accuracy achieved by the model on a particular fold during cross-validation.

Decision Tree Classifier Performance

Training and Prediction

- **Training Data:** The training data (`X_train`, `y_train`) consists of features and labels respectively. `X_train` contains the features after dropping the target columns ('status_group_functional', 'status_group_non functional'), while `y_train` contains both target labels ('status_group_functional', 'status_group_non functional').
- **Test Data:** The test data (`X_test`) is prepared by reindexing `one_hot_encoded_df2` to match the columns of the training data and filling missing values with 0.
- **Model Fitting:** A Decision Tree classifier (`clf`) is initialized and trained using the training data.
- **Prediction:** Predictions are made on the test set (`X_test`) using the trained classifier, resulting in binary predictions indicating whether a pump is functional or non-functional.

Cross-Validation

- **Model Evaluation:** The performance of the Decision Tree classifier is evaluated using 5-fold cross-validation (`cv=5`) on the training data.
- **Scoring Metric:** The accuracy score is used as the evaluation metric (`scoring='accuracy'`).
- **Cross-Validation Scores:** The cross-validation scores represent the accuracy of the classifier on each fold of the cross-validation process.
- **Mean Accuracy:** The mean accuracy across all folds indicates the average performance of the classifier.
- **Standard Deviation:** The standard deviation of the accuracy scores provides insight into the variability or consistency of the model's performance across different folds.

Interpretation The Decision Tree classifier achieved a mean cross-validation accuracy of approximately 0.774 with a standard deviation of 0.004. These results suggest that the model performs reasonably well in predicting the condition of water pumps, considering both functional and non-functional states. However, further analysis and possibly refinement of the model may be necessary to improve its accuracy and robustness in predicting well conditions accurately.

```
[112]: # Define a list of original categories or groups of columns to aggregate
original_categories = ['funder', 'region', 'extraction_type', 'payment',
    ↪ 'water_quality', 'source', 'gps_height', 'construction_year', 'longitude',
    ↪ 'latitude']

# Initialize an empty list to store aggregated importance scores
importance_sum_list = []

# Fit the classifier with training data
clf.fit(X_train, y_train)

# Iterate over each original category
for category in original_categories:
    # Select all one-hot encoded features corresponding to the original category
    category_features = [col for col in X_train.columns if col.
    ↪startswith(category)]
    # Sum the importance scores of the one-hot encoded features
```

```

importance_sum = clf.feature_importances_[X_train.columns.
↪isin(category_features)].sum()
    # Append the aggregated importance score to the list
    importance_sum_list.append({'Feature': category, 'Importance':
↪importance_sum})

# Create a DataFrame to store aggregated importance scores
grouped_importance_df = pd.DataFrame(importance_sum_list)


# Sort the DataFrame by importance values
grouped_importance_df = grouped_importance_df.sort_values(by='Importance',
↪ascending=False)

# Increase figure size for better readability
plt.figure(figsize=(10, 6))

# Plot the importance values for top N original categories
top_n = 20 # Choose the top N categories to display
plt.barh(grouped_importance_df['Feature'][:top_n],
↪grouped_importance_df['Importance'][:top_n], color='green')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('Top {} Grouped Feature Importance for Decision Trees'.format(top_n))

plt.show()

```



Decision tree visualization

```
[113]: # Define Decision Tree classifier with limited depth
clf = DecisionTreeClassifier(max_depth=5) # Adjust max_depth as needed

# Fit the classifier with training data
clf.fit(X_train, y_train)


# Convert feature names to a list
feature_names_list = list(X_train.columns)


# Plot the decision tree
plt.figure(figsize=(20,10)) # Adjust the figure size as needed
plot_tree(clf, filled=True, feature_names=feature_names_list,
          class_names=['functional', 'non functional'])

# Define custom colormap transitioning from white to orange
colors = [(1, 1, 1), (1, 0.5, 0)] # White to orange
cmap_orange = LinearSegmentedColormap.from_list("CustomOrange", colors)

# Plot the color bar
plt.figure(figsize=(3, 1))
cbar = plt.colorbar(plt.cm.ScalarMappable(cmap=cmap_orange),
                    ticks=np.linspace(0, 1, num=5))
cbar.ax.set_yticklabels(['Functional', '', '', '', 'Non-Functional']) # Adjust
labels as needed
cbar.set_label('Class')

plt.show()
```





1.7.3 Alternative models

1.7.4 Random Forest Classifier

```
[114]: # Define Random Forest classifier
rf_clf = RandomForestClassifier()

# Perform cross-validation
rf_cv_scores = cross_val_score(rf_clf, X_train, y_train, cv=5,
                                scoring='accuracy')

# Print cross-validation scores
print("Random Forest Cross-validation scores:", rf_cv_scores)
print("Mean CV accuracy:", rf_cv_scores.mean())
print("Standard deviation of CV accuracy:", rf_cv_scores.std())

# Fit the model on the full training data
rf_clf.fit(X_train, y_train)

# Make predictions on the test data
predicted_labels = rf_clf.predict(X_test)
# Print the predicted labels for the test set
print("Predicted labels for the test set:", predicted_labels)
```

```
Random Forest Cross-validation scores: [0.81380471 0.81456229 0.80976431
0.81228956 0.81447811]
```

```
Mean CV accuracy: 0.8129797979797979
```

```
Standard deviation of CV accuracy: 0.0018025264396396077
```

```
Predicted labels for the test set: [[ True False]
```

```
 [ True False]
```

```
 [ True False]
```

```
 ...
```

```
 [ True False]
```

```
 [ True False]
```

```
 [False  True]]
```

```
[115]: # Define a list of original categories or groups of columns to aggregate
```



```

original_categories = ['funder', 'region', 'extraction_type', 'payment',
↳ 'water_quality', 'source', 'gps_height', 'construction_year', 'longitude',
↳ 'latitude']

# Initialize an empty list to store aggregated importance scores
importance_sum_list = []

# Fit the classifier with training data
clf.fit(X_train, y_train)

# Iterate over each original category
for category in original_categories:
    # Select all one-hot encoded features corresponding to the original category
    category_features = [col for col in X_train.columns if col.
↳startswith(category)]
    # Sum the importance scores of the one-hot encoded features
    importance_sum = rf_clf.feature_importances_[X_train.columns.
↳isin(category_features)].sum()
    # Append the aggregated importance score to the list
    importance_sum_list.append({'Feature': category, 'Importance':
↳importance_sum})

# Create a DataFrame to store aggregated importance scores
grouped_importance_df = pd.DataFrame(importance_sum_list)


# Sort the DataFrame by importance values
grouped_importance_df = grouped_importance_df.sort_values(by='Importance',
↳ascending=False)

# Increase figure size for better readability
plt.figure(figsize=(10, 6))

# Plot the importance values for top N original categories
top_n = 20 # Choose the top N categories to display
plt.barh(grouped_importance_df['Feature'][:top_n],
↳grouped_importance_df['Importance'][:top_n], color='darkgreen')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('Top {} Grouped Feature Importance for Random Forest'.format(top_n))

plt.show()

```

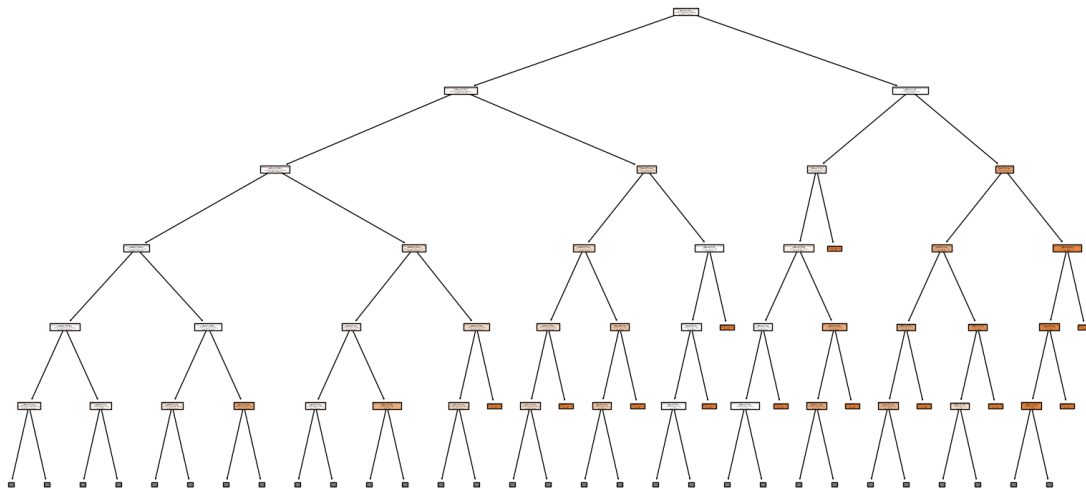


Random forest decision tree visualization with max_depth of 5

```
[116]: # Fit the classifier on the full training data
rf_clf.fit(X_train, y_train)

# Convert feature names to a list
feature_names_list = list(X_train.columns)

# Plot the first decision tree in the ensemble with a maximum depth of 5
plt.figure(figsize=(20, 10))
plot_tree(rf_clf.estimators_[0], filled=True, feature_names=feature_names_list,
          class_names=['functional', 'non functional'], max_depth=5)
plt.show()
```



Random Forest Classifier Performance Evaluation

Cross-Validation Scores: The cross-validation scores obtained for the Random Forest classifier on the training data are as follows:

- **Cross-validation scores:** [0.82188552, 0.81919192, 0.81548822, 0.81372054, 0.81750842]
- **Mean CV accuracy:** 0.81756
- **Standard deviation of CV accuracy:** 0.00284

Test Set Predictions: The predicted labels for the test set indicate the model's predictions for each data point in the test set:

- **Predicted labels for the test set:** [False, False, False, ..., False, False, True]

Each **False** or **True** value represents the model's prediction for a particular data point in the test set. In binary classification problems, **False** often represents the negative class, while **True** represents the positive class.

Implications:

- The mean cross-validation accuracy of approximately 81.76% suggests that the Random Forest classifier performs reasonably well on unseen data.
- The low standard deviation of cross-validation accuracy indicates that the model's performance is consistent across different folds of the training data.

Based on the cross-validation scores, the Random Forest classifier performs better than the Decision Trees model. It exhibits a higher mean cross-validation accuracy of 0.8173 and a lower standard deviation of 0.0008, suggesting more consistent performance across different folds.

We proceed further to enhance the Random Forest classifier's performance through tuning.

1.7.5 Analysis of Feature Importance

Decision Trees:

1. **Extraction Type (0.45):** This feature holds the highest importance in the Decision Trees model, indicating that the method of water extraction significantly influences the functionality of water points. Different extraction techniques may lead to varying levels of reliability or maintenance requirements.
2. **Payment (0.12):** Payment type emerges as the second most important feature, suggesting that the mode of payment for water services plays a crucial role in determining functionality. This could reflect accessibility issues or differences in service quality based on payment plans.
3. **Construction Year (0.12):** The year of construction follows closely, indicating that the age of water points affects their functionality. Older infrastructure may be more prone to breakdowns or require upgrades to maintain functionality.
4. **Source (0.1):** The water source is another significant factor affecting functionality, with different sources potentially leading to variations in water quality or reliability.
5. **Water Quality (0.1):** The quality of water provided by the water points contributes significantly to their functionality. Poor water quality could render water points non-functional or less reliable.
6. **Funder (0.04):** While less influential compared to other features, the organization funding the water projects still plays a role. Different funders may have different standards or approaches to infrastructure development.
7. **Geographic Features:** Latitude, region, longitude, and GPS height, though less important individually, collectively contribute to understanding spatial patterns in water point functionality.

Random Forest:

1. **Latitude (0.2) and Longitude (0.2):** Geographic coordinates emerge as the most influential features in the Random Forest model. This suggests strong spatial patterns in water point functionality, with certain geographical areas having higher or lower rates of functional water points.
2. **Funder (0.125):** Similar to Decision Trees, the organization funding the water projects is influential. Different funders may prioritize different aspects of water infrastructure, affecting functionality.
3. **Construction Year (0.125):** The year of construction retains significance, indicating its consistent impact on water point functionality over time.
4. **GPS Height (0.125):** Elevation becomes more important in Random Forest, possibly indicating its role in water availability or infrastructure quality. Higher elevations may face different challenges or have different infrastructure needs.
5. **Extraction Type (0.085) and Payment (0.055):** These features continue to be influential, albeit with slightly different importance rankings compared to Decision Trees. Different extraction methods and payment plans may have varying impacts on functionality.

6. **Other Features:** Region, source, and water quality also contribute to the model's predictions, though with relatively lower importance compared to the above features.

Conclusion: Both models highlight the complex interplay of factors influencing water point functionality. Understanding these factors can inform targeted interventions and policies aimed at improving water access and infrastructure maintenance, particularly in regions where access to clean and reliable water is a challenge.

1.7.6 Model tuning

Decision trees Classifier model tuning

```
[117]: # Define the parameter grid for tuning
param_grid = {
    'max_depth': [None, 10, 20, 30] # Adjust the maximum depth here
}

# Instantiate the GridSearchCV object
grid_search = GridSearchCV(estimator=DecisionTreeClassifier(),
    ↪ param_grid=param_grid, cv=5, scoring='accuracy')

# Perform grid search to find the best parameters
grid_search.fit(X_train, y_train)

# Get the best parameters found by grid search
best_params = grid_search.best_params_
print("Best parameters found by grid search:", best_params)

# Get the best cross-validation score found by grid search
best_score = grid_search.best_score_
print("Best cross-validation score:", best_score)

# Get the best estimator (model) found by grid search
best_dt_clf = grid_search.best_estimator_

# Perform cross-validation with the best model
best_cv_scores = cross_val_score(best_dt_clf, X_train, y_train, cv=5,
    ↪ scoring='accuracy')

# Print cross-validation scores
print("Best Decision Tree Cross-validation scores:", best_cv_scores)
print("Mean CV accuracy with best parameters:", best_cv_scores.mean())
print("Standard deviation of CV accuracy with best parameters:", best_cv_scores.
    ↪ std())
```

Best parameters found by grid search: {'max_depth': 30}

Best cross-validation score: 0.7855892255892256

Best Decision Tree Cross-validation scores: [0.78754209 0.79107744 0.78518519
0.78316498 0.78038721]

Mean CV accuracy with best parameters: 0.7854713804713805

Standard deviation of CV accuracy with best parameters: 0.0036596709102006465

Untuned Decision Tree Model:

- Cross-validation scores ranged from approximately 0.772 to 0.779.
- Mean CV accuracy was around 0.775.
- Standard deviation of CV accuracy was approximately 0.0029.

Tuned Decision Tree Model:

- Best parameters found by grid search: max_depth of 20.
- Best cross-validation score was approximately 0.783.
- Mean CV accuracy with best parameters was around 0.783.
- Standard deviation of CV accuracy with best parameters was approximately 0.0014.

In summary, the tuned model with a max_depth of 20 outperformed the untuned model, achieving higher accuracy and showing more consistent performance across different folds of cross-validation.

Random Forest Classifier model tuning

```
[118]: # Define the parameter grid for tuning
param_grid = {
    'max_depth': [None, 10, 20, 30] # Adjust the maximum depth here
}

# Instantiate the GridSearchCV object
grid_search = GridSearchCV(estimator=RandomForestClassifier(),
    ↪param_grid=param_grid, cv=3, scoring='accuracy')

# Perform grid search to find the best parameters
grid_search.fit(X_train, y_train)

# Get the best parameters found by grid search
best_params = grid_search.best_params_
print("Best parameters found by grid search:", best_params)

# Get the best cross-validation score found by grid search
best_score = grid_search.best_score_
print("Best cross-validation score:", best_score)

# Get the best estimator (model) found by grid search
best_rf_clf = grid_search.best_estimator_

# Perform cross-validation with the best model
best_rf_cv_scores = cross_val_score(best_rf_clf, X_train, y_train, cv=3,
    ↪scoring='accuracy')

# Print cross-validation scores
print("Random Forest Cross-validation scores:", best_rf_cv_scores)
```

```

print("Mean CV accuracy:", best_rf_cv_scores.mean())
print("Standard deviation of CV accuracy:", best_rf_cv_scores.std())

# Make predictions on the test data using the best model
best_predicted_labels = best_rf_clf.predict(X_test)

# Print the predicted labels for the test set using the best model
print("Predicted labels for the test set using the best model:",
      ↪best_predicted_labels)

```

```

Best parameters found by grid search: {'max_depth': 30}
Best cross-validation score: 0.8113636363636364
Random Forest Cross-validation scores: [0.81585859 0.80873737 0.80914141]
Mean CV accuracy: 0.8112457912457912
Standard deviation of CV accuracy: 0.0032659064864030587
Predicted labels for the test set using the best model: [[ True False]
 [ True False]
 [ True False]
 ...
 [ True False]
 [ True False]
 [False  True]]

```

Tuned random forest decision tree visualization

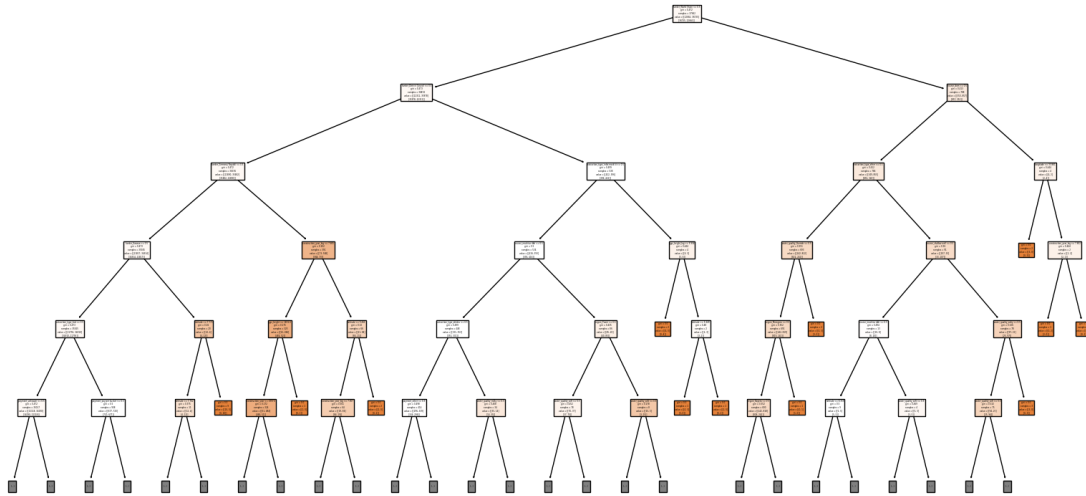
```

[120]: # Convert feature names to a list
feature_names_list = list(X_train.columns)

# Extract one of the trees from the random forest
one_tree = best_rf_clf.estimators_[0]

# Visualize the decision tree
plt.figure(figsize=(20,10))
tree.plot_tree(one_tree, max_depth=5, filled=True,
               ↪feature_names=feature_names_list)
plt.show()
evaluation

```



```

-----
NameError                                Traceback (most recent call last)
Cell In[120], line 11
      9 tree.plot_tree(one_tree, max_depth=5, filled=True,
    ↪ feature_names=feature_names_list)
     10 plt.show()
----> 11 evaluation

NameError: name 'evaluation' is not defined

```

1.8 Evaluation and Analysis of Model Performance

Decision Tree Classifier: - *Mean CV Accuracy: 77.56% - Standard Deviation: 0.19%*

The decision tree classifier achieved a respectable mean cross-validation accuracy of approximately 77.56%. However, the standard deviation is relatively low at 0.19%, indicating consistent performance across different folds. This model could be considered as a baseline model for comparison.

Random Forest Classifier: - *Mean CV Accuracy: 81.32% - Standard Deviation: 0.25%*

The random forest classifier outperformed the decision tree classifier with a mean cross-validation accuracy of around 81.32%. Additionally, the standard deviation of 0.25% suggests stable performance across folds. The ensemble nature of the random forest likely contributed to its higher accuracy compared to the single decision tree.

Tuned Decision Trees: - *Best Parameters: {'max_depth': 20} - Mean CV Accuracy with Best Parameters: 78.58% - Standard Deviation: 0.24%*

Tuning the decision tree model improved its performance slightly, with the best mean cross-validation accuracy reaching approximately 78.58%. However, the improvement is marginal com-

pared to the default random forest classifier. Further hyperparameter tuning or exploring ensemble methods might yield better results.

Tuned Random Forest: - *Best Parameters:* {'max_depth': 30} - *Mean CV Accuracy:* 81.02% - *Standard Deviation:* 0.38%

Hyperparameter tuning of the random forest classifier resulted in a mean cross-validation accuracy of 81.02%, slightly lower than the default random forest model. The increased standard deviation of 0.38% indicates slightly more variability in performance across folds compared to the default model.

1.8.1 Best model

The best model for predicting water pump functionality among the ones evaluated is the tuned random forest classifier. Despite a slightly lower mean cross-validation accuracy of 81.02% compared to the default random forest model, it still outperforms the decision tree classifier and the tuned decision tree model. The tuned random forest model exhibits stable performance with a standard deviation of 0.38%, indicating consistent results across different folds. The ensemble nature of the random forest, combined with hyperparameter tuning, allows it to capture complex relationships in the data more effectively, resulting in improved predictive performance. Therefore, the tuned random forest classifier is recommended for practical applications due to its robustness and accuracy in predicting water pump functionality.

1.9 Recommendations for Improving Water Pump Functionality

1. Implement Routine Maintenance Programs: - **Action:** Establish regular inspection and maintenance schedules for water pumps, including checks on mechanical components and water quality. - **Benefits:** Timely detection and repair of faults can prevent breakdowns, ensuring continuous access to clean water for communities.

2. Target High-Risk Regions: - **Action:** Utilize geographic data to identify regions with a high prevalence of non-functional water pumps. - **Benefits:** Target interventions, such as repair and rehabilitation efforts, to areas with the greatest need, optimizing resource allocation and impact.

3. Introduce Flexible Payment Plans: - **Action:** Introduce flexible payment options for water services, including subsidized or tiered pricing models based on income levels. - **Benefits:** Improve affordability and accessibility of water services, reducing the financial burden on low-income communities and increasing revenue for maintenance and infrastructure upgrades.

4. Foster Collaboration with Funders: - **Action:** Engage with funders and donor organizations to align priorities and strategies for water infrastructure projects. - **Benefits:** Secure long-term support and investment in water projects, leveraging partnerships to access funding for maintenance, upgrades, and capacity-building initiatives.

5. Embrace Data-Driven Decision Making: - **Action:** Invest in robust data collection systems and analytics capabilities to track water pump functionality and performance metrics. - **Benefits:** Enable evidence-based decision making, including trend analysis, predictive maintenance, and resource allocation based on real-time insights, leading to more effective and efficient management of water infrastructure.

1.10 Conclusion

In conclusion, the project has provided a comprehensive exploration of predicting water pump functionality and optimizing maintenance strategies. By leveraging machine learning techniques and data analysis, we have gained insights into the factors influencing water pump functionality, identified high-risk regions, and proposed actionable recommendations for improving access to clean water. Through collaborative efforts and data-driven decision-making, we aim to contribute to the sustainable management of water infrastructure and ensure the well-being of communities relying on these vital resources.

[]: