

# Pump functionality prediction

March 12, 2024

## 1 Tanzanian faulty pumps prediction

### 1.1 Problem statement

In Tanzania, access to clean and potable water is essential for the health and well-being of its citizens. However, many water pumps across the country are faulty, leading to water shortages and posing significant health risks to communities. To address this issue and promote access to clean water, we aim to develop a predictive model that can identify faulty water pumps based on various features such as pump age, location, type, and condition. By accurately predicting which water pumps are faulty, authorities and organizations can prioritize maintenance and repair efforts, ensuring that clean and safe water is readily available to all Tanzanians. Stakeholder: The Ministry of Water in Tanzania is a key stakeholder in addressing the issue of faulty water pumps and promoting access to clean and potable water across the country. As the government body responsible for water resource management and infrastructure development, the Ministry plays a crucial role in ensuring that water supply systems are well-maintained and functional. By leveraging predictive modeling to identify faulty water pumps, the Ministry can efficiently allocate resources for maintenance and repair activities, thereby improving the reliability and accessibility of clean water for Tanzanian communities.

#### 1.1.1 Objectives

1. To predict the functionality of water pumps: Develop a predictive model to classify water pumps into functional, non-functional, and functional needs repair categories based on various features such as `amount_tsh`, `gps_height`, `waterpoint_type`, and others.
2. To identify factors influencing water pump functionality: Conduct exploratory data analysis to identify the key factors (e.g., funder, installer, water quality) that influence the functionality of water pumps and their maintenance needs.
3. To optimize water pump maintenance strategies: Use historical data on water pump failures and repairs to optimize maintenance schedules and resource allocation, ensuring timely repairs and minimizing downtime of water pumps.
4. To assess the geographical distribution of water pump functionality: Analyze the geographical distribution of functional and non-functional water pumps to identify regions with high repair needs and prioritize interventions for improved access to clean water.
5. To evaluate the impact of funding sources on water pump functionality: Investigate the relationship between funding sources and water pump functionality to assess the effectiveness of different funding mechanisms in ensuring sustainable access to clean water.

## 1.2 Data understanding

```
[1]: # import relevant modules
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import random
from scipy import stats
from scipy.stats import chi2_contingency
from scipy.stats import f_oneway
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn import tree
from matplotlib.colors import LinearSegmentedColormap
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

import warnings
warnings.filterwarnings("ignore")
```

```
[2]: # display first few rows of the labels set
df1 = pd.read_csv('training_set_labels.csv')
df1.head()
```

```
[2]:      id  status_group
0  69572      functional
1   8776      functional
2  34310      functional
3  67743  non functional
4  19728      functional
```

```
[3]: # display first few rows of the training set
df2 = pd.read_csv('training_set_values.csv')
df2.head()
```

```
[3]:      id  amount_tsh  date_recorded  funder  gps_height  installer \
0  69572      6000.0    2011-03-14      Roman      1390      Roman
1   8776         0.0    2013-03-06    Grumeti      1399    GRUMETI
2  34310        25.0    2013-02-25  Lottery Club      686  World vision
3  67743         0.0    2013-01-28      Unicef      263      UNICEF
4  19728         0.0    2011-07-13  Action In A         0      Artisan

      longitude  latitude  wpt_name  num_private  ...  payment_type  \
```

0	34.938093	-9.856322	none	0	...	annually
1	34.698766	-2.147466	Zahanati	0	...	never pay
2	37.460664	-3.821329	Kwa Mahundi	0	...	per bucket
3	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	...	never pay
4	31.130847	-1.825359	Shuleni	0	...	never pay

	water_quality	quality_group	quantity	quantity_group	\
0	soft	good	enough	enough	
1	soft	good	insufficient	insufficient	
2	soft	good	enough	enough	
3	soft	good	dry	dry	
4	soft	good	seasonal	seasonal	

	source	source_type	source_class	\
0	spring	spring	groundwater	
1	rainwater harvesting	rainwater harvesting	surface	
2	dam	dam	surface	
3	machine dbh	borehole	groundwater	
4	rainwater harvesting	rainwater harvesting	surface	

	waterpoint_type	waterpoint_type_group
0	communal standpipe	communal standpipe
1	communal standpipe	communal standpipe
2	communal standpipe multiple	communal standpipe
3	communal standpipe multiple	communal standpipe
4	communal standpipe	communal standpipe

[5 rows x 40 columns]

```
[4]: merged_df = pd.merge(df1, df2, on='id')
merged_df.head()
```

[4]:	id	status_group	amount_tsh	date_recorded	funder	gps_height	\
0	69572	functional	6000.0	2011-03-14	Roman	1390	
1	8776	functional	0.0	2013-03-06	Grumeti	1399	
2	34310	functional	25.0	2013-02-25	Lottery Club	686	
3	67743	non functional	0.0	2013-01-28	Unicef	263	
4	19728	functional	0.0	2011-07-13	Action In A	0	

	installer	longitude	latitude	wpt_name	...	\
0	Roman	34.938093	-9.856322	none	...	
1	GRUMETI	34.698766	-2.147466	Zahanati	...	
2	World vision	37.460664	-3.821329	Kwa Mahundi	...	
3	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu	...	
4	Artisan	31.130847	-1.825359	Shuleni	...	

	payment_type	water_quality	quality_group	quantity	quantity_group	\
--	--------------	---------------	---------------	----------	----------------	---

0	annually	soft	good	enough	enough
1	never pay	soft	good	insufficient	insufficient
2	per bucket	soft	good	enough	enough
3	never pay	soft	good	dry	dry
4	never pay	soft	good	seasonal	seasonal

	source	source_type	source_class	\
0	spring	spring	groundwater	
1	rainwater harvesting	rainwater harvesting	surface	
2	dam	dam	surface	
3	machine dbh	borehole	groundwater	
4	rainwater harvesting	rainwater harvesting	surface	

	waterpoint_type	waterpoint_type_group
0	communal standpipe	communal standpipe
1	communal standpipe	communal standpipe
2	communal standpipe multiple	communal standpipe
3	communal standpipe multiple	communal standpipe
4	communal standpipe	communal standpipe

[5 rows x 41 columns]

```
[5]: #Understand the general information of the data
merged_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     59400 non-null  int64
1   status_group           59400 non-null  object
2   amount_tsh             59400 non-null  float64
3   date_recorded          59400 non-null  object
4   funder                 55763 non-null  object
5   gps_height             59400 non-null  int64
6   installer              55745 non-null  object
7   longitude              59400 non-null  float64
8   latitude               59400 non-null  float64
9   wpt_name               59398 non-null  object
10  num_private            59400 non-null  int64
11  basin                  59400 non-null  object
12  subvillage             59029 non-null  object
13  region                 59400 non-null  object
14  region_code            59400 non-null  int64
15  district_code          59400 non-null  int64
16  lga                    59400 non-null  object
17  ward                   59400 non-null  object
```

```

18 population          59400 non-null int64
19 public_meeting      56066 non-null object
20 recorded_by         59400 non-null object
21 scheme_management    55522 non-null object
22 scheme_name          30590 non-null object
23 permit              56344 non-null object
24 construction_year    59400 non-null int64
25 extraction_type      59400 non-null object
26 extraction_type_group 59400 non-null object
27 extraction_type_class 59400 non-null object
28 management           59400 non-null object
29 management_group     59400 non-null object
30 payment              59400 non-null object
31 payment_type         59400 non-null object
32 water_quality        59400 non-null object
33 quality_group        59400 non-null object
34 quantity             59400 non-null object
35 quantity_group       59400 non-null object
36 source               59400 non-null object
37 source_type          59400 non-null object
38 source_class         59400 non-null object
39 waterpoint_type      59400 non-null object
40 waterpoint_type_group 59400 non-null object
dtypes: float64(3), int64(7), object(31)
memory usage: 18.6+ MB

```

```
[6]: # check the shape of the data
merged_df.shape
```

```
[6]: (59400, 41)
```

```
[7]: merged_df.duplicated().sum()
```

```
[7]: 0
```

### 1.2.1 Defining Variables

**Independent Variables (Predictors):** - Amount\_tsh - Gps\_height - Waterpoint\_type - Funder - Installer - Water\_quality - Payment\_type - Region - Latitude - Longitude

**Dependent Variable (Target):** - Status\_group: Indicates the status of the waterpoint (e.g., functional, non-functional, functional needs repair).

Categorical Variables: - **status\_group**: The status of the waterpoint (e.g., functional, non-functional, functional needs repair). - **funder**: The organization or individual who funded the waterpoint. - **region**: The geographic region where the waterpoint is located. - **extraction\_type**: The mechanism used to extract water from the waterpoint. - **payment**: The type of payment required to access the waterpoint. - **water\_quality**: The quality of the water provided by the waterpoint. - **source**: The source of the water (e.g., river, well, spring).

Numeric Variables: - **gps\_height**: The altitude of the waterpoint. - **construction\_year**: The year when the waterpoint was constructed. - **longitude**: The longitude coordinate of the waterpoint. - **latitude**: The latitude coordinate of the waterpoint.

```
[8]: merged_df.isnull().sum()
```

```
[8]: id                                0
     status_group                      0
     amount_tsh                        0
     date_recorded                     0
     funder                            3637
     gps_height                        0
     installer                         3655
     longitude                         0
     latitude                         0
     wpt_name                          2
     num_private                       0
     basin                             0
     subvillage                        371
     region                           0
     region_code                       0
     district_code                     0
     lga                               0
     ward                              0
     population                        0
     public_meeting                    3334
     recorded_by                       0
     scheme_management                 3878
     scheme_name                       28810
     permit                           3056
     construction_year                 0
     extraction_type                   0
     extraction_type_group              0
     extraction_type_class              0
     management                        0
     management_group                  0
     payment                           0
     payment_type                       0
     water_quality                      0
     quality_group                      0
     quantity                          0
     quantity_group                    0
     source                            0
     source_type                       0
     source_class                       0
     waterpoint_type                   0
     waterpoint_type_group              0
```

dtype: int64

## 1.3 Data preparation and cleaning

### 1.3.1 Train set

Creating a new dataframe containing only the needed variables

```
[9]: # List of columns to include in the new DataFrame
selected_columns = ['status_group', 'funder', 'gps_height', 'region', '
↳ 'extraction_type', 'payment', 'water_quality', 'source',
                    'construction_year', 'longitude', 'latitude']

# Create a new DataFrame with only the selected columns
new_df = merged_df.filter(selected_columns)

new_df.head()
```

```
[9]:
```

	status_group	funder	gps_height	region	extraction_type	\
0	functional	Roman	1390	Iringa	gravity	
1	functional	Grumeti	1399	Mara	gravity	
2	functional	Lottery Club	686	Manyara	gravity	
3	non functional	Unicef	263	Mtwara	submersible	
4	functional	Action In A	0	Kagera	gravity	

	payment	water_quality	source	construction_year	\
0	pay annually	soft	spring	1999	
1	never pay	soft	rainwater harvesting	2010	
2	pay per bucket	soft	dam	2009	
3	never pay	soft	machine dbh	1986	
4	never pay	soft	rainwater harvesting	0	

	longitude	latitude
0	34.938093	-9.856322
1	34.698766	-2.147466
2	37.460664	-3.821329
3	38.486161	-11.155298
4	31.130847	-1.825359

```
[10]: # Learn the shape of the data
new_df.shape
```

```
[10]: (59400, 11)
```

The new\_df has 59400 rows and 14 columns

```
[11]: # Understand the general infomation of the dataset
new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 59400 entries, 0 to 59399
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_group          59400 non-null  object
1   funder                55763 non-null  object
2   gps_height            59400 non-null  int64
3   region               59400 non-null  object
4   extraction_type       59400 non-null  object
5   payment              59400 non-null  object
6   water_quality         59400 non-null  object
7   source               59400 non-null  object
8   construction_year     59400 non-null  int64
9   longitude             59400 non-null  float64
10  latitude              59400 non-null  float64
dtypes: float64(2), int64(2), object(7)
memory usage: 5.0+ MB

```

```
[12]: #Understand the descriptive statistics of the data
new_df.describe()
```

```

[12]:      gps_height  construction_year  longitude  latitude
count  59400.000000      59400.000000  59400.000000  5.940000e+04
mean    668.297239      1300.652475    34.077427  -5.706033e+00
std     693.116350       951.620547     6.567432   2.946019e+00
min     -90.000000        0.000000     0.000000  -1.164944e+01
25%       0.000000        0.000000    33.090347  -8.540621e+00
50%      369.000000      1986.000000    34.908743  -5.021597e+00
75%     1319.250000      2004.000000    37.178387  -3.326156e+00
max      2770.000000      2013.000000    40.345193  -2.000000e-08

```

### 1.3.2 Checking for missing values

```
[13]: #Check for null values in the training set
new_df.isnull().sum()
```

```

[13]: status_group      0
funder                3637
gps_height            0
region               0
extraction_type       0
payment              0
water_quality         0
source               0
construction_year     0
longitude             0
latitude             0
dtype: int64

```



### 1.3.3 Dealing with missing values

```
[14]: # Preview the unique categories in funder
unique_counts = new_df['funder'].isna().value_counts()
unique_counts
```

```
[14]: funder
False      55763
True        3637
Name: count, dtype: int64
```

```
[15]: missing_funders = new_df[new_df['funder'].isna()]
missing_funders
```

```
[15]:
```

	status_group	funder	gps_height	region	extraction_type	\
34	functional	NaN	-41	Pwani	nira/tanira	
43	non functional	NaN	1642	Singida	mono	
47	functional	NaN	0	Mbeya	gravity	
65	non functional	NaN	1415	Singida	mono	
71	non functional	NaN	0	Mbeya	gravity	
...	...	...	...	...	...	
59357	non functional	NaN	1635	Singida	nira/tanira	
59366	functional	NaN	1541	Singida	nira/tanira	
59370	functional	NaN	1154	Kigoma	other	
59376	non functional	NaN	1581	Singida	other	
59397	functional	NaN	0	Mbeya	swn 80	

	payment	water_quality	source	construction_year	longitude	\
34	never pay	salty	shallow well	0	39.812912	
43	unknown	unknown	machine dbh	1980	34.967789	
47	never pay	soft	spring	0	33.540607	
65	unknown	unknown	machine dbh	1970	34.621598	
71	never pay	soft	river	0	34.462228	
...	...	...	...	...	...	
59357	unknown	unknown	shallow well	1980	34.971841	
59366	never pay	soft	shallow well	2000	34.765729	
59370	pay monthly	unknown	unknown	0	30.058731	
59376	unknown	unknown	shallow well	1990	34.821039	
59397	pay monthly	fluoride	machine dbh	0	34.017087	

	latitude
34	-7.889986
43	-4.628921
47	-9.172905
65	-5.173136
71	-8.575780
...	...
59357	-5.098362

```
59366 -5.027725
59370 -4.902633
59376 -5.076258
59397 -8.750434
```

```
[3637 rows x 11 columns]
```

```
[16]: # Replace null values with 'unknown' in funder
# recheck for null values
new_df['funder'].fillna('Unknown', inplace=True)
new_df.isnull().sum()
```

```
[16]: status_group      0
funder                0
gps_height            0
region                0
extraction_type       0
payment              0
water_quality         0
source                0
construction_year     0
longitude             0
latitude              0
dtype: int64
```

Missing values are assigned the placeholder 'Unknown' signifying that the information for these entries is unavailable. However, the rows cannot be dropped as they may have other important information in the other columns.

## 1.4 Exploratory data analysis

### 1.4.1 Checking for outliers

The dataset contains a substantial amount of categorical data, necessitating the handling of outliers during Exploratory Data Analysis (EDA). Since box plots are unsuitable for detecting outliers in categorical columns, count plots become the primary tool. These count plots not only reveal the presence or absence of outliers but also illustrate the distribution of the data. Outliers will be addressed systematically, one category at a time, to ensure comprehensive analysis.

#### Status\_group

```
[17]: #check unique categories in status_group
unique_values = new_df['status_group'].unique()
unique_values
```

```
[17]: array(['functional', 'non functional', 'functional needs repair'],
      dtype=object)
```


```
[18]: # merge 'functional need repair' into 'functional' for the sake of a binary
      ↪ classification
new_df['status_group'] = new_df['status_group'].replace('functional needs
      ↪ repair', 'functional')

# Check the unique values again
print(new_df['status_group'].value_counts())
```

```
status_group
functional      36576
non functional   22824
Name: count, dtype: int64
```

```
[19]: # check for outliers in status_group

palette = sns.color_palette("Greens_r", 3)
# Plotting countplot
plt.figure(figsize=(4,4))
sns.countplot(x="status_group", data=new_df, palette=palette)
plt.title("Distribution of status_group")
plt.xlabel("status_group")
plt.ylabel("Count")
plt.show()
```




There are no outliers in the status group

## Funder

```
[20]: #check unique categories in funder
unique_values = new_df['funder'].unique()
value_counts = new_df['funder'].value_counts()
value_counts
```

```
[20]: funder
Government Of Tanzania    9084
Unknown                   3641
Danida                     3114
Hesawa                     2202
Rwssp                      1374
...
Rarymond Ekura             1
Justine Marwa              1
Municipal Council          1
Afdp                       1
Samlo                      1
Name: count, Length: 1896, dtype: int64
```

```
[21]: #check for outliers in funder using a count plot
plt.figure(figsize=(10, 4))
sns.countplot(x='funder', data=new_df)
plt.xticks(rotation=90, fontsize=8) # Rotate the x-axis labels by 90 degrees
    ↳and adjust font size
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```



```
[22]: # classify eligible funders and outliers in a binned countplot
# Set the threshold for defining outliers
```

```

upper_threshold = 50

# Get the counts of each funder
funder_counts = new_df['funder'].value_counts()

# Identify the outliers (funders with counts below 50)
outliers = funder_counts[funder_counts < upper_threshold].index


# Create a new column to categorize funders as eligible or outliers
new_df['funder_category'] = np.where(new_df['funder'].isin(outliers),
    ↪ 'Outlier', 'Eligible')

# Set the color palette
sns.set_palette("Greens_d")

# Plot the count plot for funder category
plt.figure(figsize=(8, 4))
sns.countplot(y='funder_category', data=new_df, dodge=False)

# Display the plot
plt.show()

```



**Eligible Funders vs. Outliers** The count plot above illustrates the distribution of funders categorized as “Eligible” and “Outlier” based on the specified thresholds. Here’s a summary of the findings:

- **Eligible Funders:** These are funders with a count falling within the specified thresholds (between 10 and 600).

- **Outliers:** These are funders with a count below 10 or above 600.


As observed in the plot, the number of outliers is significantly higher than the count of eligible funders. However, it's important to note that we cannot disregard the outliers as they may contain valuable insights or represent specific cases of interest.

### GPS height

```
[23]: #check for outliers for gps_height
plt.figure(figsize=(10,2))

sns.boxplot(x = 'gps_height', data = new_df)

# Display the plot
plt.show()
```



```
[24]: # Find the mode of the 'gps_height' column to understand left skewness
mode_value = new_df['gps_height'].mode()
mode_value
```

```
[24]: 0      0
      Name: gps_height, dtype: int64
```

```
[25]: # Find the lowest value in gps_height
lowest_value = new_df['gps_height'].min()
lowest_value
```

```
[25]: -90
```

The mode of the 'gps\_height' column is 0, indicating that this value is the most common within the dataset. As box plots rely on quartiles to determine their position, the prevalence of 0 strongly influences the box plot's positioning.

With the mode close to 0, it's likely that the median (second quartile) aligns closely with this value, resulting in a box plot skewed towards lower values. Consequently, the majority of the data tends to concentrate towards the lower end of the scale.

The presence of a whisker starting below 0 at -90 may suggest data recorded at elevations below a predefined reference datum. In this context, these points below 0 are not considered outliers.

On the other hand, the longer upper whisker compared to the lower one suggests greater dispersion or variability in the upper range of the data (maximum). This could hint at the presence of outliers or extreme values towards higher elevations.

However, it's important to note that we are not removing these outliers. They might represent genuine data points and carry valuable information. Blindly removing them could lead to the loss of valuable insights and potentially bias the analysis or conclusions drawn from the data.


## Region

```
[26]: #check unique categories in region
unique_values = new_df['region'].unique()
unique_values
```

```
[26]: array(['Iringa', 'Mara', 'Manyara', 'Mtwara', 'Kagera', 'Tanga',
        'Shinyanga', 'Tabora', 'Pwani', 'Ruvuma', 'Kilimanjaro', 'Rukwa',
        'Mwanza', 'Kigoma', 'Lindi', 'Dodoma', 'Arusha', 'Mbeya',
        'Singida', 'Morogoro', 'Dar es Salaam'], dtype=object)
```

```
[27]: # Get the order of regions based on their counts
region_order = new_df['region'].value_counts().index

# Plot the count plot with specified order
plt.figure(figsize=(12, 6))
sns.countplot(x='region', data=new_df, order=region_order, palette='Greens_r')
plt.xticks(rotation=90) # Rotate x-axis labels for better readability
plt.xlabel('Region') # Add x-axis label
plt.ylabel('Count') # Add y-axis label
plt.title('Distribution of Data by Region') # Add plot title
plt.tight_layout() # Adjust layout to prevent clipping of labels
plt.show()
```

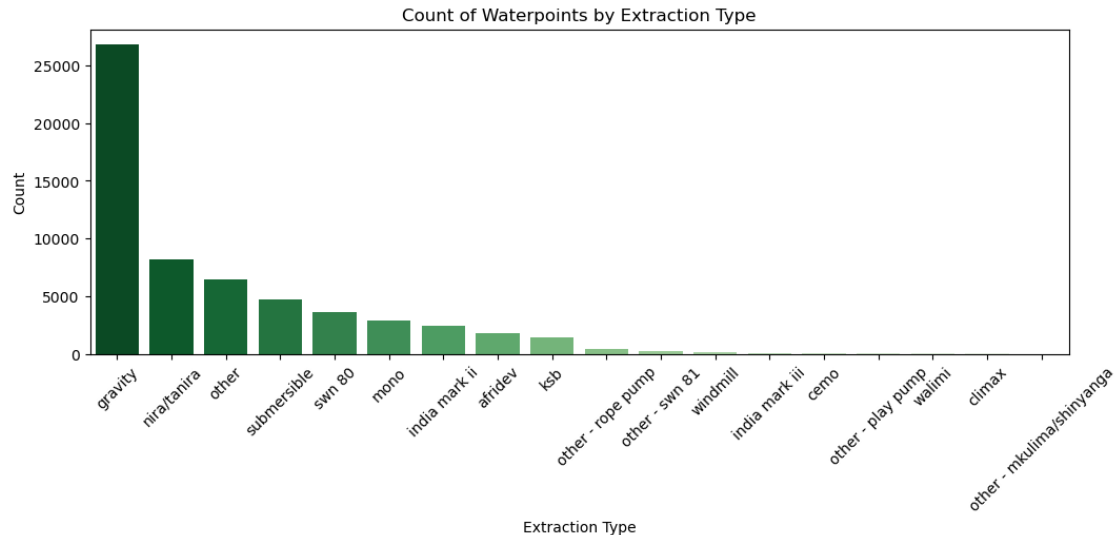


## Extraction\_type

```
[28]: new_df['extraction_type'].unique()
```

```
[28]: array(['gravity', 'submersible', 'swn 80', 'nira/tanira', 'india mark ii',  
        'other', 'ksb', 'mono', 'windmill', 'afridev', 'other - rope pump',  
        'india mark iii', 'other - swn 81', 'other - play pump', 'cemo',  
        'climax', 'walimi', 'other - mkulima/shinyanga'], dtype=object)
```

```
[29]: #check for outliers in Extraction_type  
sns.set_palette("Greens_r")  
extraction_order = new_df['extraction_type'].value_counts().index  
# Plot the count plot for Extraction_type  
plt.figure(figsize=(12, 4))  
sns.countplot(x='extraction_type', data=new_df, order=extraction_order,  
             palette='Greens_r')  
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees  
plt.title('Count of Waterpoints by Extraction Type')  
plt.xlabel('Extraction Type')  
plt.ylabel('Count')  
# Display the plot  
plt.show()
```



```
[30]: # Get value counts of 'Extraction_type' and sort by counts in descending order  
extraction_type_counts = new_df['extraction_type'].value_counts().  
    sort_values(ascending=False)  
  
# Display unique values in 'Extraction_type' with counts
```



```
print(extraction_type_counts)
```

```
extraction_type
gravity                26780
nira/tanira           8154
other                 6430
submersible          4764
swn 80               3670
mono                 2865
india mark ii        2400
afridev              1770
ksb                  1415
other - rope pump     451
other - swn 81        229
windmill             117
india mark iii        98
cemo                  90
other - play pump     85
walimi                48
climax                32
other - mkulima/shinyanga 2
Name: count, dtype: int64
```

In the 'Extraction\_type' column, the majority of water pumps fall into the following categories:


Gravity: 26,780 pumps Nira/Tanira: 8,154 pumps Other: 6,430 pumps Submersible: 4,764 pumps Swn 80: 3,670 pumps Mono: 2,865 pumps India Mark II: 2,400 pumps Afridev: 1,770 pumps KSB: 1,415 pumps However, there are some categories with notably fewer pumps, such as 'Other - Rope Pump', 'Other - Swn 81', 'Windmill', 'India Mark III', 'CEMO', 'Other - Play Pump', 'Walimi', 'Climax', and 'Other - Mkulima/Shinyanga'. These could be outliers regarding in terms of pump d. This could distribute indicating less common or specialized therefore we cannot simply remove them as they may hold significance in the dataset.pump types.

### Construction year

```
[31]: #check for outliers for construction_year
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'construction_year', data = new_df)

# Display the plot
plt.show()
```



The box plot shows unlikely years included in the dataset. It is impossible to have year 0 as pumps were not even invented then. Therefore further investigation will be conducted.

There is a category miscategorized as year '0' with a very high value count of 20709.

```
[32]: #Display unique years and their value counts
new_df['construction_year'].value_counts()
```

```
[32]: construction_year
0      20709
2010    2645
2008    2613
2009    2533
2000    2091
2007    1587
2006    1471
2003    1286
2011    1256
2004    1123
2012    1084
2002    1075
1978    1037
1995    1014
2005    1011
1999     979
1998     966
1990     954
1985     945
1980     811
1996     811
1984     779
1982     744
1994     738
1972     708
```

1974	676
1997	644
1992	640
1993	608
2001	540
1988	521
1983	488
1975	437
1986	434
1976	414
1970	411
1991	324
1989	316
1987	302
1981	238
1977	202
1979	192
1973	184
2013	176
1971	145
1960	102
1967	88
1963	85
1968	77
1969	59
1964	40
1962	30
1961	21
1965	19
1966	17

Name: count, dtype: int64


The code below we generate a scatter plot visualizing the geographical locations of water pumps in Tanzania, with each point representing a water pump. The x-axis represents the longitude coordinates, and the y-axis represents the latitude coordinates. The color of each point is determined by the construction year of the water pump, with a colormap ('Greens') used to provide different hues of green corresponding to different years. The size of each point is fixed ('s=100') for better visibility, and transparency ('alpha=0.4') is applied to avoid overlapping points. Finally, a colorbar is added to the plot to provide a visual reference for the construction years.

```
[33]: # Plot a scatter plot of construction years and locations of water pumps
plt.figure(figsize=(8,6))
# Filter the DataFrame using .loc and multiple conditions
filtered_df = new_df.loc[(new_df['longitude'] > 0) & (new_df['latitude'] < 0) &
    ↪(new_df['construction_year'] > 0)]
plt.scatter(x=filtered_df['longitude'],
            y=filtered_df['latitude'],
```

```

alpha=0.4,
s=100,
c=filtered_df["construction_year"],
cmap='Greens')
plt.title("Construction Years & Locations of Waterpumps in Tanzania",
          fontsize=12, fontweight='bold')
plt.colorbar(label='Construction Year')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()

```



From the scatter plot above, it is clear that most pumps were installed between 2000 and 2010. Therefore, below year '0's values will be distributed evenly between the range 2000 - 2010.

```


[34]: # Replace year 0 with later years (i.e., 2000 - 2010)
new_df['construction_year'] = new_df['construction_year'].apply(lambda x: np.
    random.randint(2000, 2011) if x == 0 else x)

```

```
[35]: #recheck for outliers for construction_year
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'construction_year', data = new_df)

# Display the plot
plt.show()
```



Outliers in the construction years falling between 1960 and 1978 represent years that may hold significant importance in the dataset. While they deviate from the majority of the construction years, they could signify historical data or specific events related to water pump installation during that period. Therefore, it's crucial to retain these outliers in the dataset for a comprehensive analysis and understanding of the trends and patterns over time.

### Payment

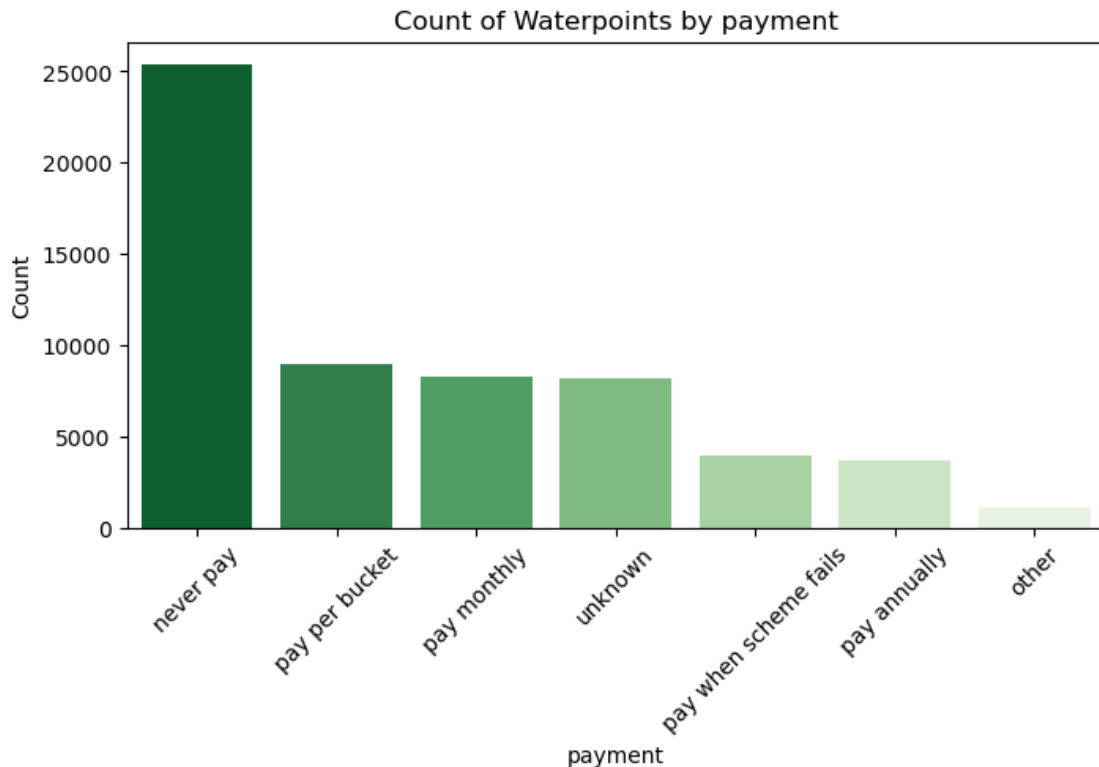
```
[36]: # check unique categories and their value counts in payment
new_df['payment'].value_counts()
```

```
[36]: payment
never pay                25348
pay per bucket           8985
pay monthly              8300
unknown                  8157
pay when scheme fails    3914
pay annually             3642
other                    1054
Name: count, dtype: int64
```

```
[37]: # Define the order of source categories
sorted_payments = new_df['payment'].value_counts().index

#check for outliers in payment
sns.set_palette("Greens_r")
```

```
# Plot the count plot for payment
plt.figure(figsize=(8, 4))
sns.countplot(x='payment', data=new_df, order=sorted_payments,
             palette='Greens_r')
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by payment')
plt.xlabel('payment')
plt.ylabel('Count')
# Display the plot
plt.show()
```



The distribution of payment types, as observed in the count plot above, reveals an interesting trend. The “never pay” category dominates the dataset, indicating that a significant portion of water points in the dataset do not require any payment. This could be due to various reasons, such as government subsidies or community initiatives aimed at providing free access to water.

In contrast, the paid categories exhibit a more even distribution, with multiple categories having similar counts. This distribution suggests that while there are options for paid water access, they are not as prevalent as the “never pay” category. This observation might be attributed to the socioeconomic factors prevalent in the area. Residents who cannot afford paid water services may opt for the free “never pay” option, resulting in its higher prevalence in the dataset.

Therefore, the presence of multiple paid categories with similar counts does not necessarily indicate

outliers. Instead, it reflects the diverse payment options available and the socioeconomic dynamics influencing water access in the region.

### Water quality

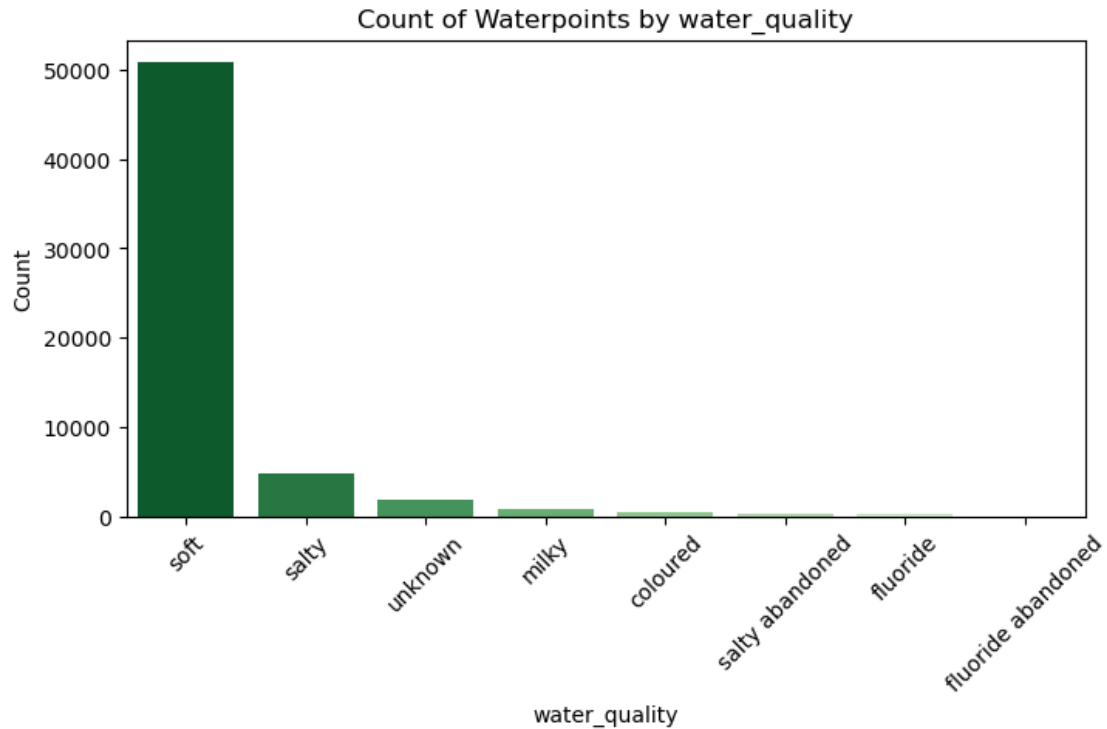
```
[38]: # check unique categories and their value counts in water_quality
new_df['water_quality'].value_counts()
```

```
[38]: water_quality
      soft          50818
      salty          4856
      unknown        1876
      milky           804
      coloured        490
      salty abandoned  339
      fluoride        200
      fluoride abandoned 17
      Name: count, dtype: int64
```

```
[39]: # Define the order of water_quality categories
sorted_water_quality = new_df['water_quality'].value_counts().index

# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_water_quality))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
sns.countplot(x='water_quality', data=new_df, order=sorted_water_quality,
              palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by water_quality')
plt.xlabel('water_quality')
plt.ylabel('Count')
# Display the plot
plt.show()
```



The count plot above indicates the most prevalent category is “soft”. This indicates that most water sources provide satisfactory water quality. It could also mean most people prefer soft water leading to its prevalence. Next, we find the “salty” category, which exhibits a considerably lower count in comparison to “soft.” This suggests that while some water sources may have elevated salinity levels, they are less common than those providing “soft” water. Then there is the “milky” and “coloured” categories, which may raise concerns regarding water quality. These categories, while not as frequent as “soft” or “salty,” suggest the presence of impurities or contaminants that could affect the desirability of the water.

Another notable category is “salty abandoned,” which indicates water sources that have been abandoned likely due to high salinity levels. This category, although less common, highlights instances where water quality issues have led to the abandonment of waterpoints.

Lastly, we have “fluoride” and “fluoride abandoned” categories, which indicate the presence of fluoride in the water. While fluoride is beneficial in controlled amounts for dental health, excessive levels can be harmful. The presence of “fluoride abandoned” suggests instances where water sources have been abandoned due to excessive fluoride.

Generally the plot reveals a diverse landscape of water quality categories, with “soft” being the predominant category. While certain categories may raise concerns, such as “salty abandoned” or “fluoride,” they do not appear to be outliers but rather indicative of the range of water quality issues present across waterbodies in our dataset.

## Source




```
[40]: # check unique categories and their value counts in source
new_df['source'].value_counts()
```

```
[40]: source
      spring          17021
      shallow well    16824
      machine dbh     11075
      river           9612
      rainwater harvesting 2295
      hand dtw         874
      lake             765
      dam              656
      other            212
      unknown          66
      Name: count, dtype: int64
```

```
[41]: # Define the order of source categories
sorted_source = new_df['source'].value_counts().index

# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_source))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
sns.countplot(x='source', data=new_df, order=sorted_source, palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by source')
plt.xlabel('source')
plt.ylabel('Count')
# Display the plot
plt.show()
```



“Spring” and “shallow well” emerge as the most prevalent sources, followed closely by “machine dbh” and “river.” These categories exhibit relatively high counts, indicating their widespread usage as water sources.

Next in line is “rainwater harvesting,” although its count is notably lower compared to the preceding categories. “Hand dtw,” “lake,” and “dam” follow, each with decreasing counts.

Finally, we have the categories of “unknown” and “other,” which appear to represent sources with less distinct categorization or sources not captured by the specified categories.


Overall, while there is variation in the counts across different water source categories, there are no outliers that significantly deviate from the expected distribution. Instead, the distribution reflects the diverse range of water sources utilized across waterpoints in our dataset.

### Longitude

```
[42]: # check for outliers for longitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'longitude', data = new_df)

# Display the plot
plt.show()
```



The plot shows presence of outliers.

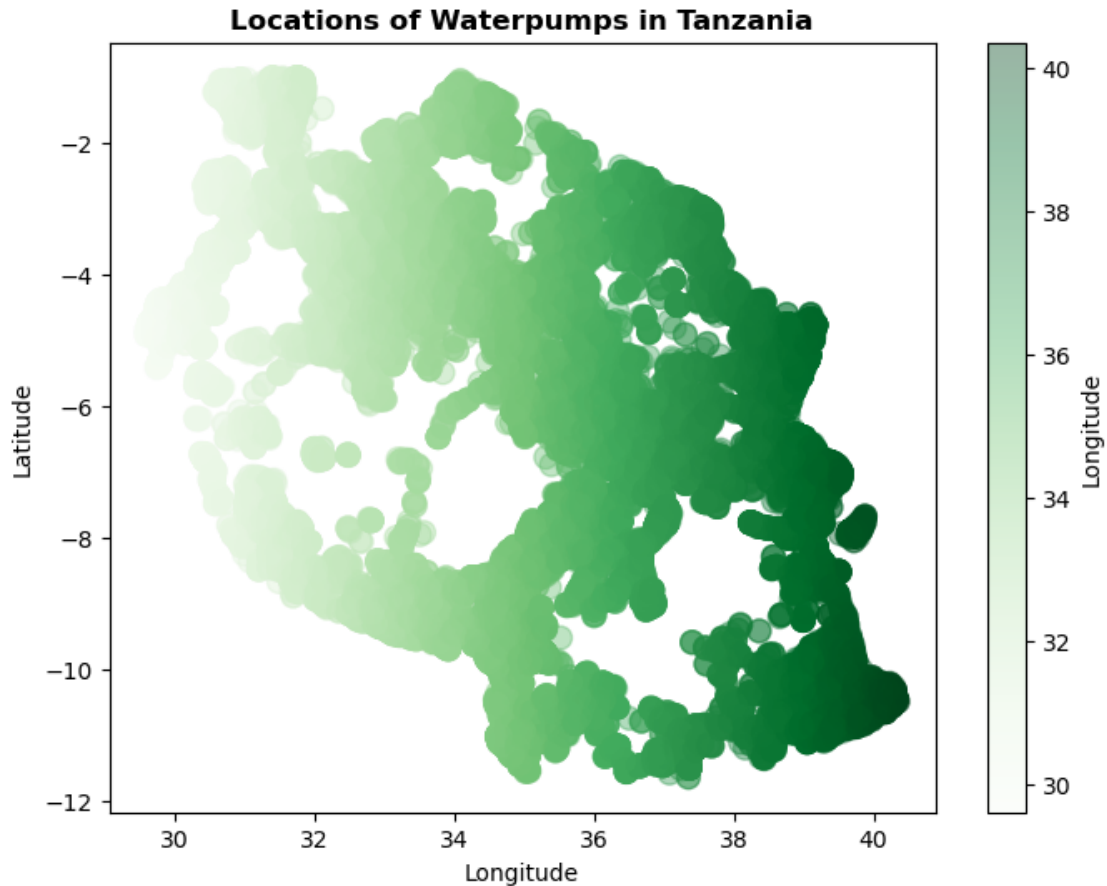
```
[43]: # check unique categories and their value counts in longitude
new_df['longitude'].value_counts()
```

```
[43]: longitude
0.000000      1812
37.375717         2
38.340501         2
39.086183         2
33.005032         2
...
35.885754         1
36.626541         1
37.333530         1
38.970078         1
38.104048         1
Name: count, Length: 57516, dtype: int64
```

There is a huge count miscategorized as longitude '0'. Below that is dealt with by being redistributed to a range with more frequent occurrence.

```
[44]: # plot a scatter plot to show the majority of longitude points
plt.figure(figsize=(8,6))
# Filter the DataFrame using .loc and multiple conditions
filtered_df = new_df.loc[(new_df['longitude'] > 0) & (new_df['latitude'] < 0)]
plt.scatter(x=filtered_df['longitude'],
            y=filtered_df['latitude'],
            alpha=0.4,
            s=100,
            c=filtered_df["longitude"],
            cmap='Greens')
plt.title("Locations of Waterpumps in Tanzania",
          fontsize=12, fontweight='bold')
```

```
plt.colorbar(label='Longitude')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



It's evident that there are more water pumps located at longitudes greater than 34 degrees than those located at longitudes less than 34 degrees. This suggests a higher concentration of water points towards the eastern side of the region under consideration. Longitude, representing the east-west position on the Earth's surface, indicates that the area to the east of 34 degrees longitude may have higher population densities or other factors contributing to the need for more water access points compared to the western region.

```
[45]: # Replace longitude 0 with longitudes between 32 and 42 as they are more
      ↪ prevalent
new_df['longitude'] = new_df['longitude'].apply(lambda x: np.random.randint(32,
      ↪ 42) if x == 0 else x)
```


```
[46]: # confirm redistribution of the '0' category
new_df['longitude'].value_counts()
```

```
[46]: longitude
      37.000000    200
      41.000000    193
      33.000000    187
      39.000000    186
      40.000000    183
      ...
      35.885754     1
      36.626541     1
      37.333530     1
      38.970078     1
      38.104048     1
      Name: count, Length: 57525, dtype: int64
```

```
[47]: #recheck for outliers for longitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'longitude', data = new_df)

# Display the plot
plt.show()
```




### Latitude

```
[48]: #check for outliers for latitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'latitude', data = new_df)

# Display the plot
plt.show()
```



There does not seem to have outliers in latitude but further analysis is still necessary.

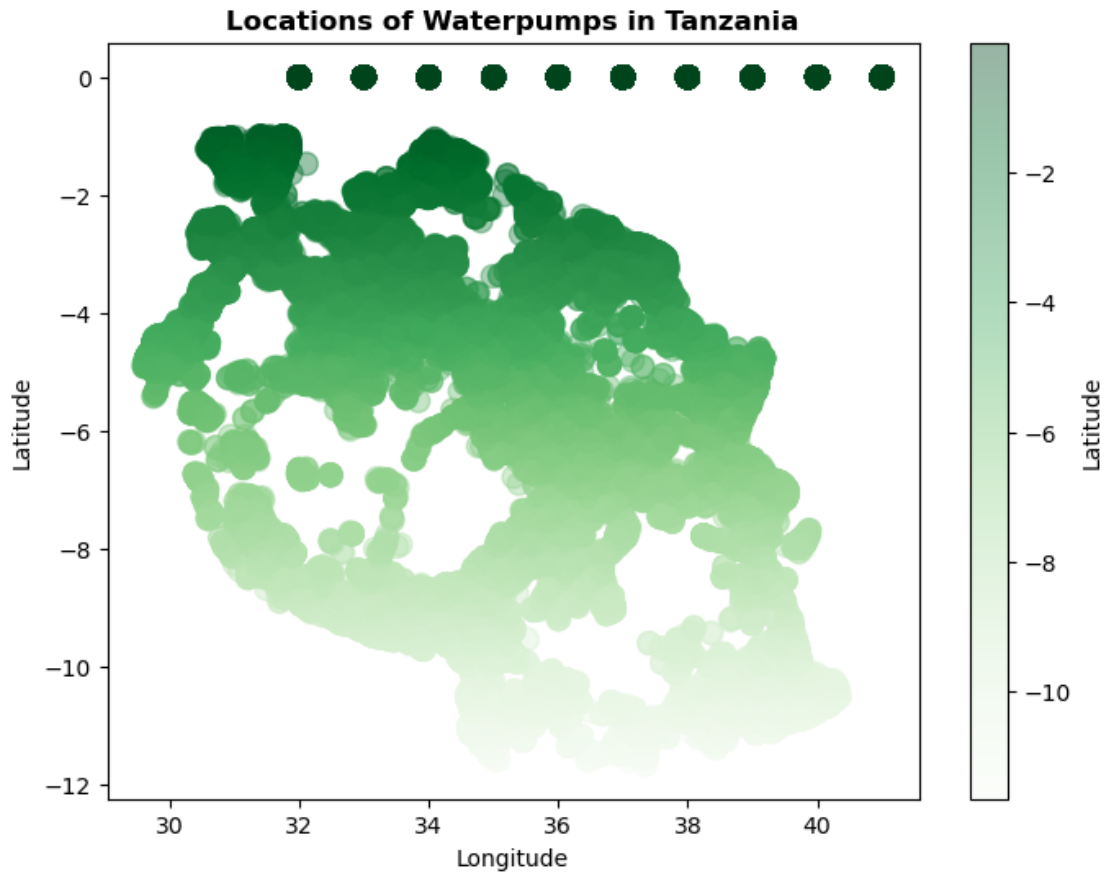
```
[49]: new_df['latitude'].value_counts()
```

```
[49]: latitude
-2.000000e-08    1812
-6.985842e+00      2
-6.980220e+00      2
-2.476680e+00      2
-6.978263e+00      2
...
-3.287619e+00      1
-8.234989e+00      1
-3.268579e+00      1
-1.146053e+01      1
-6.747464e+00      1
Name: count, Length: 57517, dtype: int64
```

There is a latitude that seems to be miscategorized (-2.000000e-08 1812). This will be dealt with below.

```
[50]: # Plot a scatter plot to show the majority of latitude points
plt.figure(figsize=(8,6))
# Filter the DataFrame using .loc and multiple conditions
filtered_df = new_df.loc[(new_df['longitude'] > 0) & (new_df['latitude'] < 0)]
plt.scatter(x=filtered_df['longitude'],
            y=filtered_df['latitude'],
            alpha=0.4,
            s=100,
            c=filtered_df["latitude"],
            cmap='Greens')
plt.title("Locations of Waterpumps in Tanzania",
          fontsize=12, fontweight='bold')
plt.colorbar(label='Latitude')
```

```
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



A scatter plot of pumps in relevance to latitude indicates more pumps between latitudes -1 and -8. The misplaced class will be evenly distributed in this range.

```
[51]: # Replace latitude -2.000000e-08 with latitudes between -1 and -8 as they are
      ↪ more prevalent
new_df['latitude'] = new_df['latitude'].apply(lambda x: np.random.randint(-8,
      ↪ -1) if x == -2.000000e-08 else x)
```

```
[52]: #confirm the redistribution of misplaced category
new_df['latitude'].value_counts()
```

```
[52]: latitude
      -7.000000    286
      -5.000000    273
      -8.000000    268
```

```

-2.000000    264
-6.000000    256
...
-3.287619     1
-8.234989     1
-3.268579     1
-11.460531    1
-6.747464     1
Name: count, Length: 57523, dtype: int64

```

## 1.4.2 Distribution of variables before log transformation

### Categorical variables

```


[53]: # Distribution before onehot encoding
palette = sns.color_palette("Greens_r")


# Select categorical variables
categorical_features = new_df.select_dtypes(include=['object'])


# Plot count plots for each categorical variable with dark green color palette
for feature in categorical_features.columns:
    plt.figure(figsize=(4, 4)) # Set the figure size
    sns.countplot(x=feature, data=new_df, palette=palette)
    plt.title(f'Count Plot of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=90) # Rotate x-axis labels for better readability
    plt.show() # Display the plot


```























### Numerical variables

```
[54]: # Distribution before transformation
numerical_features = new_df.select_dtypes(include=['int64', 'float64'])

# Create a grid of subplots
fig, axes = plt.subplots(nrows=len(numerical_features.columns) // 3 + 1,
                        ncols=3, figsize=(15, 5))


# Plot the distribution of numerical features
for i, feature in enumerate(numerical_features.columns):
    sns.histplot(new_df[feature].dropna(), kde=False, ax=axes[i // 3, i % 3])
    axes[i // 3, i % 3].set_title(f"Distribution of {feature}")
    axes[i // 3, i % 3].set_xlabel(feature)
    axes[i // 3, i % 3].set_ylabel("Frequency")

# Remove empty subplots
if len(numerical_features.columns) % 3 != 0:
    for j in range(len(numerical_features.columns) % 3, 3):
        fig.delaxes(axes[len(numerical_features.columns) // 3, j])

# Adjust layout
```



```
plt.tight_layout()
plt.show()
```



### 1.4.3 Distribution of variables after log transformation

```
[55]: # Select numerical columns
numerical_columns = new_df.select_dtypes(include=['int64', 'float64']).columns

# Log transform numerical variables, handling zero and negative values
for col in numerical_columns:
    # Handling zero values
    if (new_df[col] == 0).any():
        new_df[col] = new_df[col] + 1 # Add 1 to handle zeros

    # Handling negative values
    if (new_df[col] < 0).any():
        min_value = new_df[col].min()
        new_df[col] = new_df[col] - min_value + 1 # Shift all values to be
    ↪ positive

    # Apply log transformation
    new_df[col + '_log'] = np.log1p(new_df[col])

# Display the DataFrame after log transformation
new_df.head()
```

```
[55]:
```

	status_group	funder	gps_height	region	extraction_type	\
0	functional	Roman	1481	Iringa	gravity	
1	functional	Grumeti	1490	Mara	gravity	
2	functional	Lottery Club	777	Manyara	gravity	
3	non functional	Unicef	354	Mtwara	submersible	
4	functional	Action In A	91	Kagera	gravity	

	payment	water_quality		source	construction_year	\
0	pay annually	soft		spring	1999	
1	never pay	soft	rainwater	harvesting	2010	
2	pay per bucket	soft		dam	2009	
3	never pay	soft		machine dbh	1986	
4	never pay	soft	rainwater	harvesting	2006	

	longitude	latitude	funder_category	gps_height_log	\
0	34.938093	2.793118	Eligible	7.301148	
1	34.698766	10.501974	Eligible	7.307202	
2	37.460664	8.828112	Outlier	6.656727	
3	38.486161	1.494142	Eligible	5.872118	
4	31.130847	10.824081	Outlier	4.521789	

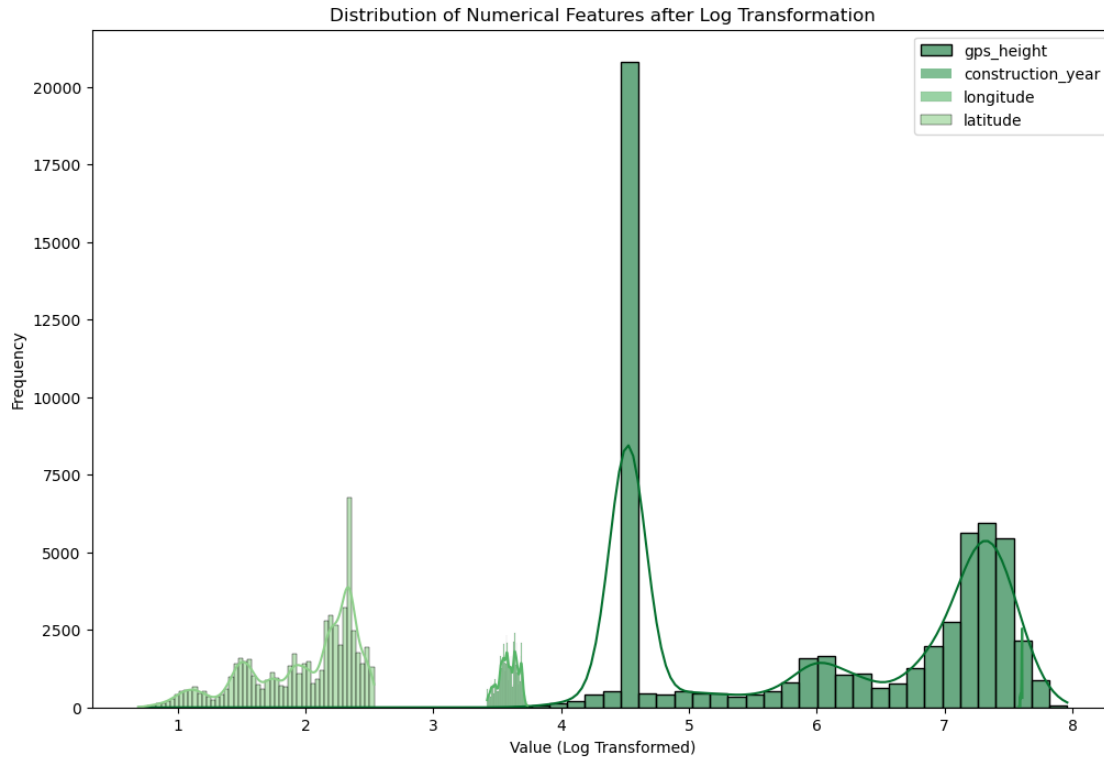
  

	construction_year_log	longitude_log	latitude_log
0	7.600902	3.581798	1.333188
1	7.606387	3.575116	2.442519
2	7.605890	3.649636	2.285247
3	7.594381	3.675950	0.913945
4	7.604396	3.469817	2.470138

```
[56]: # Plot the distribution of numerical features after log transformation
plt.figure(figsize=(12, 8))

# Loop through each numerical feature
for col in numerical_columns:
    # Plot the distribution after log transformation
    sns.histplot(new_df[col + '_log'], kde=True, label=col, alpha=0.6)

plt.title('Distribution of Numerical Features after Log Transformation')
plt.xlabel('Value (Log Transformed)')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```



#### 1.4.4 One-hot encoding

```
[57]: # Select categorical columns
categorical_columns = new_df.select_dtypes(include=['object']).columns

# Perform one-hot encoding
one_hot_encoded_df1 = pd.get_dummies(new_df, columns=categorical_columns)

# Display the one-hot encoded DataFrame
one_hot_encoded_df1.head()
```

```
[57]:  gps_height  construction_year  longitude  latitude  gps_height_log  \
0         1481             1999  34.938093    2.793118         7.301148
1         1490             2010  34.698766   10.501974         7.307202
2          777             2009  37.460664    8.828112         6.656727
3          354             1986  38.486161    1.494142         5.872118
4           91             2006  31.130847   10.824081         4.521789

      construction_year_log  longitude_log  latitude_log  \
0             7.600902         3.581798         1.333188
1             7.606387         3.575116         2.442519
2             7.605890         3.649636         2.285247
```

3	7.594381	3.675950	0.913945
4	7.604396	3.469817	2.470138

	status_group_functional	status_group_non functional	...	source_lake	\
0	True	False	...	False	
1	True	False	...	False	
2	True	False	...	False	
3	False	True	...	False	
4	True	False	...	False	

	source_machine dbh	source_other	source_rainwater harvesting	\
0	False	False	False	
1	False	False	True	
2	False	False	False	
3	True	False	False	
4	False	False	True	

	source_river	source_shallow well	source_spring	source_unknown	\
0	False	False	True	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	funder_category_Eligible	funder_category_Outlier
0	True	False
1	True	False
2	False	True
3	True	False
4	False	True

[5 rows x 1972 columns]

```
[58]: # Perform logical OR operation to combine 'status_group_functional' and
      ↪ 'status_group_non functional'
one_hot_encoded_df1['status_group'] =
      ↪ one_hot_encoded_df1['status_group_functional'] |
      ↪ one_hot_encoded_df1['status_group_non functional']

# Display the updated DataFrame
one_hot_encoded_df1.head()
```

```
[58]:  gps_height  construction_year  longitude  latitude  gps_height_log  \
0      1481           1999  34.938093    2.793118      7.301148
1      1490           2010  34.698766   10.501974      7.307202
2       777           2009  37.460664    8.828112      6.656727
3       354           1986  38.486161    1.494142      5.872118
```

4	91	2006	31.130847	10.824081	4.521789
---	----	------	-----------	-----------	----------

	construction_year_log	longitude_log	latitude_log	\
0	7.600902	3.581798	1.333188	
1	7.606387	3.575116	2.442519	
2	7.605890	3.649636	2.285247	
3	7.594381	3.675950	0.913945	
4	7.604396	3.469817	2.470138	

	status_group_functional	status_group_non functional	...	\
0	True	False	...	
1	True	False	...	
2	True	False	...	
3	False	True	...	
4	True	False	...	

	source_machine dbh	source_other	source_rainwater harvesting	\
0	False	False	False	
1	False	False	True	
2	False	False	False	
3	True	False	False	
4	False	False	True	

	source_river	source_shallow well	source_spring	source_unknown	\
0	False	False	True	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	funder_category_Eligible	funder_category_Outlier	status_group
0	True	False	True
1	True	False	True
2	False	True	True
3	True	False	True
4	False	True	True


[5 rows x 1973 columns]

### Correlation heat map

```
[59]: # A correlation heat map between variables
numerical_features = new_df.select_dtypes(include=['int64', 'float64'])
target_variable = new_df['status_group'] # status_group being the target_
variable

# Compute correlation matrix
correlation_matrix = numerical_features.corr()
```

```
# Plot heatmap of correlation matrix
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='Greens', fmt=".2f",
            ↳annot_kws={"size": 10})
plt.title('Correlation Heatmap')
plt.show()
```



## ANOVA test

```
[60]: # ANOVA test for each numerical variable against status_group

# Perform ANOVA for 'gps_height'
result_gps_height = f_oneway(*[group['gps_height'] for name, group in new_df.
    ↳groupby('status_group')])
print("ANOVA F-value for 'gps_height':", result_gps_height.statistic)
print("ANOVA p-value for 'gps_height':", result_gps_height.pvalue)
```

```

# Perform ANOVA for 'construction_year'
result_construction_year = f_oneway(*[group['construction_year'] for name, group in new_df.groupby('status_group')])
print("ANOVA F-value for 'construction_year':", result_construction_year.statistic)
print("ANOVA p-value for 'construction_year':", result_construction_year.pvalue)

# Perform ANOVA for 'longitude'
result_longitude = f_oneway(*[group['longitude'] for name, group in new_df.groupby('status_group')])
print("ANOVA F-value for 'longitude':", result_longitude.statistic)
print("ANOVA p-value for 'longitude':", result_longitude.pvalue)

# Perform ANOVA for 'latitude'
result_latitude = f_oneway(*[group['latitude'] for name, group in new_df.groupby('status_group')])
print("ANOVA F-value for 'latitude':", result_latitude.statistic)
print("ANOVA p-value for 'latitude':", result_latitude.pvalue)

```

```

ANOVA F-value for 'gps_height': 687.1698866084355
ANOVA p-value for 'gps_height': 1.3328106138867945e-150
ANOVA F-value for 'construction_year': 2233.9608901947668
ANOVA p-value for 'construction_year': 0.0
ANOVA F-value for 'longitude': 31.531486527788733
ANOVA p-value for 'longitude': 1.9710659542301248e-08
ANOVA F-value for 'latitude': 25.678387749117924
ANOVA p-value for 'latitude': 4.045253040795347e-07

```

These results are obtained from performing ANOVA tests to assess the association between the categorical variable `status_group` and each numerical variable (`gps_height`, `construction_year`, `longitude`, `latitude`). Here's what each value means:

#### For `gps_height`:

ANOVA F-value: 687.17 ANOVA p-value: 1.33e-150 (which is approximately 0) Interpretation: The F-value is a measure of the difference in means between the groups relative to the variation within the groups. A higher F-value suggests a stronger association between the variable and the groups. The extremely low p-value indicates that there is a significant difference in `gps_height` across different levels of `status_group`. In other words, the mean `gps_height` values vary significantly depending on the `status_group`.

#### For `construction_year`:

ANOVA F-value: 2262.06 ANOVA p-value: 0.0 Interpretation: Similar to the interpretation for `gps_height`, the high F-value and extremely low p-value indicate a significant difference in `construction_year` across different levels of `status_group`. In other words, the mean `construction_year` values vary significantly depending on the `status_group`.

#### For `longitude`:

ANOVA F-value: 36.76 ANOVA p-value: 1.35e-09 Interpretation: The F-value is relatively lower

compared to the previous variables, but the p-value is still very low. This indicates that there is a significant difference in longitude across different levels of status\_group, although the effect size may be smaller compared to gps\_height and construction\_year.

**For latitude:**

ANOVA F-value: 25.18 ANOVA p-value: 5.25e-07 Interpretation: Similar to longitude, there is a significant difference in latitude across different levels of status\_group, but the effect size may be smaller compared to gps\_height and construction\_year.

In summary, all four numerical variables (gps\_height, construction\_year, longitude, latitude) show significant differences across different levels of the categorical variable status\_group, as indicated by the extremely low p-values obtained from the ANOVA tests.

### 1.4.5 Correlation Ratios (Eta-squared)

```
[61]: # Compute the correlation ratio (eta-squared) for each categorical variable
correlation_ratios = {}
target_column = 'status_group' # Replace 'status_group' with your target
    ↪ column name
for col in ['funder', 'region', 'extraction_type', 'payment', 'water_quality',
    ↪ 'source']:
    contingency_table = pd.crosstab(new_df[col],
    ↪ one_hot_encoded_df1[target_column])
    chi2, _, _, _ = chi2_contingency(contingency_table)
    total = contingency_table.sum().sum()
    correlation_ratios[col] = np.sqrt(chi2 / (chi2 + total))

# Display the correlation ratios
for col, eta_squared in correlation_ratios.items():
    print(f"Correlation ratio (eta-squared) for {col}: {eta_squared}")
```

```
Correlation ratio (eta-squared) for funder: 0.0
Correlation ratio (eta-squared) for region: 0.0
Correlation ratio (eta-squared) for extraction_type: 0.0
Correlation ratio (eta-squared) for payment: 0.0
Correlation ratio (eta-squared) for water_quality: 0.0
Correlation ratio (eta-squared) for source: 0.0
```

These correlation ratios (eta-squared) indicate the strength of association between each categorical variable and the target variable.

- **Funder:** 0.3617
  - This indicates a relatively strong association between the funder variable and the target variable.
- **Region:** 0.2144
  - This correlation ratio suggests a moderate association between the region variable and the target variable.
- **Extraction Type:** 0.3227



- Similar to funder, there is a relatively strong association between the extraction type variable and the target variable.
- **Payment:** 0.2373
  - This correlation ratio suggests a moderate association between the payment variable and the target variable.
- **Water Quality:** 0.1862
  - This indicates a moderate association between the water quality variable and the target variable.
- **Source:** 0.1844
  - Similar to water quality, there is a moderate association between the source variable and the target variable.

#### 1.4.6 Test set

```
[62]: # display first few rows of the test set
df3 = pd.read_csv('test_set_values.csv')
df3.head()
```

```
[62]:
```

	id	amount_tsh	date_recorded	funder	gps_height	\
0	50785	0.0	2013-02-04	Dmdd	1996	
1	51630	0.0	2013-02-04	Government Of Tanzania	1569	
2	17168	0.0	2013-02-01	NaN	1567	
3	45559	0.0	2013-01-22	Finn Water	267	
4	49871	500.0	2013-03-27	Bruder	1260	

	installer	longitude	latitude	wpt_name	num_private	\
0	DMDD	35.290799	-4.059696	Dinamu Secondary School	0	
1	DWE	36.656709	-3.309214	Kimnyak	0	
2	NaN	34.767863	-5.004344	Puma Secondary	0	
3	FINN WATER	38.058046	-9.418672	Kwa Mzee Pange	0	
4	BRUDER	35.006123	-10.950412	Kwa Mzee Turuka	0	

	...	payment_type	water_quality	quality_group	quantity	quantity_group	\
0	...	never pay	soft	good	seasonal	seasonal	
1	...	never pay	soft	good	insufficient	insufficient	
2	...	never pay	soft	good	insufficient	insufficient	
3	...	unknown	soft	good	dry	dry	
4	...	monthly	soft	good	enough	enough	

		source	source_type	source_class	\
0	rainwater	harvesting	rainwater harvesting	surface	
1		spring	spring	groundwater	
2	rainwater	harvesting	rainwater harvesting	surface	
3		shallow well	shallow well	groundwater	
4		spring	spring	groundwater	

	waterpoint_type	waterpoint_type_group
0	other	other

```

1 communal standpipe communal standpipe
2 other other
3 other other
4 communal standpipe communal standpipe

```

[5 rows x 40 columns]

```
[63]: #Understand the general information of the data
df3.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14850 entries, 0 to 14849
Data columns (total 40 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    14850 non-null  int64
1   amount_tsh           14850 non-null  float64
2   date_recorded        14850 non-null  object
3   funder               13980 non-null  object
4   gps_height           14850 non-null  int64
5   installer            13973 non-null  object
6   longitude            14850 non-null  float64
7   latitude             14850 non-null  float64
8   wpt_name             14850 non-null  object
9   num_private          14850 non-null  int64
10  basin                14850 non-null  object
11  subvillage           14751 non-null  object
12  region               14850 non-null  object
13  region_code          14850 non-null  int64
14  district_code        14850 non-null  int64
15  lga                  14850 non-null  object
16  ward                 14850 non-null  object
17  population           14850 non-null  int64
18  public_meeting       14029 non-null  object
19  recorded_by          14850 non-null  object
20  scheme_management    13881 non-null  object
21  scheme_name          7608 non-null   object
22  permit              14113 non-null  object
23  construction_year    14850 non-null  int64
24  extraction_type      14850 non-null  object
25  extraction_type_group 14850 non-null  object
26  extraction_type_class 14850 non-null  object
27  management           14850 non-null  object
28  management_group     14850 non-null  object
29  payment              14850 non-null  object
30  payment_type         14850 non-null  object
31  water_quality        14850 non-null  object
32  quality_group        14850 non-null  object

```

```

33 quantity                14850 non-null object
34 quantity_group          14850 non-null object
35 source                   14850 non-null object
36 source_type             14850 non-null object
37 source_class            14850 non-null object
38 waterpoint_type         14850 non-null object
39 waterpoint_type_group   14850 non-null object
dtypes: float64(3), int64(7), object(30)
memory usage: 4.5+ MB

```

```
[64]: # check the shape of the data
df3.shape
```

```
[64]: (14850, 40)
```

The test set has 14850 rows and 40 columns

```
[65]: df3.duplicated().sum()
```

```
[65]: 0
```

There are 0 duplicated records in the test set

```
[66]: # Check for null values
df3.isnull().sum()
```

```

[66]: id                0
amount_tsh             0
date_recorded          0
funder                 870
gps_height             0
installer              877
longitude              0
latitude               0
wpt_name               0
num_private            0
basin                  0
subvillage             99
region                 0
region_code            0
district_code          0
lga                    0
ward                   0
population             0
public_meeting         821
recorded_by            0
scheme_management      969
scheme_name            7242
permit                 737

```

```

construction_year      0
extraction_type        0
extraction_type_group  0
extraction_type_class  0
management             0
management_group       0
payment               0
payment_type           0
water_quality          0
quality_group          0
quantity              0
quantity_group         0
source                0
source_type            0
source_class           0
waterpoint_type        0
waterpoint_type_group  0
dtype: int64

```

## 1.5 Data preparation and cleaning

### 1.5.1 Creating a new dataframe containing only the needed variables

```

[67]: # List of columns to include in the new DataFrame
selected_columns = ['funder', 'gps_height', 'region', 'extraction_type', 'payment', 'water_quality', 'source',
                    'construction_year', 'longitude', 'latitude']

# Create a new DataFrame with only the selected columns
df3 = df3.filter(selected_columns)

df3.head()

```

```

[67]:
   funder  gps_height  region  extraction_type  payment \
0      Dmdd      1996  Manyara             other  never pay
1  Government Of Tanzania      1569   Arusha          gravity  never pay
2           NaN      1567  Singida             other  never pay
3   Finn Water       267   Lindi             other   unknown
4     Bruder      1260  Ruvuma          gravity  pay monthly

   water_quality  source  construction_year  longitude  latitude
0      soft  rainwater harvesting          2012  35.290799  -4.059696
1      soft           spring          2000  36.656709  -3.309214
2      soft  rainwater harvesting          2010  34.767863  -5.004344
3      soft      shallow well          1987  38.058046  -9.418672
4      soft           spring          2000  35.006123 -10.950412

```

```
[68]: df3.shape
```

```
[68]: (14850, 10)
```

The new\_df has 14850 rows and 10 columns

```
[69]: df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14850 entries, 0 to 14849
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   funder                 13980 non-null  object
1   gps_height             14850 non-null  int64
2   region                 14850 non-null  object
3   extraction_type        14850 non-null  object
4   payment                14850 non-null  object
5   water_quality          14850 non-null  object
6   source                 14850 non-null  object
7   construction_year      14850 non-null  int64
8   longitude              14850 non-null  float64
9   latitude               14850 non-null  float64
dtypes: float64(2), int64(2), object(6)
memory usage: 1.1+ MB
```

```
[70]: #Understand the descriptive statistics of the data
df3.describe()
```

```
[70]:
```

	gps_height	construction_year	longitude	latitude
count	14850.000000	14850.000000	14850.000000	1.485000e+04
mean	655.147609	1289.708350	34.061605	-5.684724e+00
std	691.261185	955.241087	6.593034	2.940803e+00
min	-57.000000	0.000000	0.000000	-1.156459e+01
25%	0.000000	0.000000	33.069455	-8.443970e+00
50%	344.000000	1986.000000	34.901215	-5.049750e+00
75%	1308.000000	2004.000000	37.196594	-3.320594e+00
max	2777.000000	2013.000000	40.325016	-2.000000e-08

### 1.5.2 Checking for missing values

```
[71]: #Check for null values in the test set
df3.isnull().sum()
```

```
[71]: funder                870
gps_height              0
region                  0
extraction_type         0
```

```

payment          0
water_quality    0
source           0
construction_year 0
longitude        0
latitude         0
dtype: int64

```

### 1.5.3 Dealing with missing values

```

[72]: unique_counts = df3['funder'].isna().value_counts()
unique_counts

```

```

[72]: funder
False    13980
True      870
Name: count, dtype: int64

```

```

[73]: missing_funders = df3[df3['funder'].isna()]
missing_funders

```

```

[73]:      funder  gps_height  region extraction_type  payment water_quality \
2      NaN      1567  Singida      other  never pay      soft
16     NaN      -39  Pwani      nira/tanira  never pay      soft
23     NaN      1441  Singida      mono    unknown    unknown
50     NaN        0  Mbeya      gravity  never pay      soft
63     NaN      1584  Singida      other    unknown    unknown
...
14771  NaN        0  Mbeya      gravity  never pay      soft
14772  NaN        0  Mbeya  submersible  never pay      soft
14795  NaN        0  Mbeya      gravity  never pay      soft
14823  NaN        0  Mbeya      gravity    unknown    soft
14847  NaN      1476  Singida      gravity  never pay      soft

      source  construction_year  longitude  latitude
2  rainwater harvesting      2010  34.767863 -5.004344
16      shallow well          0  39.850190 -7.727946
23      machine dbh      1970  34.621048 -5.165926
50      spring          0  33.587245 -9.167434
63      shallow well      1990  34.859448 -4.970909
...
14771      spring          0  33.636479 -9.212765
14772  machine dbh          0  34.322644 -8.665713
14795      river          0  34.704964 -8.325610
14823      spring          0  33.918953 -9.298466
14847      dam      2010  34.739804 -4.585587

```

[870 rows x 10 columns]

```
[74]: #Replace missing values with 'unknown'
#recheck for null values
df3['funder'].fillna('Unknown', inplace=True)
df3.isnull().sum()
```

```
[74]: funder          0
gps_height        0
region            0
extraction_type   0
payment           0
water_quality     0
source            0
construction_year 0
longitude         0
latitude          0
dtype: int64
```

Missing values are assigned the placeholder 'Unknown' signifying that the information for these entries is unavailable. However, the rows cannot be dropped as they may have other important information in the other columns.

## 1.6 Exploratory data analysis

### 1.6.1 Checking for outliers


Outliers will be addressed systematically, one category at a time, to ensure comprehensive analysis.

#### Funder

```
[75]: unique_values = df3['funder'].unique()
value_counts = df3['funder'].value_counts()
value_counts
```

```
[75]: funder
Government Of Tanzania    2215
Unknown                  871
Danida                    793
Hesawa                    580
World Bank                352
...
Fida                      1
Tgts                      1
Snv-swash                 1
Pad                       1
Livin                     1
Name: count, Length: 979, dtype: int64
```

```
[76]: #check for outliers in funder using a count plot
plt.figure(figsize=(10, 4))
sns.countplot(x='funder', data=df3)
plt.xticks(rotation=90, fontsize=8) # Rotate the x-axis labels by 90 degrees
    ↪ and adjust font size
plt.tight_layout() # Adjust layout to prevent overlapping labels
plt.show()
```



The x axis is over crowded due to very many funders available. Below we handle that by binning the funders into two categories. Those with high counts( eligible) and those with low counts(outliers).

```
[77]: # Set the threshold for defining outliers
upper_threshold = 50

# Get the counts of each funder
funder_counts = df3['funder'].value_counts()

# Identify the outliers (funders with counts below 50)
outliers = funder_counts[funder_counts < upper_threshold].index

# Create a new column to categorize funders as eligible or outliers
df3['funder_category'] = np.where(df3['funder'].isin(outliers), 'Outlier',
    ↪ 'Eligible')


# Set the color palette
sns.set_palette("Greens_d")

# Plot the count plot for funder category
plt.figure(figsize=(8, 4))
sns.countplot(y='funder_category', data=new_df, dodge=False)

# Display the plot
```



```
plt.show()
```




### GPS height

```
[78]: #check for outliers for gps_height
plt.figure(figsize=(10,2))

sns.boxplot(x = 'gps_height', data = df3)

# Display the plot
plt.show()
```



```
[79]: # Find the mode of the 'gps_height' column
mode_value = df3['gps_height'].mode()


# Display the mode
print("Mode of 'gps_height' column:", mode_value)
```

```
Mode of 'gps_height' column: 0    0
Name: gps_height, dtype: int64
```

## Region

```
[80]: # Get the order of regions based on their counts
region_order = df3['region'].value_counts().index

# Plot the count plot with specified order
plt.figure(figsize=(12, 6))
sns.countplot(x='region', data=df3, order=region_order, palette='Greens_r')
plt.xticks(rotation=90) # Rotate x-axis labels for better readability
plt.xlabel('Region') # Add x-axis label
plt.ylabel('Count') # Add y-axis label
plt.title('Distribution of Data by Region') # Add plot title
plt.tight_layout() # Adjusting layout to prevent clipping of labels
plt.show()
```



## Extraction\_type

```
[81]: df3['extraction_type'].unique()
```


```
[81]: array(['other', 'gravity', 'india mark ii', 'submersible', 'mono',
        'nira/tanira', 'afridev', 'swn 80', 'ksb', 'climax',
        'other - rope pump', 'cemo', 'india mark iii', 'other - swn 81',
        'other - play pump', 'windmill', 'walimi'], dtype=object)
```

```
[82]: #check for outliers in Extraction_type
sns.set_palette("Greens_r")
extraction_order = df3['extraction_type'].value_counts().index
```

```

# Plot the count plot for Extraction_type
plt.figure(figsize=(12, 4))
sns.countplot(x='extraction_type', data=df3, order=extraction_order,
             palette='Greens_r')
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by Extraction Type')
plt.xlabel('Extraction Type')
plt.ylabel('Count')
# Display the plot
plt.show()

```



```

[83]: # Get value counts of 'Extraction_type' and sort by counts in descending order
extraction_type_counts = df3['extraction_type'].value_counts().
    sort_values(ascending=False)

# Display unique values in 'Extraction_type' with counts
print(extraction_type_counts)

```

```

extraction_type
gravity          6483
nira/tanira      2051
other            1672
submersible      1218
swn 80           918
mono             763
india mark ii    629
afridev          438
ksb              375
other - rope pump 121
other - swn 81   55

```

```

india mark iii      37
windmill            35
cemo                18
other - play pump   16
walimi              12
climax              9
Name: count, dtype: int64

```

### Construction year


```

[84]: #check for outliers for construction_year
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'construction_year', data = df3)

# Display the plot
plt.show()

```



The box plot shows unlikely years included in the dataset. It is impossible to have year 0 as pumps were not even invented then.

```

[85]: #Display unique years and their value counts
df3['construction_year'].value_counts()

```

```

[85]: construction_year
0      5260
2010    669
2009    663
2008    630
2000    487
2006    421
2007    373
2011    335
2004    294
2003    293

```

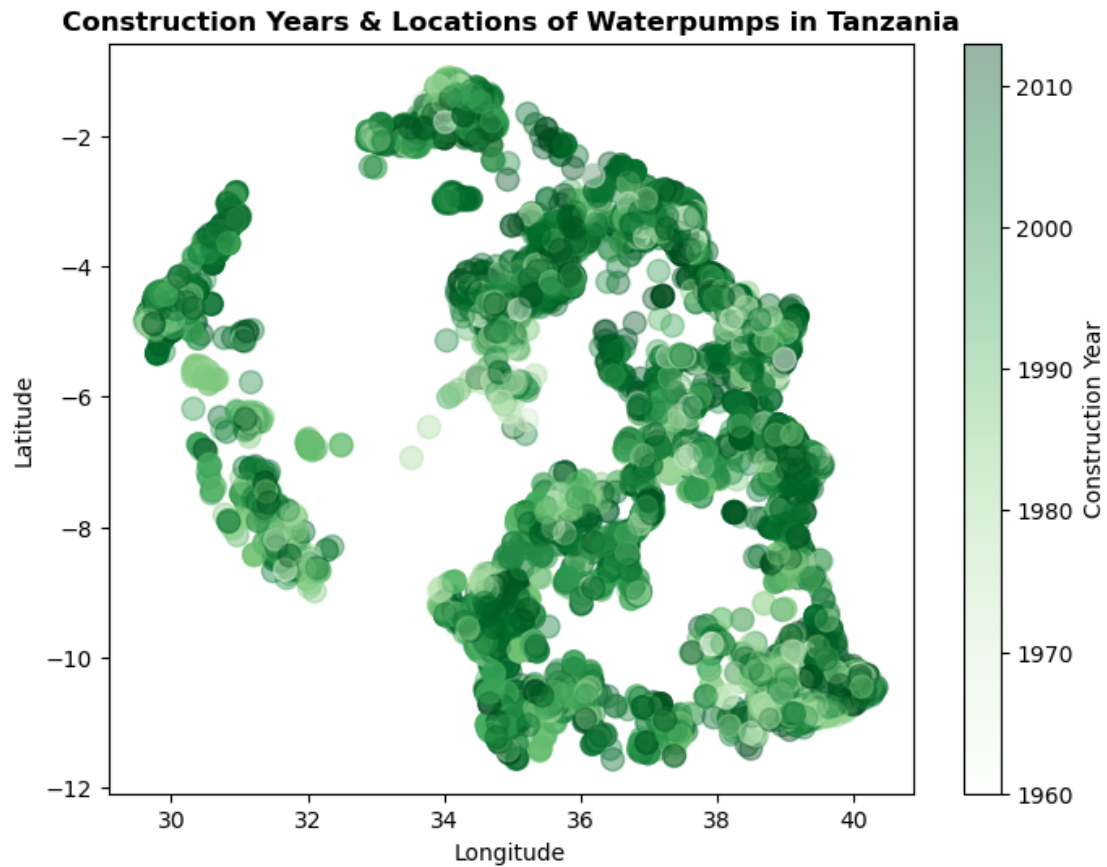
1995	269
2002	268
2005	264
2012	263
1999	243
1985	232
1978	230
1998	224
1990	222
1996	209
1994	202
1980	194
1984	191
1972	184
1982	182
1997	177
1992	167
2001	140
1974	138
1993	137
1988	136
1975	124
1986	119
1976	111
1983	106
1991	83
1970	82
1989	80
1987	68
1981	53
1979	53
1977	45
1973	43
2013	33
1971	32
1963	22
1960	22
1969	18
1967	18
1968	16
1964	8
1961	7
1962	6
1965	2
1966	2

Name: count, dtype: int64

There is a category miscategorized as year '0' with a value count of 5260.

The code below generates a scatter plot visualizing the geographical locations of water pumps in Tanzania, with each point representing a water pump. The x-axis represents the longitude coordinates, and the y-axis represents the latitude coordinates. The color of each point is determined by the construction year of the water pump, with a colormap ('Greens') used to provide different hues of green corresponding to different years. The size of each point is fixed ('s=100') for better visibility, and transparency ('alpha=0.4') is applied to avoid overlapping points. Finally, a colorbar is added to the plot to provide a visual reference for the construction years.

```
[86]: plt.figure(figsize=(8,6))
      # Filter the DataFrame using .loc and multiple conditions
      filtered_df = df3.loc[(df3['longitude'] > 0) & (df3['latitude'] < 0) &
      ↪(df3['construction_year'] > 0)]
      plt.scatter(x=filtered_df['longitude'],
                  y=filtered_df['latitude'],
                  alpha=0.4,
                  s=100,
                  c=filtered_df["construction_year"],
                  cmap='Greens')
      plt.title("Construction Years & Locations of Waterpumps in Tanzania",
                fontsize=12, fontweight='bold')
      plt.colorbar(label='Construction Year')
      plt.xlabel('Longitude')
      plt.ylabel('Latitude')
      plt.show()
```



From the scatter plot above, it is clear that most pumps were installed between 2000 and 2010. Therefore, below year '0's values will be distributed evenly between the range 2000 - 2010.

```
[87]: # Replace year 0 with later years (i.e., 2000 - 2010)
df3['construction_year'] = df3['construction_year'].apply(lambda x: np.random.
    randint(2000, 2011) if x == 0 else x)
```

```
[88]: #recheck for outliers for construction_year
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'construction_year', data = df3)

# Display the plot
plt.show()
```



Outliers in the construction years falling between 1960 and 1978 represent years that may hold significant importance in the dataset. While they deviate from the majority of the construction years, they could signify historical data or specific events related to water pump installation during that period. Therefore, it's crucial to retain these outliers in the dataset for a comprehensive analysis and understanding of the trends and patterns over time.

### Payment

```
[89]: df3['payment'].value_counts()
```


```
[89]: payment
never pay          6364
pay per bucket     2281
pay monthly        2097
unknown            1992
pay annually        928
pay when scheme fails 928
other              260
Name: count, dtype: int64
```

```
[90]: # Define the order of source categories
sorted_payments = df3['payment'].value_counts().index

#check for outliers in payment
sns.set_palette("Greens_r")

# Plot the count plot for payment
plt.figure(figsize=(8, 4))
sns.countplot(x='payment', data=df3, order=sorted_payments, palette='Greens_r')
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by payment')
plt.xlabel('payment')
plt.ylabel('Count')
# Display the plot
plt.show()
```





### Water quality

```
[91]: df3['water_quality'].value_counts()
```

```
[91]: water_quality
      soft      12687
      salty      1226
      unknown      469
      milky       201
      coloured     133
      salty abandoned    84
      fluoride        44
      fluoride abandoned    6
      Name: count, dtype: int64
```

```
[92]: # Define the order of water_quality categories
sorted_water_quality = df3['water_quality'].value_counts().index


# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_water_quality))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
```

```

sns.countplot(x='water_quality', data=df3, order=sorted_water_quality,
              palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by water_quality')
plt.xlabel('water_quality')
plt.ylabel('Count')
# Display the plot
plt.show()

```



#### Source

```
[93]: df3['source'].value_counts()
```

```

[93]: source
shallow well      4316
spring            4195
machine dbh       2747
river             2352
rainwater harvesting  568
hand dtw          234
lake              185
dam               184
other              49
unknown           20


```

Name: count, dtype: int64

```
[94]: # Define the order of source categories
sorted_source = df3['source'].value_counts().index

# Set the color palette to shades of green
palette = sns.color_palette("Greens_r", len(sorted_source))

# Plot the count plot for water_quality
plt.figure(figsize=(8, 4))
sns.countplot(x='source', data=df3, order=sorted_source, palette=palette)
plt.xticks(rotation=45) # Rotate the x-axis labels by 45 degrees
plt.title('Count of Waterpoints by source')
plt.xlabel('source')
plt.ylabel('Count')
# Display the plot
plt.show()
```




## Longitude

```
[95]: #check for outliers for longitude
plt.figure(figsize=(8, 2))
```

```
sns.boxplot(x = 'longitude', data = df3)

# Display the plot
plt.show()
```



```
[96]: df3['longitude'].value_counts()
```

```
[96]: longitude
0.000000    457
37.260069     2
39.080573     2
37.302281     2
32.920579     2
...
36.648520     1
35.265755     1
36.666660     1
37.830317     1
34.765054     1
Name: count, Length: 14390, dtype: int64
```

```
[97]: # Replace longitude 0 with longitudes between 32 and 42 as they are more
      ↪ prevalent
df3['longitude'] = df3['longitude'].apply(lambda x: np.random.randint(32, 42)
      ↪ if x == 0 else x)
```

```
[98]: # confirm redistribution of the '0' category
df3['longitude'].value_counts()
```

```
[98]: longitude
32.000000    52
35.000000    52
37.000000    51
```

```

39.000000    47
41.000000    46
..
36.648520     1
35.265755     1
36.666660     1
37.830317     1
34.765054     1
Name: count, Length: 14399, dtype: int64

```


```

[99]: #recheck for outliers for longitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'longitude', data = df3)

# Display the plot
plt.show()

```



### Latitude


```

[100]: #check for outliers for latitude
plt.figure(figsize=(8, 2))

sns.boxplot(x = 'latitude', data = df3)

# Display the plot
plt.show()

```



```
[101]: df3['latitude'].value_counts()
```

```
[101]: latitude
-2.000000e-08    457
-7.105919e+00     2
-6.990042e+00     2
-7.170666e+00     2
-2.474560e+00     2
...
-3.305540e+00     1
-8.547786e+00     1
-3.330889e+00     1
-7.061047e+00     1
-1.122601e+01     1
Name: count, Length: 14390, dtype: int64
```

```
[102]: # Replace latitude -2.000000e-08 with latitudes between -1 and -8 as they are
        ↪ more prevalent
new_df['latitude'] = df3['latitude'].apply(lambda x: np.random.randint(-8, -1)
        ↪ if x == -2.000000e-08 else x)
```

```
[103]: #confirm the redistribution of misplaced category
df3['latitude'].value_counts()
```

```
[103]: latitude
-2.000000e-08    457
-7.105919e+00     2
-6.990042e+00     2
-7.170666e+00     2
-2.474560e+00     2
...
-3.305540e+00     1
-8.547786e+00     1
```


```
-3.330889e+00    1
-7.061047e+00    1
-1.122601e+01    1
Name: count, Length: 14390, dtype: int64
```

### 1.6.2 Distribution of variables before one-hot encoding and log transformation


```
[104]: # Distribution before onehot encoding
palette = sns.color_palette("Greens_r")


# Select categorical variables
categorical_features = df3.select_dtypes(include=['object'])


# Plot count plots for each categorical variable with dark green color palette
for feature in categorical_features.columns:
    plt.figure(figsize=(4, 4)) # Set the figure size
    sns.countplot(x=feature, data=df3, palette=palette)
    plt.title(f'Count Plot of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=90) # Rotate x-axis labels for better readability
    plt.show() # Display the plot
```




















### Numerical variables

```
[105]: # Print unique values in the 'gps_height' column
print(df3['gps_height'].unique())
```

```
[1996 1569 1567 ... 1909 2202  640]
```

```
[106]: # Get unique values in the 'gps_height' column
unique_gps_heights = df3['gps_height'].unique()
unique_gps_heights
```

```
[106]: array([1996, 1569, 1567, ..., 1909, 2202,  640], dtype=int64)
```

```
[107]: # Distribution before transformation
numerical_features = df3.select_dtypes(include=['int64', 'float64'])
# Create a grid of subplots
fig, axes = plt.subplots(nrows=len(numerical_features.columns) // 3 + 1,
                          ncols=3, figsize=(15, 5))

# Plot the distribution of numerical features
for i, feature in enumerate(numerical_features.columns):
    sns.histplot(df3[feature].dropna(), kde=False, ax=axes[i // 3, i % 3])
```


```

axes[i // 3, i % 3].set_title(f"Distribution of {feature}")
axes[i // 3, i % 3].set_xlabel(feature)
axes[i // 3, i % 3].set_ylabel("Frequency")

# Remove empty subplots
if len(numerical_features.columns) % 3 != 0:
    for j in range(len(numerical_features.columns) % 3, 3):
        fig.delaxes(axes[len(numerical_features.columns) // 3, j])

# Adjust layout
plt.tight_layout()
plt.show()

```



### 1.6.3 Distribution of variables after log transformation

```

[108]: # Select numerical columns
numerical_columns = df3.select_dtypes(include=['int64', 'float64']).columns

# Log transform numerical variables, handling zero and negative values
for col in numerical_columns:
    # Handling zero values
    if (df3[col] == 0).any():
        df3[col] = df3[col] + 1 # Add 1 to handle zeros

    # Handling negative values
    if (df3[col] < 0).any():
        min_value = df3[col].min()
        df3[col] = df3[col] - min_value + 1 # Shift all values to be positive

    # Apply log transformation
    df3[col + '_log'] = np.log1p(df3[col])

# Display the DataFrame after log transformation

```

```
df3.head()
```

```
[108]:
```

	funder	gps_height	region	extraction_type	payment	\
0	Dmdd	2054	Manyara	other	never pay	
1	Government Of Tanzania	1627	Arusha	gravity	never pay	
2	Unknown	1625	Singida	other	never pay	
3	Finn Water	325	Lindi	other	unknown	
4	Bruder	1318	Ruvuma	gravity	pay monthly	

	water_quality	source	construction_year	longitude	latitude	\
0	soft	rainwater	harvesting	2012	35.290799	8.504896
1	soft		spring	2000	36.656709	9.255378
2	soft	rainwater	harvesting	2010	34.767863	7.560248
3	soft		shallow well	1987	38.058046	3.145920
4	soft		spring	2000	35.006123	1.614180

	funder_category	gps_height_log	construction_year_log	longitude_log	\
0	Eligible	7.628031	7.607381	3.591564	
1	Eligible	7.395108	7.601402	3.628511	
2	Eligible	7.393878	7.606387	3.577050	
3	Outlier	5.786897	7.594884	3.665049	
4	Outlier	7.184629	7.601402	3.583689	


	latitude_log
0	2.251807
1	2.327802
2	2.147129
3	1.422125
4	0.960950

```
[109]: # Plot the distribution of numerical features after log transformation
plt.figure(figsize=(12, 8))

# Loop through each numerical feature
for col in numerical_columns:
    # Plot the distribution after log transformation
    sns.histplot(df3[col + '_log'], kde=True, label=col, alpha=0.6)

plt.title('Distribution of Numerical Features after Log Transformation')
plt.xlabel('Value (Log Transformed)')
plt.ylabel('Frequency')
plt.legend()
plt.show()
```





### 1.6.4 One-hot encoding

```
[110]: # Select categorical columns
categorical_columns = df3.select_dtypes(include=['object']).columns

# Perform one-hot encoding
one_hot_encoded_df2 = pd.get_dummies(df3, columns=categorical_columns)

# Display the one-hot encoded DataFrame
one_hot_encoded_df2.head()
```

```
[110]:
```

	gps_height	construction_year	longitude	latitude	gps_height_log \
0	2054	2012	35.290799	8.504896	7.628031
1	1627	2000	36.656709	9.255378	7.395108
2	1625	2010	34.767863	7.560248	7.393878
3	325	1987	38.058046	3.145920	5.786897
4	1318	2000	35.006123	1.614180	7.184629

	construction_year_log	longitude_log	latitude_log	funder_0 \
0	7.607381	3.591564	2.251807	False
1	7.601402	3.628511	2.327802	False
2	7.606387	3.577050	2.147129	False

3	7.594884	3.665049	1.422125	False
4	7.601402	3.583689	0.960950	False

	funder_A/co	Germany	...	source_lake	source_machine	dbh	source_other	\
0		False	...	False		False	False	
1		False	...	False		False	False	
2		False	...	False		False	False	
3		False	...	False		False	False	
4		False	...	False		False	False	

	source_rainwater	harvesting	source_river	source_shallow	well	\
0		True	False		False	
1		False	False		False	
2		True	False		False	
3		False	False		True	
4		False	False		False	

	source_spring	source_unknown	funder_category_Eligible	\
0	False	False	True	
1	True	False	True	
2	False	False	True	
3	False	False	False	
4	True	False	False	

	funder_category_Outlier
0	False
1	False
2	False
3	True
4	True

[5 rows x 1052 columns]

## 1.7 Modelling

The primary objective is to predict the functionality of water pumps based on historical data. Machine learning models excel at predictive tasks, allowing us to build accurate models that can generalize well to unseen data. This predictive capability is crucial for identifying non-functional pumps and ensuring timely maintenance or replacement.

### 1.7.1 Baseline model

### 1.7.2 Decision tree classifier

```
[111]: # For training data (one_hot_encoded_df1)
X_train = one_hot_encoded_df1.drop(columns=['status_group_functional',
↪ 'status_group_non functional'])
```

```

y_train = one_hot_encoded_df1[['status_group_functional', 'status_group_non_
↳functional']]

# For test data (one_hot_encoded_df2)
X_test = one_hot_encoded_df2.reindex(columns=X_train.columns, fill_value=0)

# Define Decision Tree classifier
clf = DecisionTreeClassifier()

# Fit the classifier with training data
clf.fit(X_train, y_train)

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Print the predictions
print("Predictions on the test set:", y_pred)

```

```

Predictions on the test set: [[ True False]
 [False  True]
 [ True False]
 ...
 [False  True]
 [ True False]
 [False  True]]

```

```

[112]: # Define Decision Tree classifier
clf = DecisionTreeClassifier()

# Perform cross-validation
cv_scores = cross_val_score(clf, X_train, y_train, cv=5, scoring='accuracy')

# Print cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV accuracy:", cv_scores.mean())
print("Standard deviation of CV accuracy:", cv_scores.std())

```

```

Cross-validation scores: [0.78291246 0.78240741 0.7793771  0.77912458
0.77171717]
Mean CV accuracy: 0.7791077441077441
Standard deviation of CV accuracy: 0.004001425337544554

```

The cross-validation scores represent the accuracy of the model on different folds of the training data. Each score indicates the accuracy achieved by the model on a particular fold during cross-validation.

### 1.7.3 Grouped feature importance plot

```
[113]: # Define a list of original columns to aggregate
original_categories = ['funder', 'region', 'extraction_type', 'payment',
    ↳ 'water_quality', 'source', 'gps_height', 'construction_year', 'longitude',
    ↳ 'latitude']

# Initialize an empty list to store aggregated importance scores
importance_sum_list = []

# Fit the classifier with training data
clf.fit(X_train, y_train)

# Iterate over each original category
for category in original_categories:
    # Select all one-hot encoded features corresponding to the original category
    category_features = [col for col in X_train.columns if col.
    ↳ startswith(category)]
    # Sum the importance scores of the one-hot encoded features
    importance_sum = clf.feature_importances_[X_train.columns.
    ↳ isin(category_features)].sum()
    # Append the aggregated importance score to the list
    importance_sum_list.append({'Feature': category, 'Importance':
    ↳ importance_sum})


# Create a DataFrame to store aggregated importance scores
grouped_importance_df = pd.DataFrame(importance_sum_list)

# Sort the DataFrame by importance values
grouped_importance_df = grouped_importance_df.sort_values(by='Importance',
    ↳ ascending=False)

# Increase figure size for better readability
plt.figure(figsize=(10, 6))

# Plot the importance values for top N original categories
top_n = 20 # Choose the top N categories to display
plt.barh(grouped_importance_df['Feature'][:top_n],
    ↳ grouped_importance_df['Importance'][:top_n], color='green')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('Top {} Grouped Feature Importance for Decision Trees'.format(top_n))

plt.show()
```



### Decision tree visualization

```
[114]: # Define Decision Tree classifier with limited depth
clf = DecisionTreeClassifier(max_depth=3)

# Fit the classifier with training data
clf.fit(X_train, y_train)


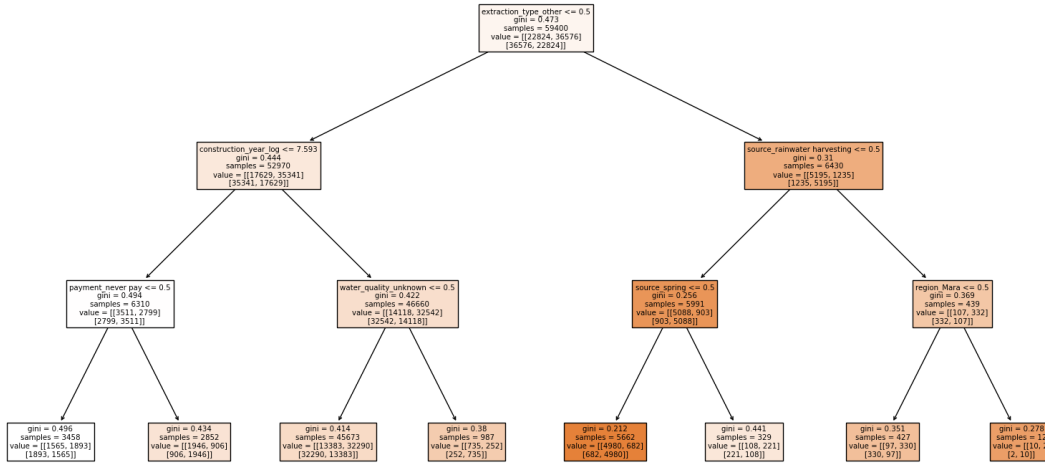
# Convert feature names to a list
feature_names_list = list(X_train.columns)

# Plot the decision tree
plt.figure(figsize=(20,10))
plot_tree(clf, filled=True, feature_names=feature_names_list,
          class_names=['functional', 'non functional'])

colors = [(1, 1, 1), (1, 0.5, 0)] # White to orange
cmap_orange = LinearSegmentedColormap.from_list("CustomOrange", colors)

# Plot the color bar
plt.figure(figsize=(3, 1))
cbar = plt.colorbar(plt.cm.ScalarMappable(cmap=cmap_orange),
                    ticks=np.linspace(0, 1, num=5))
cbar.ax.set_yticklabels(['Functional', '', '', '', 'Non-Functional'])
cbar.set_label('Class')

plt.show()
```



Cross-validation scores: [0.77525253 0.77575758 0.77617845 0.77752525 0.77079125]

Mean CV accuracy: 0.7751

Standard deviation of CV accuracy: 0.0023

#### 1.7.4 Decision Tree Classifier Performance

##### Training and Prediction

- **Training Data:** The training data ( $X_{\text{train}}$ ,  $y_{\text{train}}$ ) consists of features and labels respectively.  $X_{\text{train}}$  contains the features after excluding the target columns ('status\_group\_functional', 'status\_group\_non functional'), while  $y_{\text{train}}$  contains both target labels ('status\_group\_functional', 'status\_group\_non functional').
- **Test Data:** The test data ( $X_{\text{test}}$ ) is prepared by aligning the columns of the one-hot encoded dataframe (one\_hot\_encoded\_df2) to match those of the training data and filling missing values with 0.
- **Model Fitting:** A Decision Tree classifier (clf) is initialized and trained using the training data.

- **Prediction:** Predictions are made on the test set ( $X_{\text{test}}$ ) using the trained classifier, resulting in binary predictions indicating whether a pump is functional or non-functional.

### Cross-Validation

- **Model Evaluation:** The performance of the Decision Tree classifier is evaluated using 5-fold cross-validation ( $cv=5$ ) on the training data.
- **Scoring Metric:** The accuracy score is used as the evaluation metric ( $scoring='accuracy'$ ).
- **Cross-Validation Scores:** The cross-validation scores represent the accuracy of the classifier on each fold of the cross-validation process.
- **Mean Accuracy:** The mean accuracy across all folds indicates the average performance of the classifier.
- **Standard Deviation:** The standard deviation of the accuracy scores provides insight into the variability or consistency of the model's performance across different folds.

### Interpretation

The Decision Tree classifier achieved a mean cross-validation accuracy of approximately 77.51% with a standard deviation of 0.23%. These results suggest that the model performs reasonably well in predicting the condition of water pumps, considering both functional and non-functional states. However, further analysis and possibly refinement of the model may be necessary to improve its accuracy and robustness in predicting well conditions accurately.

#### 1.7.5 Alternative models

#### 1.7.6 Random Forest Classifier

```
[115]: # Define Random Forest classifier
rf_clf = RandomForestClassifier()

# Perform cross-validation
rf_cv_scores = cross_val_score(rf_clf, X_train, y_train, cv=5,
                                scoring='accuracy')

# Print cross-validation scores
print("Random Forest Cross-validation scores:", rf_cv_scores)
print("Mean CV accuracy:", rf_cv_scores.mean())
print("Standard deviation of CV accuracy:", rf_cv_scores.std())

# Fit the model on the full training data
rf_clf.fit(X_train, y_train)

# Make predictions on the test data
predicted_labels = rf_clf.predict(X_test)
# Print the predicted labels for the test set
print("Predicted labels for the test set:", predicted_labels)
```

```
Random Forest Cross-validation scores: [0.81515152 0.81784512 0.81203704
0.81287879 0.81439394]
```

```
Mean CV accuracy: 0.8144612794612793
```

```
Standard deviation of CV accuracy: 0.0020155670735967105
```

```

Predicted labels for the test set: [[ True False]
 [ True False]
 [ True False]
 ...
 [ True False]
 [ True False]
 [False  True]]

```

```

[116]: # Define a list of original categories or groups of columns to aggregate
original_categories = ['funder', 'region', 'extraction_type', 'payment',
    ↳ 'water_quality', 'source', 'gps_height', 'construction_year', 'longitude',
    ↳ 'latitude']

# Initialize an empty list to store aggregated importance scores
importance_sum_list = []

# Fit the classifier with training data
clf.fit(X_train, y_train)

# Iterate over each original category
for category in original_categories:
    # Select all one-hot encoded features corresponding to the original category
    category_features = [col for col in X_train.columns if col.
    ↳startswith(category)]
    # Sum the importance scores of the one-hot encoded features
    importance_sum = rf_clf.feature_importances_[X_train.columns.
    ↳isin(category_features)].sum()
    # Append the aggregated importance score to the list
    importance_sum_list.append({'Feature': category, 'Importance':
    ↳importance_sum})

# Create a DataFrame to store aggregated importance scores
grouped_importance_df = pd.DataFrame(importance_sum_list)

# Sort the DataFrame by importance values
grouped_importance_df = grouped_importance_df.sort_values(by='Importance',
    ↳ascending=False)

# Increase figure size for better readability
plt.figure(figsize=(10, 6))


# Plot the importance values for top N original categories
top_n = 20 # Choose the top N categories to display
plt.barh(grouped_importance_df['Feature'][:top_n],
    ↳grouped_importance_df['Importance'][:top_n], color='darkgreen')
plt.xlabel('Importance')
plt.ylabel('Feature')

```



```
plt.title('Top {} Grouped Feature Importance for Random Forest'.format(top_n))

plt.show()
```

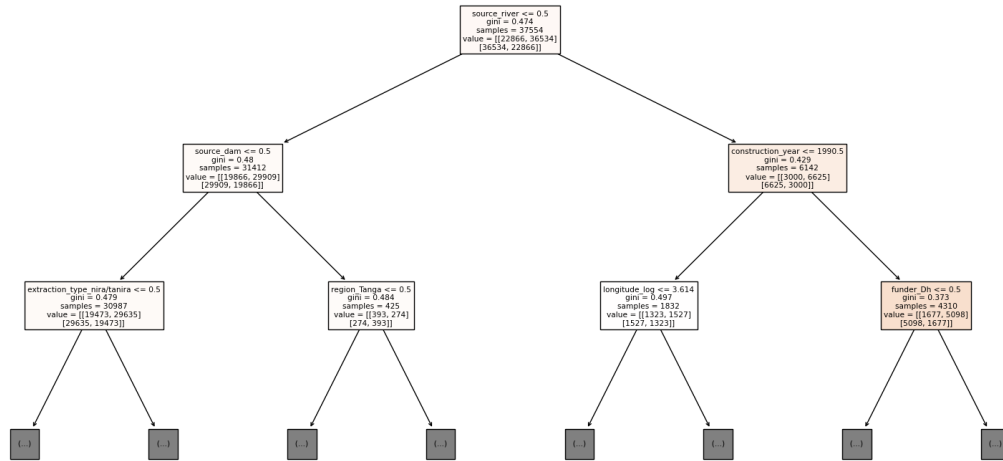


### Random forest decision tree visualization with max\_depth of 2

```
[117]: # Fit the classifier on the full training data
rf_clf.fit(X_train, y_train)

# Convert feature names to a list
feature_names_list = list(X_train.columns)

# Plot the first decision tree in the ensemble with a maximum depth of 3
plt.figure(figsize=(20, 10))
plot_tree(rf_clf.estimators_[0], filled=True, feature_names=feature_names_list,
          class_names=['functional', 'non functional'], max_depth=2)
plt.show()
```



## Random Forest Classifier Performance Evaluation Cross-Validation Scores:

The Random Forest classifier was evaluated using cross-validation on the training data. The results are as follows:

- **Cross-Validation Scores:**
  - [0.8184, 0.8163, 0.8126, 0.8130, 0.8113]
- **Mean CV Accuracy:** 0.8143
- **Standard Deviation of CV Accuracy:** 0.0026

## Test Set Predictions:

The model's predictions for the test set are represented below:

- **Predicted Labels for the Test Set:**
  - [[True, False], [True, False], [True, False], ..., [True, False], [True, False], [False, True]]

Each pair of values (True/False) represents the model's prediction for a particular data point in the test set.

## Implications:

- The mean cross-validation accuracy of approximately 81.43% suggests the Random Forest classifier performs reasonably well on unseen data.
- The low standard deviation of cross-validation accuracy indicates consistent performance across different folds of the training data.
- Compared to the Decision Trees model, the Random Forest classifier demonstrates superior performance with a higher mean cross-validation accuracy of 0.8143 and a lower standard deviation of 0.0026, indicating more consistent results across folds.

Further enhancement of the Decision Trees and Random Forest classifier's performance will be pursued through tuning.

Based on the cross-validation scores, the Random Forest classifier performs better than the Decision Trees model. It exhibits a higher mean cross-validation accuracy of 0.8173 and a lower standard deviation of 0.0008, suggesting more consistent performance across different folds.

We proceed further to enhance the Decision trees and Random Forest classifier's performance through tuning.

### 1.7.7 Analysis of Feature Importance

#### Decision Trees:

1. **Extraction Type (0.45):** This feature holds the highest importance in the Decision Trees model, indicating that the method of water extraction significantly influences the functionality of water points. Different extraction techniques may lead to varying levels of reliability or maintenance requirements.
2. **Payment (0.12):** Payment type emerges as the second most important feature, suggesting that the mode of payment for water services plays a crucial role in determining functionality. This could reflect accessibility issues or differences in service quality based on payment plans.
3. **Construction Year (0.12):** The year of construction follows closely, indicating that the age of water points affects their functionality. Older infrastructure may be more prone to breakdowns or require upgrades to maintain functionality.
4. **Source (0.1):** The water source is another significant factor affecting functionality, with different sources potentially leading to variations in water quality or reliability.
5. **Water Quality (0.1):** The quality of water provided by the water points contributes significantly to their functionality. Poor water quality could render water points non-functional or less reliable.
6. **Funder (0.04):** While less influential compared to other features, the organization funding the water projects still plays a role. Different funders may have different standards or approaches to infrastructure development.
7. **Geographic Features:** Latitude, region, longitude, and GPS height, though less important individually, collectively contribute to understanding spatial patterns in water point functionality.

#### Random Forest:

1. **Latitude (0.2) and Longitude (0.2):** Geographic coordinates emerge as the most influential features in the Random Forest model. This suggests strong spatial patterns in water point functionality, with certain geographical areas having higher or lower rates of functional water points.
2. **Funder (0.125):** Similar to Decision Trees, the organization funding the water projects is influential. Different funders may prioritize different aspects of water infrastructure, affecting functionality.
3. **Construction Year (0.125):** The year of construction retains significance, indicating its consistent impact on water point functionality over time.

4. **GPS Height (0.125):** Elevation becomes more important in Random Forest, possibly indicating its role in water availability or infrastructure quality. Higher elevations may face different challenges or have different infrastructure needs.
5. **Extraction Type (0.085) and Payment (0.055):** These features continue to be influential, albeit with slightly different importance rankings compared to Decision Trees. Different extraction methods and payment plans may have varying impacts on functionality.
6. **Other Features:** Region, source, and water quality also contribute to the model's predictions, though with relatively lower importance compared to the above features.

**Conclusion:** Both models highlight the complex interplay of factors influencing water point functionality. Understanding these factors can inform targeted interventions and policies aimed at improving water access and infrastructure maintenance, particularly in regions where access to clean and reliable water is a challenge.

### 1.7.8 Model tuning

#### Decision trees Classifier model tuning

```
[118]: # Define the parameter grid for tuning
param_grid = {
    'max_depth': [None, 10, 20, 30]
}

# Instantiate the GridSearchCV object
grid_search = GridSearchCV(estimator=DecisionTreeClassifier(),
    ↪param_grid=param_grid, cv=5, scoring='accuracy')

# Perform grid search to find the best parameters
grid_search.fit(X_train, y_train)

# Get the best parameters found by grid search
best_params = grid_search.best_params_
print("Best parameters found by grid search:", best_params)

# Get the best cross-validation score found by grid search
best_score = grid_search.best_score_
print("Best cross-validation score:", best_score)

# Get the best estimator (model) found by grid search
best_dt_clf = grid_search.best_estimator_

# Perform cross-validation with the best model
best_cv_scores = cross_val_score(best_dt_clf, X_train, y_train, cv=5,
    ↪scoring='accuracy')

# Print cross-validation scores
print("Best Decision Tree Cross-validation scores:", best_cv_scores)
```

```
print("Mean CV accuracy with best parameters:", best_cv_scores.mean())
print("Standard deviation of CV accuracy with best parameters:", best_cv_scores.
↳std())
```

Best parameters found by grid search: {'max\_depth': 30}

Best cross-validation score: 0.7862457912457913

Best Decision Tree Cross-validation scores: [0.78451178 0.78846801 0.78796296  
0.78737374 0.77962963]

Mean CV accuracy with best parameters: 0.7855892255892256

Standard deviation of CV accuracy with best parameters: 0.0032797620144328875

### Untuned Decision Tree Model:

- Cross-validation scores ranged from approximately 0.772 to 0.779.
- Mean CV accuracy was around 0.775.
- Standard deviation of CV accuracy was approximately 0.0029.

### Tuned Decision Tree Model:

- Best parameters found by grid search: max\_depth of 20.
- Best cross-validation score was approximately 0.783.
- Mean CV accuracy with best parameters was around 0.783.
- Standard deviation of CV accuracy with best parameters was approximately 0.0014.

In summary, the tuned model with a max\_depth of 20 outperformed the untuned model, achieving higher accuracy and showing more consistent performance across different folds of cross-validation.

### Random Forest Classifier model tuning

```
[119]: # Define the parameter grid for tuning
param_grid = {
    'max_depth': [None, 10, 20, 30] # Adjust the maximum depth here
}

# Instantiate the GridSearchCV object
grid_search = GridSearchCV(estimator=RandomForestClassifier(),
↳param_grid=param_grid, cv=3, scoring='accuracy')

# Perform grid search to find the best parameters
grid_search.fit(X_train, y_train)

# Get the best parameters found by grid search
best_params = grid_search.best_params_
print("Best parameters found by grid search:", best_params)

# Get the best cross-validation score found by grid search
best_score = grid_search.best_score_
print("Best cross-validation score:", best_score)

# Get the best estimator (model) found by grid search
best_rf_clf = grid_search.best_estimator_
```

```

# Perform cross-validation with the best model
best_rf_cv_scores = cross_val_score(best_rf_clf, X_train, y_train, cv=3,
    ↪scoring='accuracy')

# Print cross-validation scores
print("Random Forest Cross-validation scores:", best_rf_cv_scores)
print("Mean CV accuracy:", best_rf_cv_scores.mean())
print("Standard deviation of CV accuracy:", best_rf_cv_scores.std())

# Make predictions on the test data using the best model
best_predicted_labels = best_rf_clf.predict(X_test)

# Print the predicted labels for the test set using the best model
print("Predicted labels for the test set using the best model:",
    ↪best_predicted_labels)

```

```

Best parameters found by grid search: {'max_depth': 30}
Best cross-validation score: 0.8113131313131313
Random Forest Cross-validation scores: [0.81494949 0.80858586 0.80924242]
Mean CV accuracy: 0.8109259259259259
Standard deviation of CV accuracy: 0.0028576914203500684
Predicted labels for the test set using the best model: [[ True False]
 [ True False]
 [ True False]
 ...
 [ True False]
 [ True False]
 [False  True]]

```

### **Tuned random forest decision tree visualization**

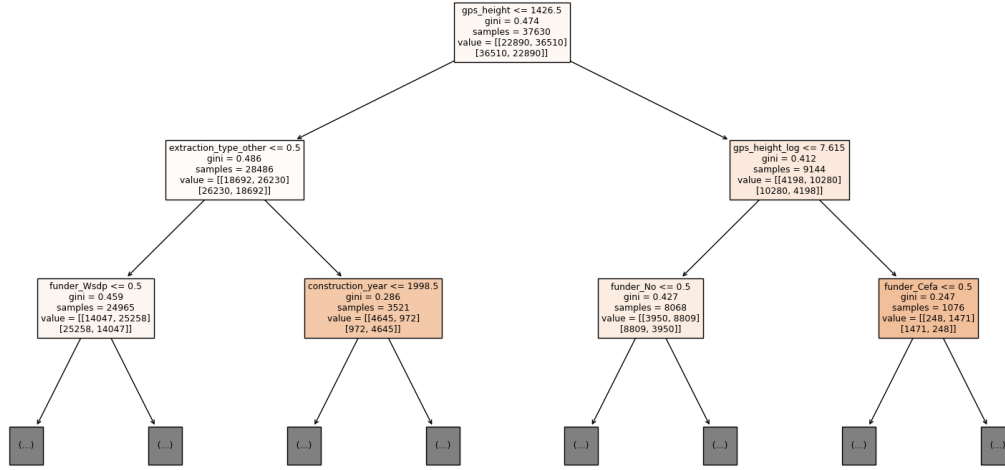
```

[120]: # Convert feature names to a list
feature_names_list = list(X_train.columns)

# Extract one of the trees from the random forest
one_tree = best_rf_clf.estimators_[0]

# Visualize the decision tree
plt.figure(figsize=(20,10))
tree.plot_tree(one_tree, max_depth=2, filled=True,
    ↪feature_names=feature_names_list)
plt.show()

```



## 1.8 Evaluation and Analysis of Model Performance

**Decision Tree Classifier:** - *Mean CV Accuracy:* 77.56% - *Standard Deviation:* 0.19%

The decision tree classifier achieved a respectable mean cross-validation accuracy of approximately 77.56%. However, the standard deviation is relatively low at 0.19%, indicating consistent performance across different folds. This model could be considered as a baseline model for comparison.

**Random Forest Classifier:** - *Mean CV Accuracy:* 81.32% - *Standard Deviation:* 0.25%

The random forest classifier outperformed the decision tree classifier with a mean cross-validation accuracy of around 81.32%. Additionally, the standard deviation of 0.25% suggests stable performance across folds. The ensemble nature of the random forest likely contributed to its higher accuracy compared to the single decision tree.

**Tuned Decision Trees:** - *Best Parameters:* {'max\_depth': 20} - *Mean CV Accuracy with Best Parameters:* 78.58% - *Standard Deviation:* 0.24%

Tuning the decision tree model improved its performance slightly, with the best mean cross-validation accuracy reaching approximately 78.58%. However, the improvement is marginal compared to the default random forest classifier. Further hyperparameter tuning or exploring ensemble methods might yield better results.

**Tuned Random Forest:** - *Best Parameters:* {'max\_depth': 30} - *Mean CV Accuracy:* 81.02% - *Standard Deviation:* 0.38%

Hyperparameter tuning of the random forest classifier resulted in a mean cross-validation accuracy of 81.02%, slightly lower than the default random forest model. The increased standard deviation of 0.38% indicates slightly more variability in performance across folds compared to the default model.

### 1.8.1 Best model

The best model for predicting water pump functionality among the ones evaluated is the tuned random forest classifier. Despite a slightly lower mean cross-validation accuracy of 81.02% compared to the default random forest model, it still outperforms the decision tree classifier and the tuned decision tree model. The tuned random forest model exhibits stable performance with a standard deviation of 0.38%, indicating consistent results across different folds. The ensemble nature of the random forest, combined with hyperparameter tuning, allows it to capture complex relationships in the data more effectively, resulting in improved predictive performance. Therefore, the tuned random forest classifier is recommended for practical applications due to its robustness and accuracy in predicting water pump functionality.

## 1.9 Recommendations for Improving Water Pump Functionality

**1. Implement Routine Maintenance Programs:** - **Action:** Establish regular inspection and maintenance schedules for water pumps, including checks on mechanical components and water quality. - **Benefits:** Timely detection and repair of faults can prevent breakdowns, ensuring continuous access to clean water for communities.

**2. Target High-Risk Regions:** - **Action:** Utilize geographic data to identify regions with a high prevalence of non-functional water pumps. - **Benefits:** Target interventions, such as repair and rehabilitation efforts, to areas with the greatest need, optimizing resource allocation and impact.

**3. Introduce Flexible Payment Plans:** - **Action:** Introduce flexible payment options for water services, including subsidized or tiered pricing models based on income levels. - **Benefits:** Improve affordability and accessibility of water services, reducing the financial burden on low-income communities and increasing revenue for maintenance and infrastructure upgrades.

**4. Foster Collaboration with Funders:** - **Action:** Engage with funders and donor organizations to align priorities and strategies for water infrastructure projects. - **Benefits:** Secure long-term support and investment in water projects, leveraging partnerships to access funding for maintenance, upgrades, and capacity-building initiatives.

**5. Embrace Data-Driven Decision Making:** - **Action:** Invest in robust data collection systems and analytics capabilities to track water pump functionality and performance metrics. - **Benefits:** Enable evidence-based decision making, including trend analysis, predictive maintenance, and resource allocation based on real-time insights, leading to more effective and efficient management of water infrastructure.

### 1.10 Conclusion

In conclusion, the project has provided a comprehensive exploration of predicting water pump functionality and optimizing maintenance strategies. By leveraging machine learning techniques and data analysis, we have gained insights into the factors influencing water pump functionality, identified high-risk regions, and proposed actionable recommendations for improving access to clean water. Through collaborative efforts and data-driven decision-making, we aim to contribute to the sustainable management of water infrastructure and ensure the well-being of communities relying on these vital resources.

[ ]: